


# Computer Architecture

Lec 4a: Single-Cycle Datapath

# This Lecture: Single-Cycle Datapath


---



- Overview of ISAs
- Datapath storage elements
- MIPS Datapath
- MIPS Control


# Review: ISA

---



- App/OS are software ... execute on hardware
- HW/SW interface is **ISA (instruction set architecture)**
  - A **"contract"** between SW and HW
  - Encourages compatibility, allows SW/HW to evolve independently
  - **Functional definition** of HW storage locations & operations
    - Storage locations: registers, memory
    - Operations: add, multiply, branch, load, store, etc.
  - **Precise description** of how to invoke & access them
    - Instructions (bit-patterns hardware interprets as commands)

# MIPS: A Real ISA





- **MIPS**: example of real ISA
  - 32/64-bit operations
  - 32-bit insns
  - 64 registers
    - 32 integer, 32 floating point
  - ~100 different insns
  - Full OS support

**Example code is MIPS, but  
all ISAs are pretty similar**

```
.data
array: .space 100
sum:   .word 0
.text

array_sum:
    li $5, 0
    la $1, array
    la $2, sum
array_sum_loop:
    lw $3, 0($1)
    lw $4, 0($2)
    add $4, $3, $4
    sw $4, 0($2)
    addi $1, $1, 1
    addi $5, $5, 1
    li $6, 100
    blt $5, $6, array_sum_loop
```

# Review: Insn Execution Model

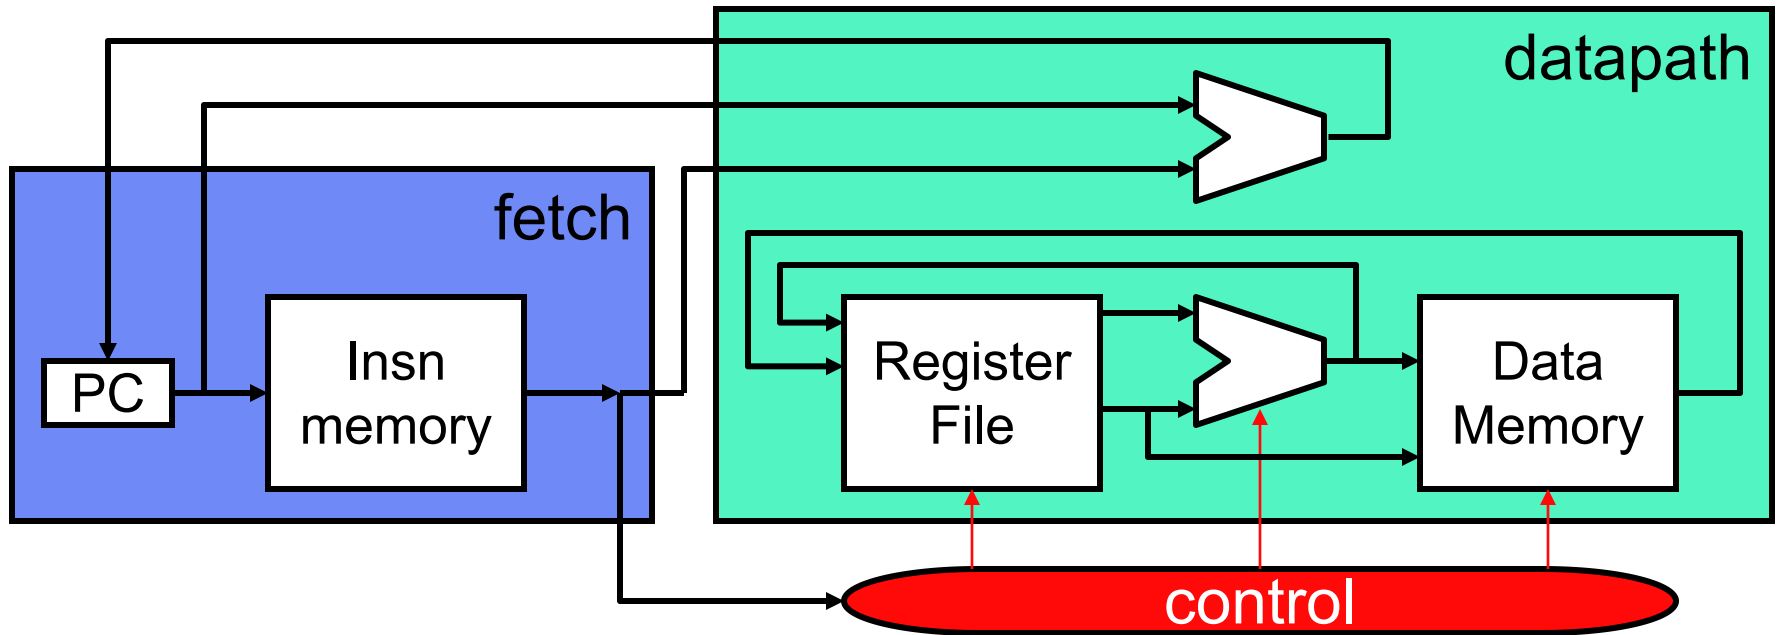


Instruction → Insn

- The computer is just finite state machine
  - **Registers** (few of them, but fast)
  - **Memory** (lots of memory, but slower)
  - **Program counter** (next insn to execute)
    - Sometimes called “instruction pointer”
- A computer executes **instructions**
  - **Fetches** next instruction from memory
  - **Decodes** it (figure out what it does)
  - **Reads** its **inputs** (registers & memory)
  - **Executes** it (adds, multiply, etc.)
  - **Write** its **outputs** (registers & memory)
  - **Next insn** (adjust the program counter)
- **Program is just “data in memory”**
  - Makes computers programmable (“universal”)

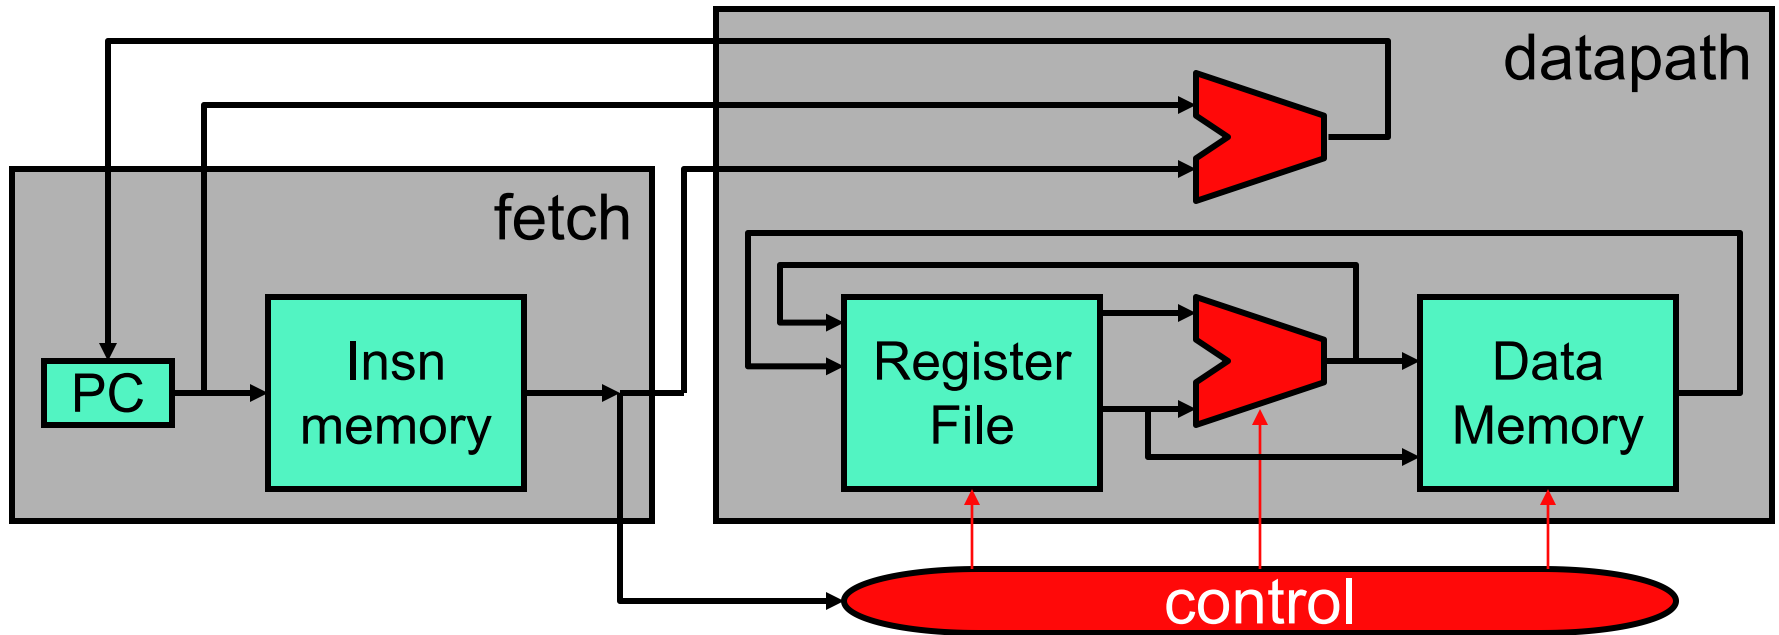
# Implementing an ISA

# Implementing an ISA



- **Datapath**: performs computation (registers, ALUs, etc.)
  - ISA specific: can implement every insn (single-cycle: in one pass!)
- **Control**: determines which computation is performed
  - Routes data through datapath (which regs, which ALU op)
- **Fetch**: get insn, translate opcode into control
- **Fetch** → **Decode** → **Execute** "cycle"

# Two Types of Components



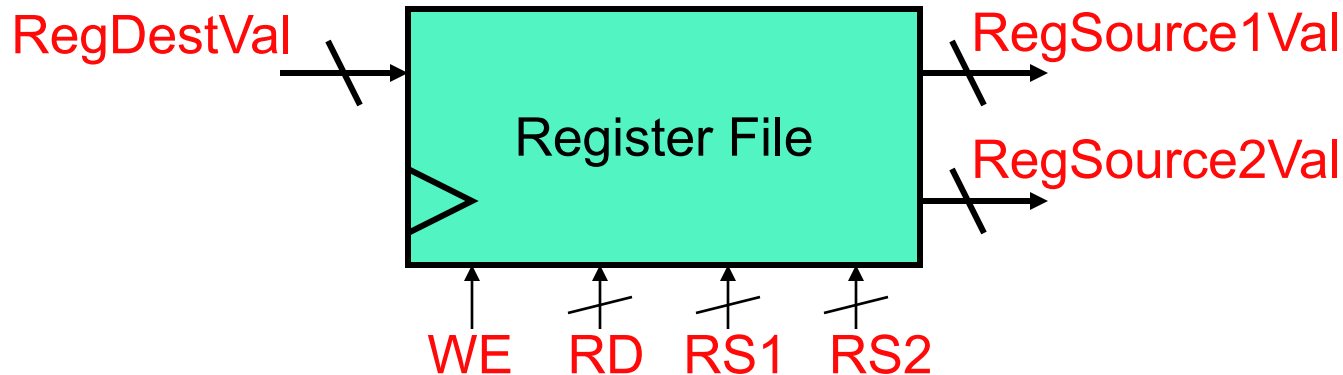
- **Purely combinational**: stateless computation
  - ALUs, muxes, control
  - Arbitrary Boolean functions
- **Combinational/sequential**: storage
  - PC, insn/data memories, register file
  - Internally contain some combinational components



# Datapath Storage Elements

# Register File


---



- **Register file**: M N-bit storage words
  - Multiplexed input/output: data buses write/read "random" word
- **"Port"**: set of buses for accessing a random word in array
  - Data bus (N-bits) + address bus ( $\log_2 M$ -bits) + optional WE bit
  - P ports = P parallel and independent accesses
- MIPS integer register file
  - 32 32-bit words, two read ports + one write port

# Decoder

- **Decoder**: converts binary integer to “1-hot” representation
  - Binary representation of  $0 \dots 2^N - 1$ :  $N$  bits
  - 1 hot representation of  $0 \dots 2^N - 1$ :  $2^N$  bits
    - $J$  represented as  $J^{\text{th}}$  bit 1, all other bits zero
  - Example below: 2-to-4 decoder



# Decoder in Verilog (1 of 2)

---

```
module decoder_2_to_4 (binary_in, onehot_out);  
    input [1:0] binary_in;  
    output [3:0] onehot_out;  
    assign onehot_out[0] = (~binary_in[0] & ~binary_in[1]);  
    assign onehot_out[1] = (~binary_in[0] & binary_in[1]);  
    assign onehot_out[2] = (binary_in[0] & ~binary_in[1]);  
    assign onehot_out[3] = (binary_in[0] & binary_in[1]);  
endmodule
```

- Is there a simpler way?

# Decoder in Verilog (2 of 2)

---

```
module decoder_2_to_4 (binary_in, onehot_out);  
    input [1:0] binary_in;  
    output [3:0] onehot_out;  
    assign onehot_out[0] = (binary_in == 2'd0);  
    assign onehot_out[1] = (binary_in == 2'd1);  
    assign onehot_out[2] = (binary_in == 2'd2);  
    assign onehot_out[3] = (binary_in == 2'd3);  
endmodule
```

# Register File Interface


---



- Inputs:
  - RS1, RS2 (reg. sources to read), RD (reg. destination to write)
  - WE (write enable), RDestVal (value to write)
- Outputs: RSrc1Val, RSrc2Val (value of RS1 & RS2 registers)


# Register File: Four Registers

---



- Register file with four registers

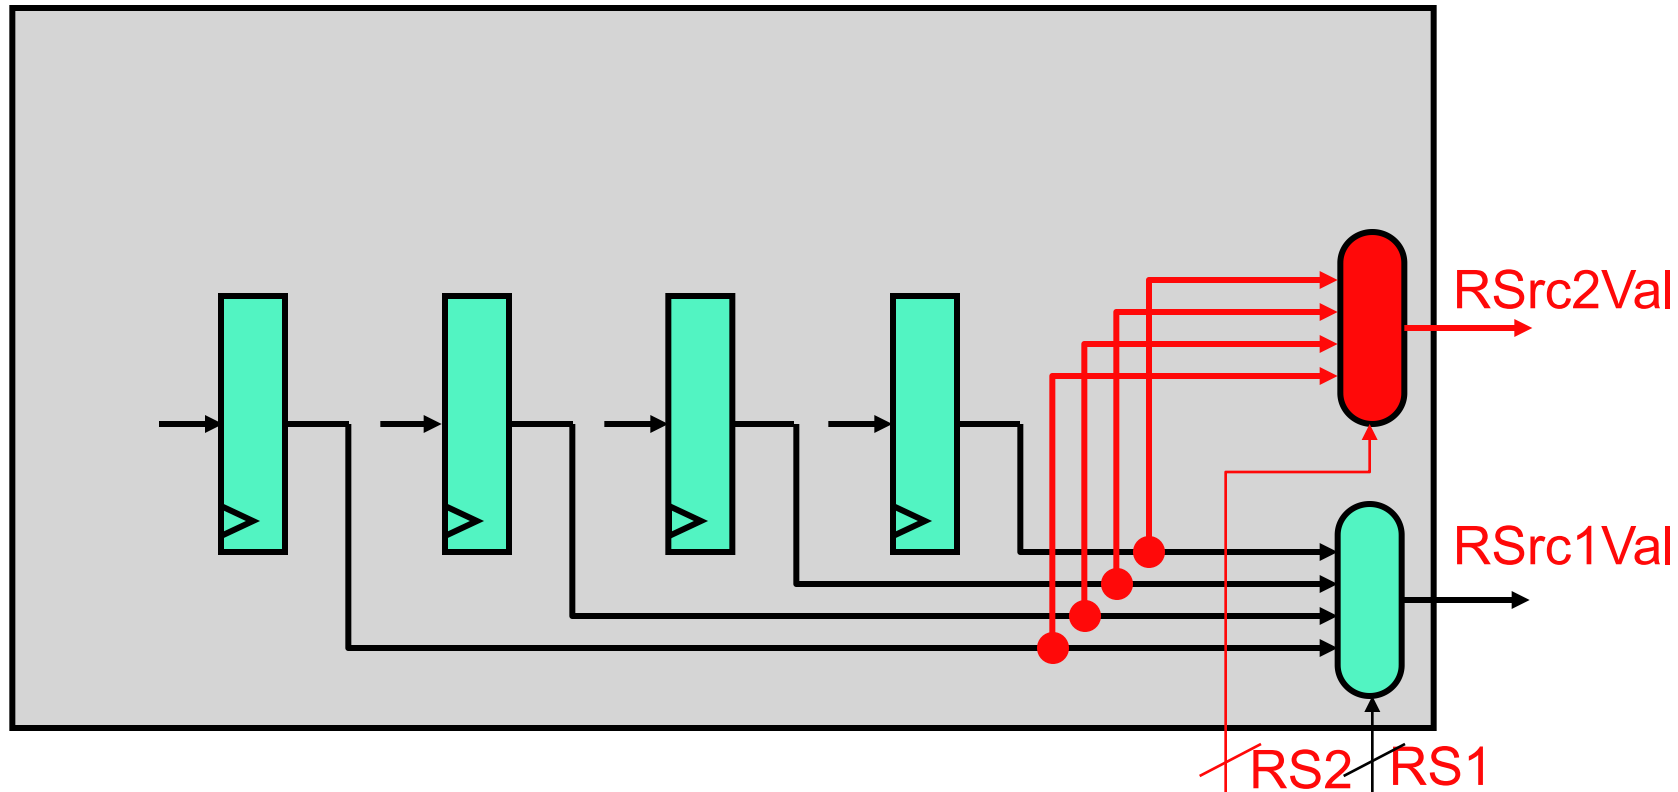
# Add a Read Port



- Output of each register into 4to1 mux (RSrc1Val)
  - RS1 is select input of RSrc1Val mux




# Add Another Read Port



- Output of each register into another 4to1 mux (RSrc2Val)
  - RS2 is select input of RSrc2Val mux


# Add a Write Port



- Input RegDestVal into each register
  - Enable only one register's WE: (Decoded RD) & (WE)
- What if we needed two write ports?

# Another Useful Component: Memory

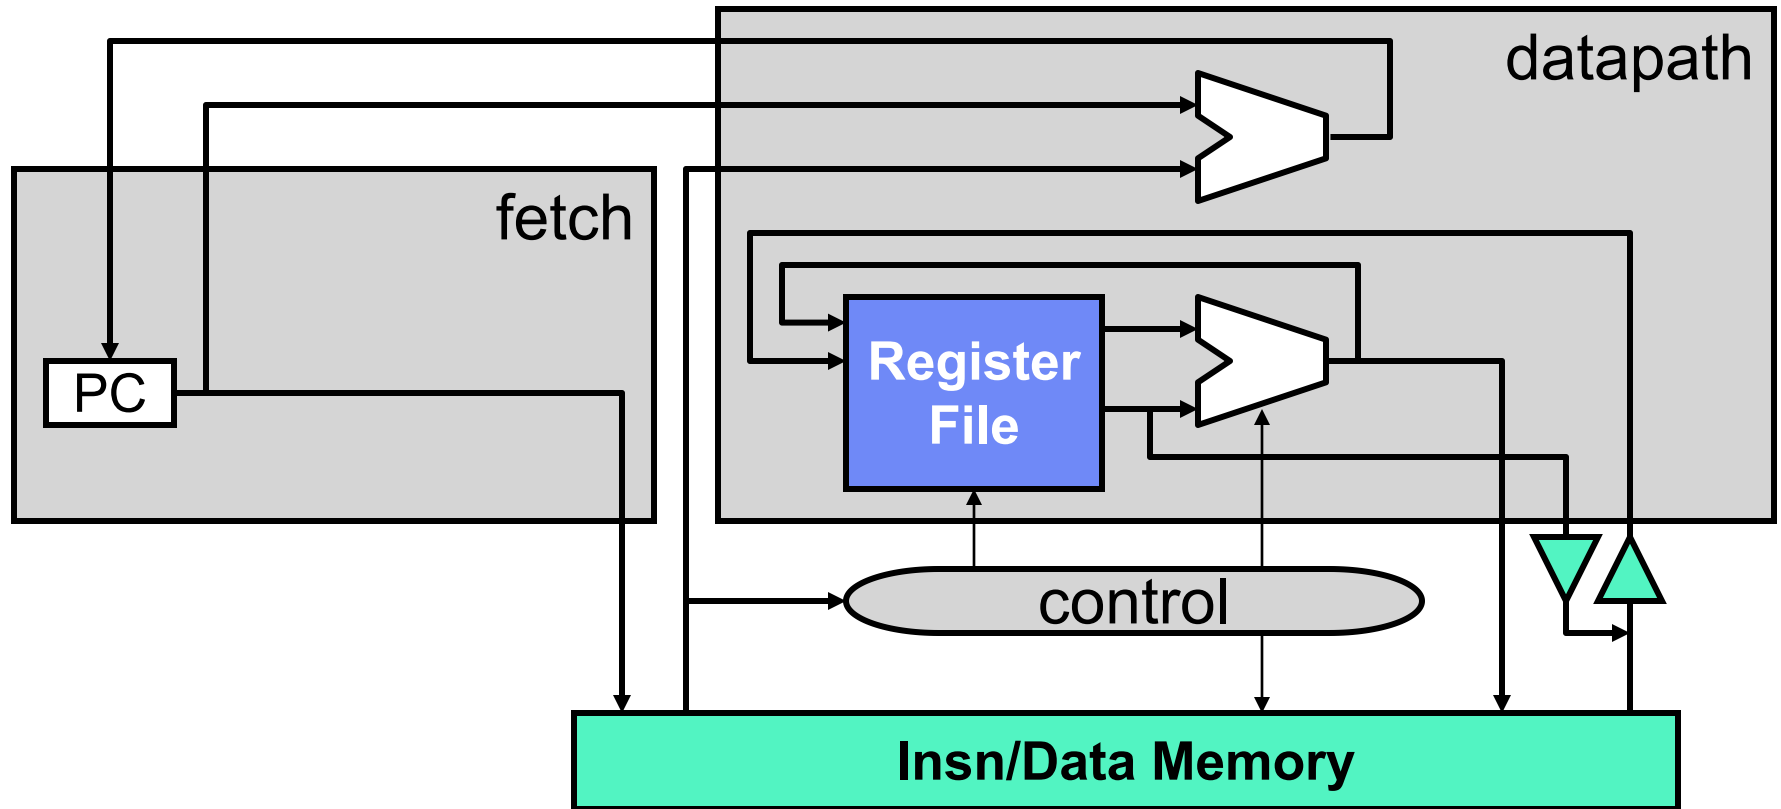
---



- Register file: M N-bit storage words
  - Few words ( $< 256$ ), many ports, dedicated read and write ports
- **Memory**: M N-bit storage words, yet not a register file
  - Many words ( $> 1024$ ), few ports (1, 2), shared read/write ports
- Leads to different implementation choices
  - Lots of circuit tricks and such
  - Larger memories typically only 6 transistors per bit

# MIPS Datapath

# Unified vs Split Memory Architecture



- **Unified architecture:** unified insn/data memory
- **“Harvard” architecture:** split insn/data memories

# Datapath for MIPS ISA

---

- MIPS: 32-bit instructions, registers are \$0, \$2... \$31


- Consider only the following instructions

<code>add \$1,\$2,\$3</code>	<code>\$1 = \$2 + \$3</code>	(add)
<code>addi \$1,\$2,3</code>	<code>\$1 = \$2 + 3</code>	(add immed)
<code>lw \$1,4(\$3)</code>	<code>\$1 = Memory[4+\$3]</code>	(load)
<code>sw \$1,4(\$3)</code>	<code>Memory[4+\$3] = \$1</code>	(store)
<code>beq \$1,\$2,PC_relative_target</code>		(branch equal)
<code>j absolute_target</code>		(unconditional jump)

- Why only these?
  - Most other instructions are the same from datapath viewpoint
  - The ones that aren't are left for you to figure out 😊

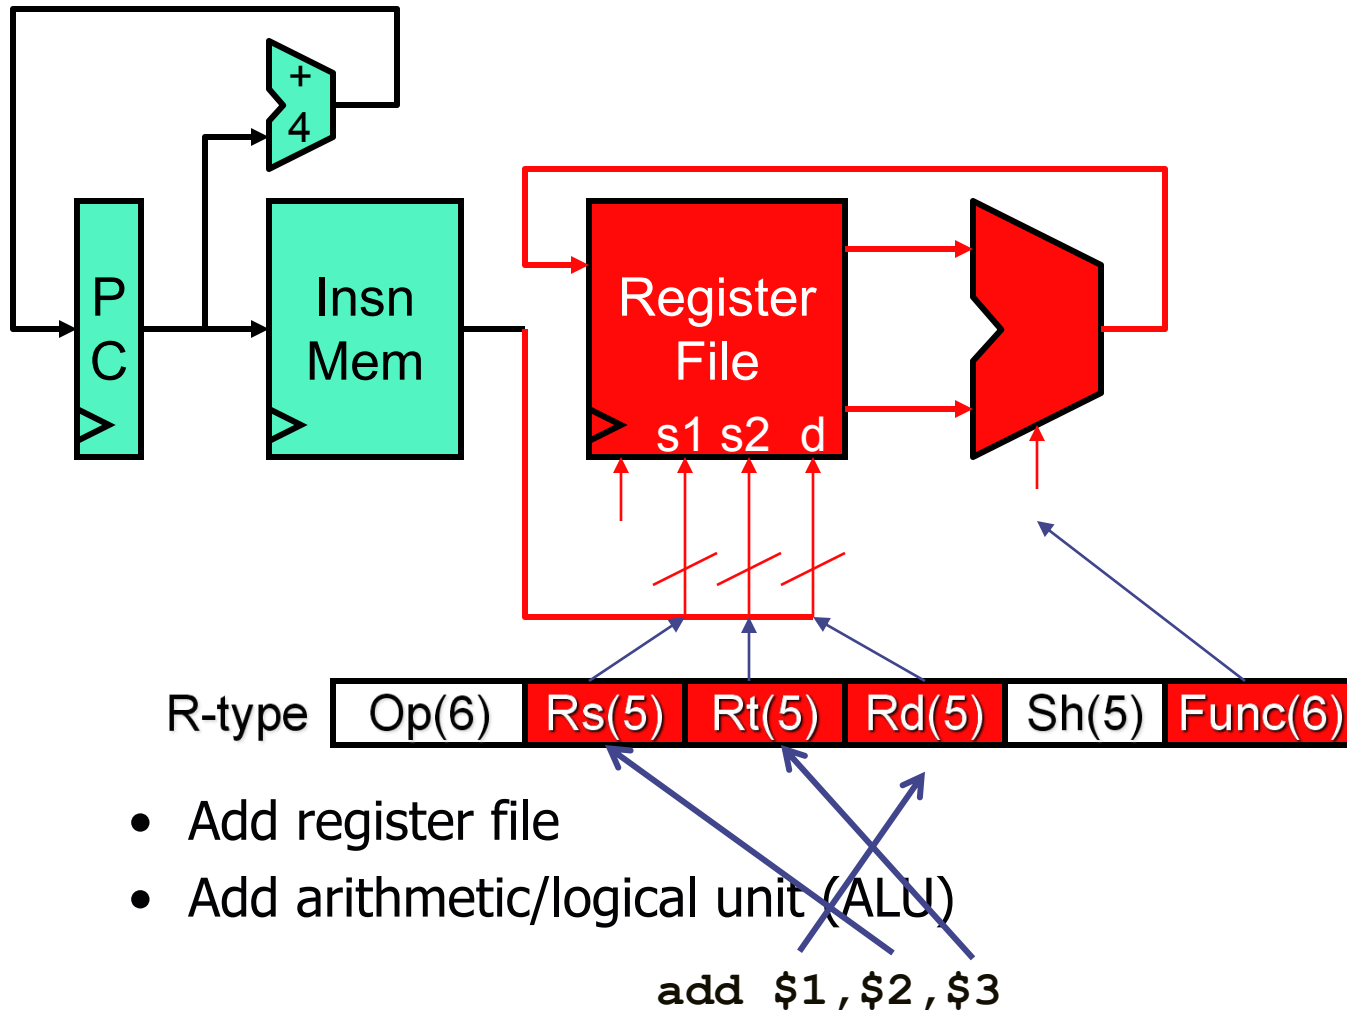
# Start With Fetch

---



- PC and instruction memory (split insn/data architecture, for now)
- A +4 incrementer computes default next instruction PC

# First Instruction: **add**






# Wire Select in Verilog

---

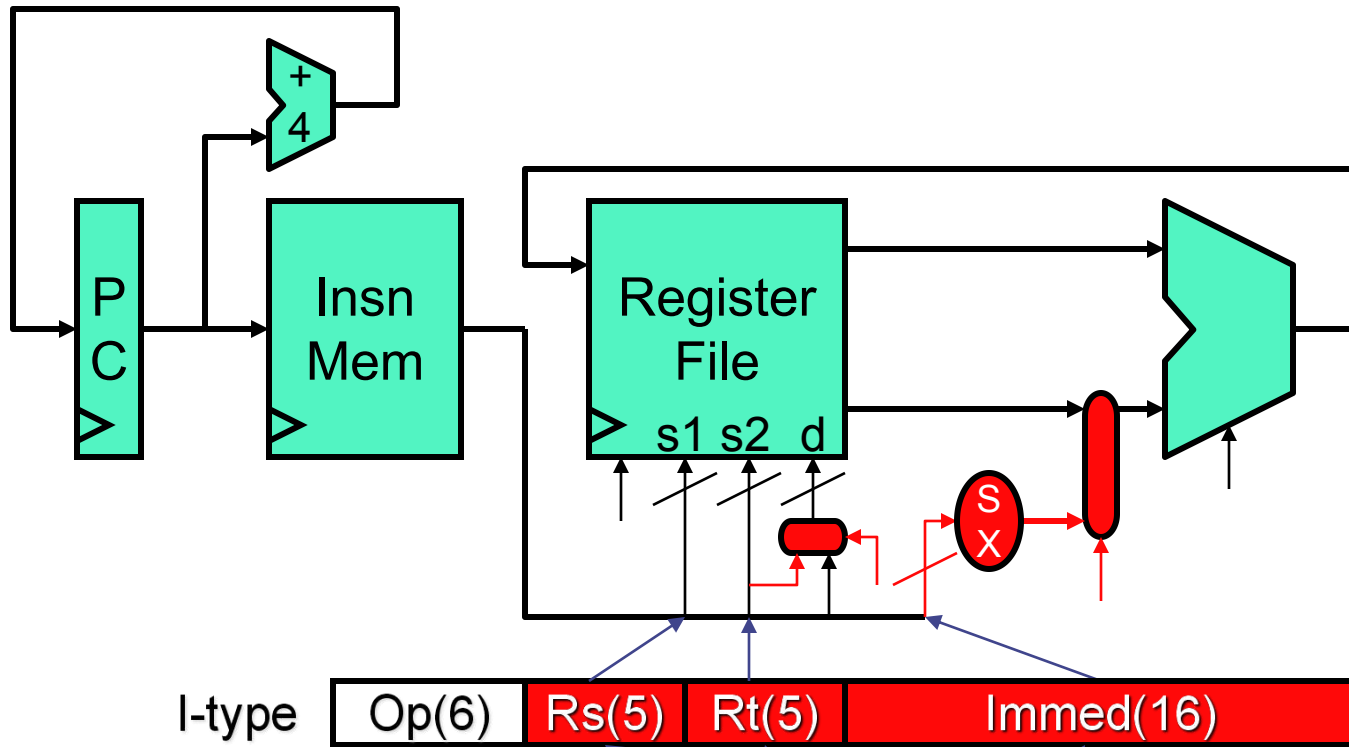
- How to pull out individual fields of an insn? **Wire select**

```
wire [31:0] insn;  
wire [5:0] op = insn[31:26];  
wire [4:0] rs = insn[25:21];  
wire [4:0] rt = insn[20:16];  
wire [4:0] rd = insn[15:11];  
wire [4:0] sh = insn[10:6];  
wire [5:0] func = insn[5:0];
```

R-type



# Second Instruction: **addi**



- Destination register can now be either Rd or Rt
- Add sign extension unit and mux into second ALU input

**addi** \$1, \$2, 8

# Verilog Wire Concatenation

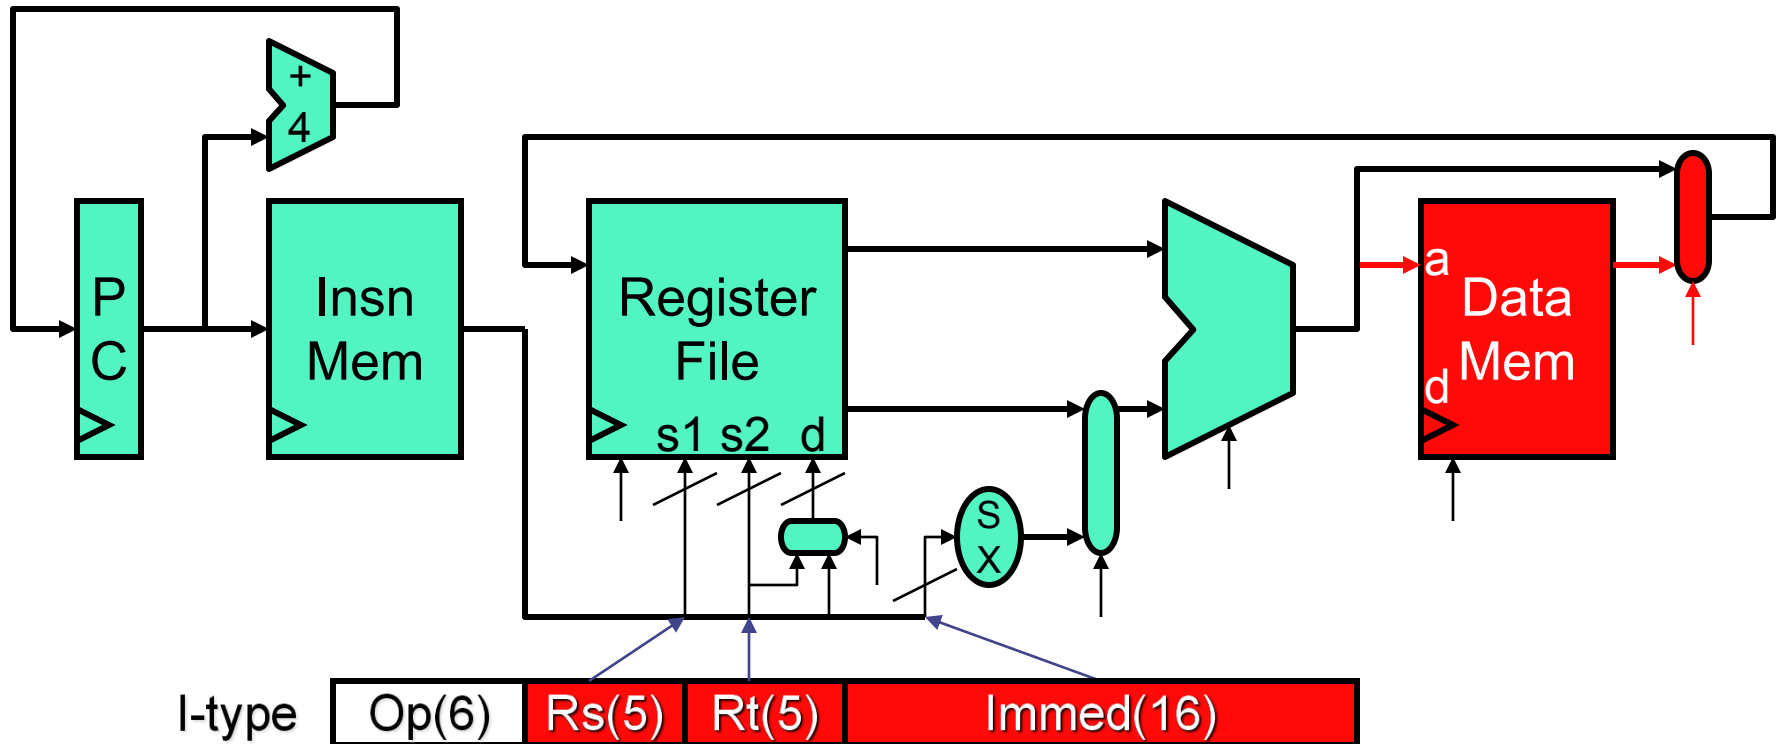
---

- Recall two Verilog constructs
  - Wire concatenation:** {bus0, bus1, ... , busn}
  - Wire repeat:** {repeat\_x\_times{w0}}
- How do you specify sign extension? **Wire concatenation**

```
wire [31:0] insn;  
wire [15:0] imm16 = insn[15:0];  
wire [31:0] sximm16 = {{16{imm16[15]}}}, imm16};
```



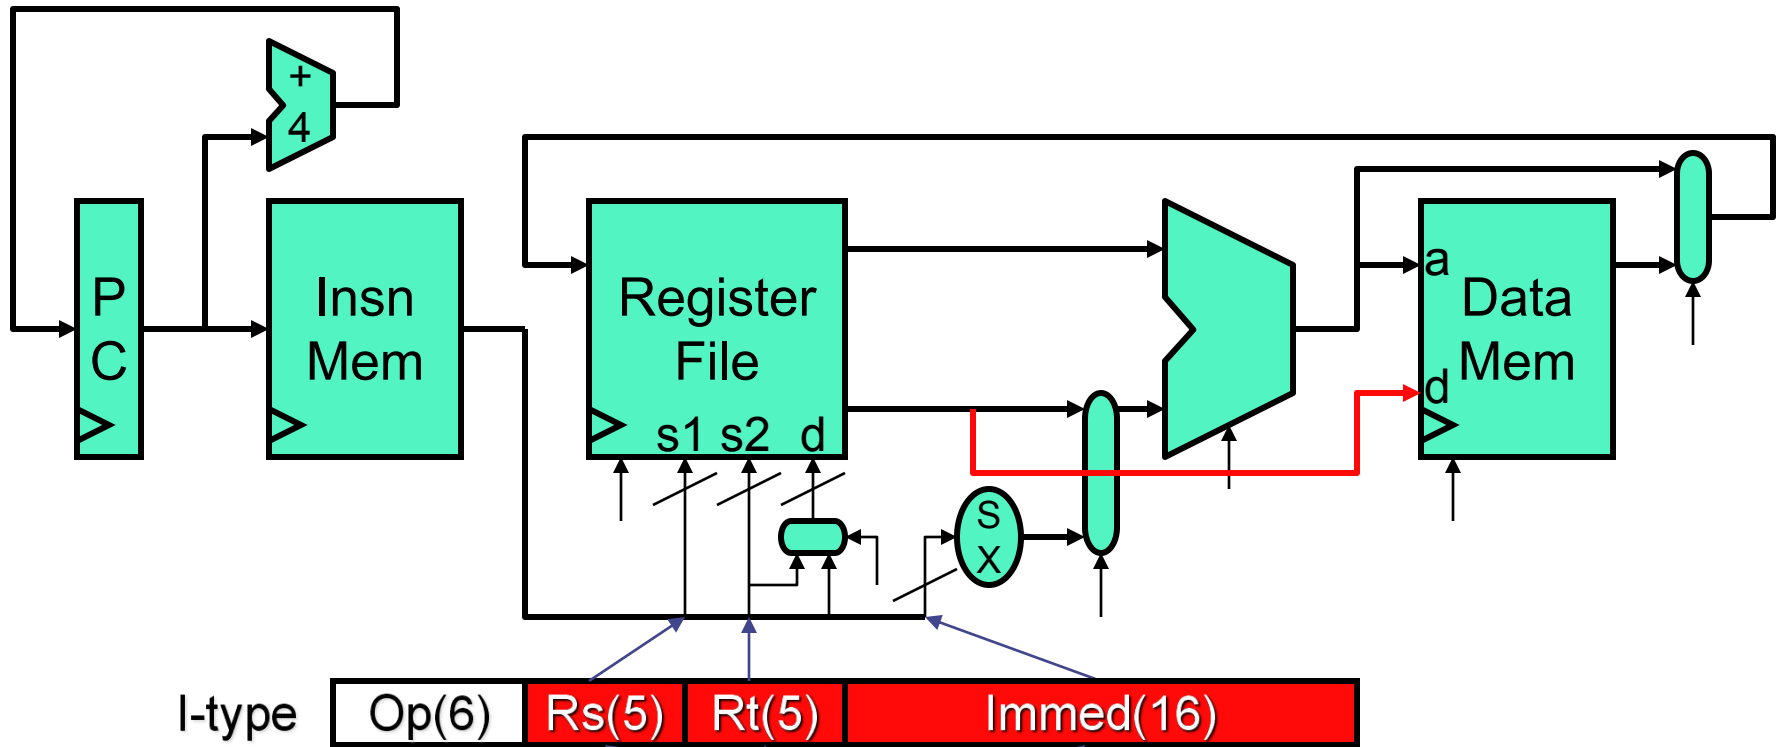
# Third Instruction: **lw**



- Add data memory, address is ALU output
- Add register write data mux to select memory output or ALU output

**lw** \$1, 8(\$2)

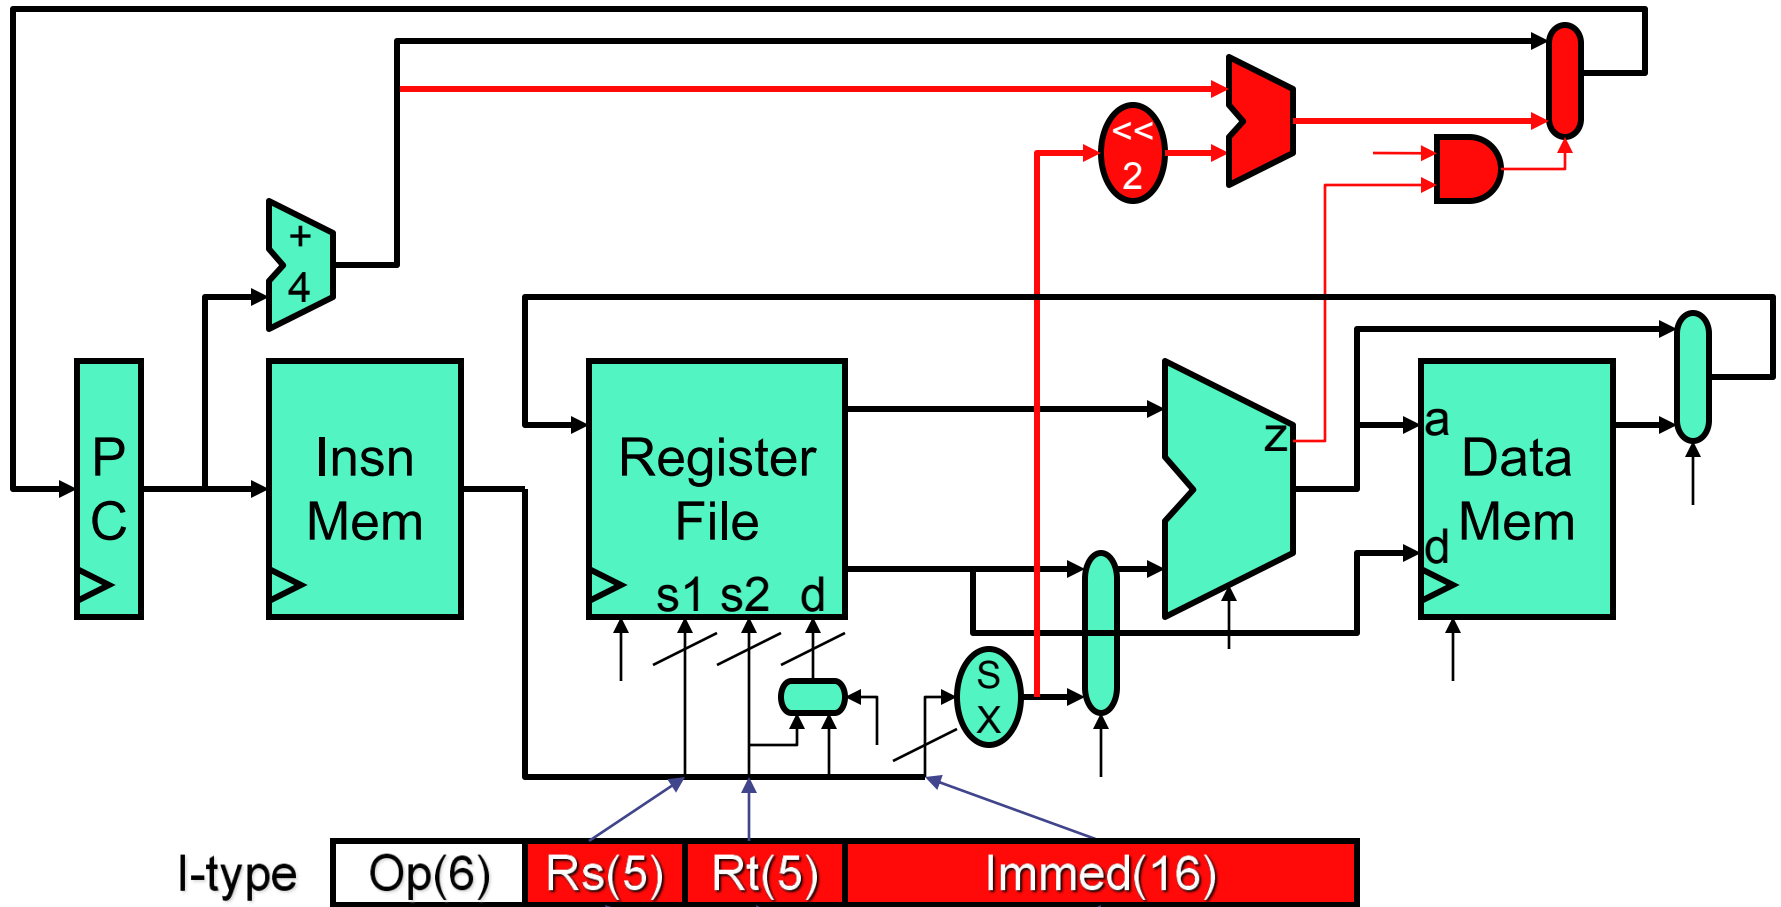
# Fourth Instruction: **SW**



- Add path from second input register to data memory data input

sw \$1, 8 (\$2)

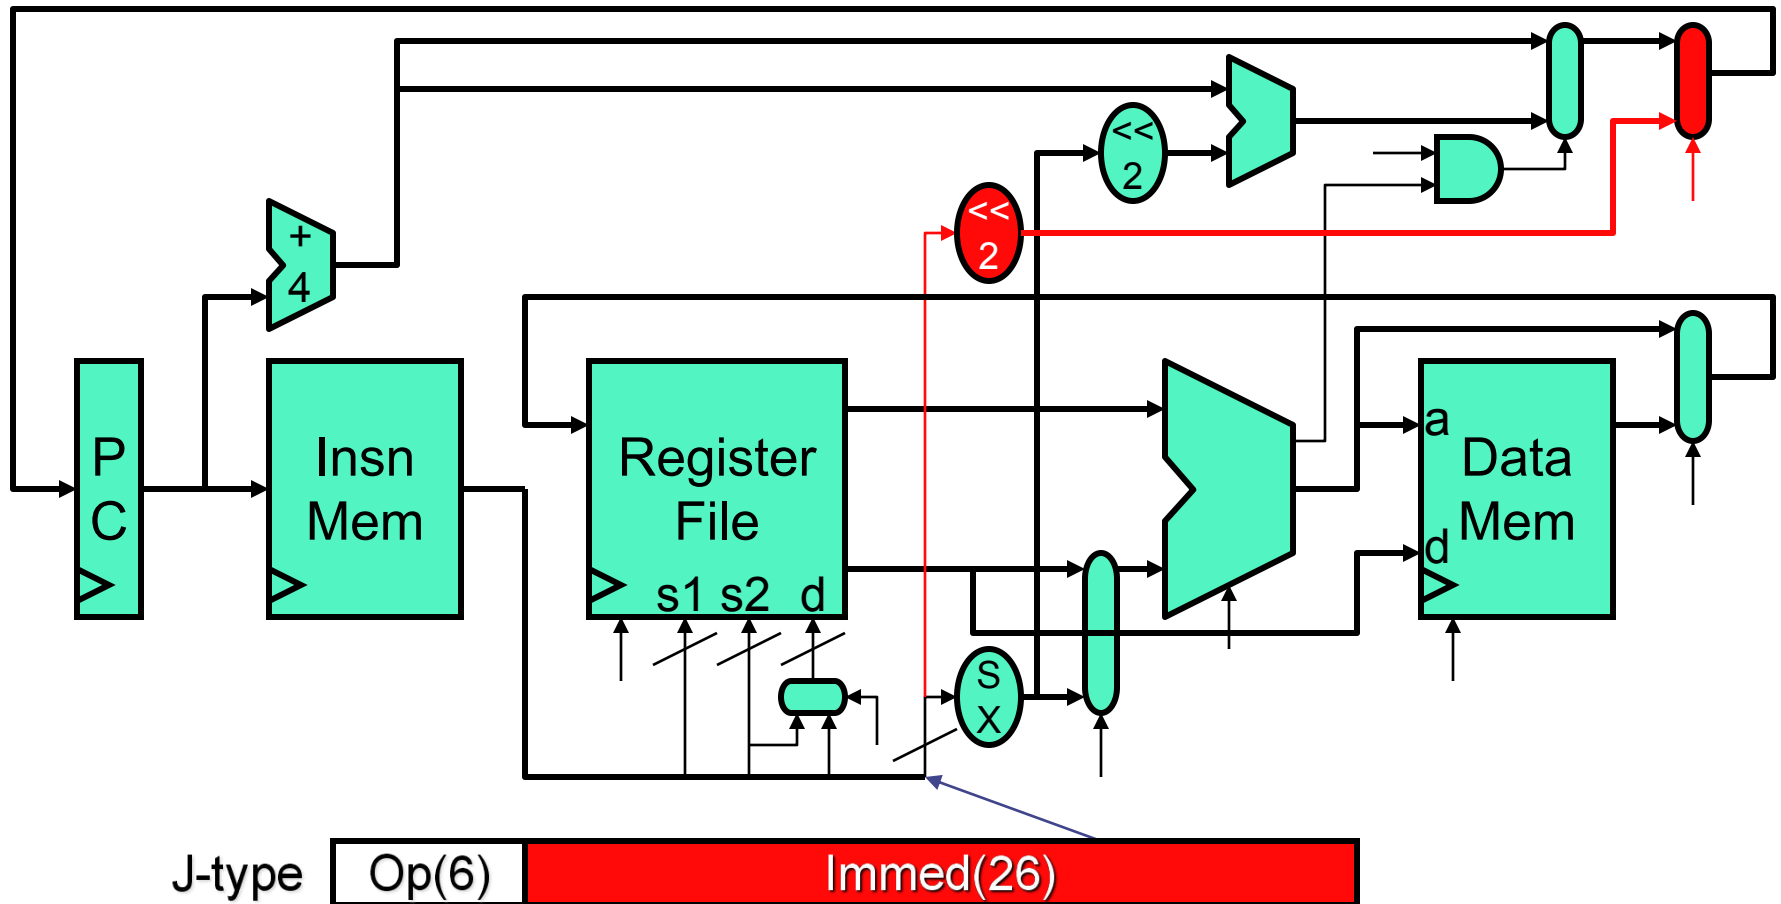
# Fifth Instruction: **beq**



- Add left shift unit and adder to compute PC-relative branch target
- Add PC input mux to select PC+4 or branch target

**beq** \$1, \$2, 8

# Sixth Instruction: **j**

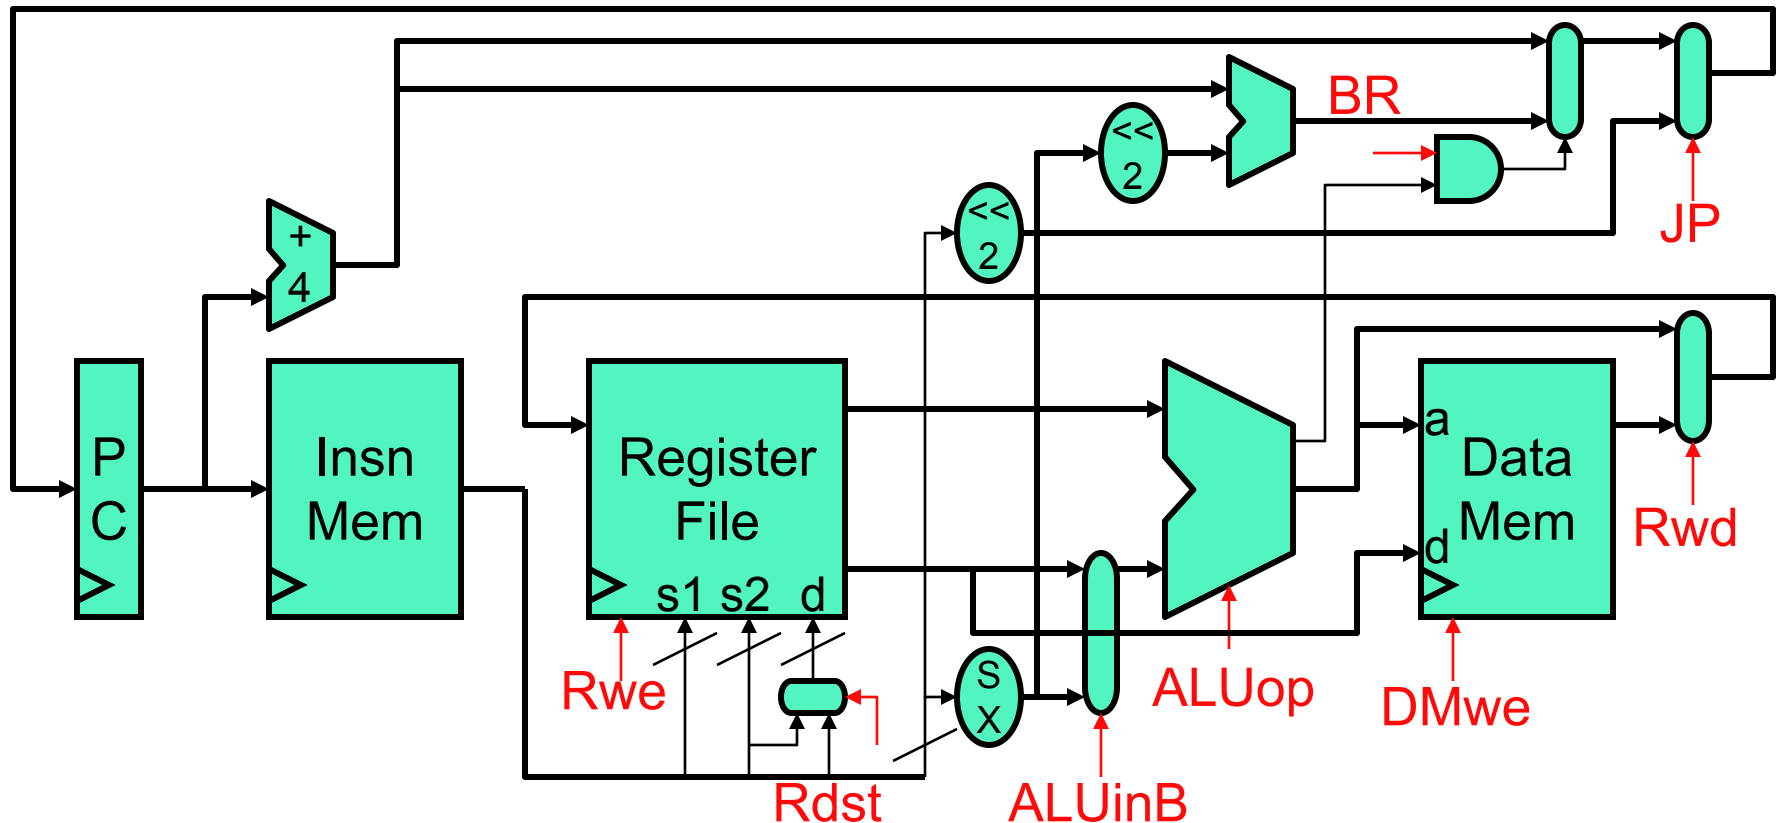


- Add shifter to compute left shift of 26-bit immediate
- Add additional PC input mux for jump target

# MIPS Control

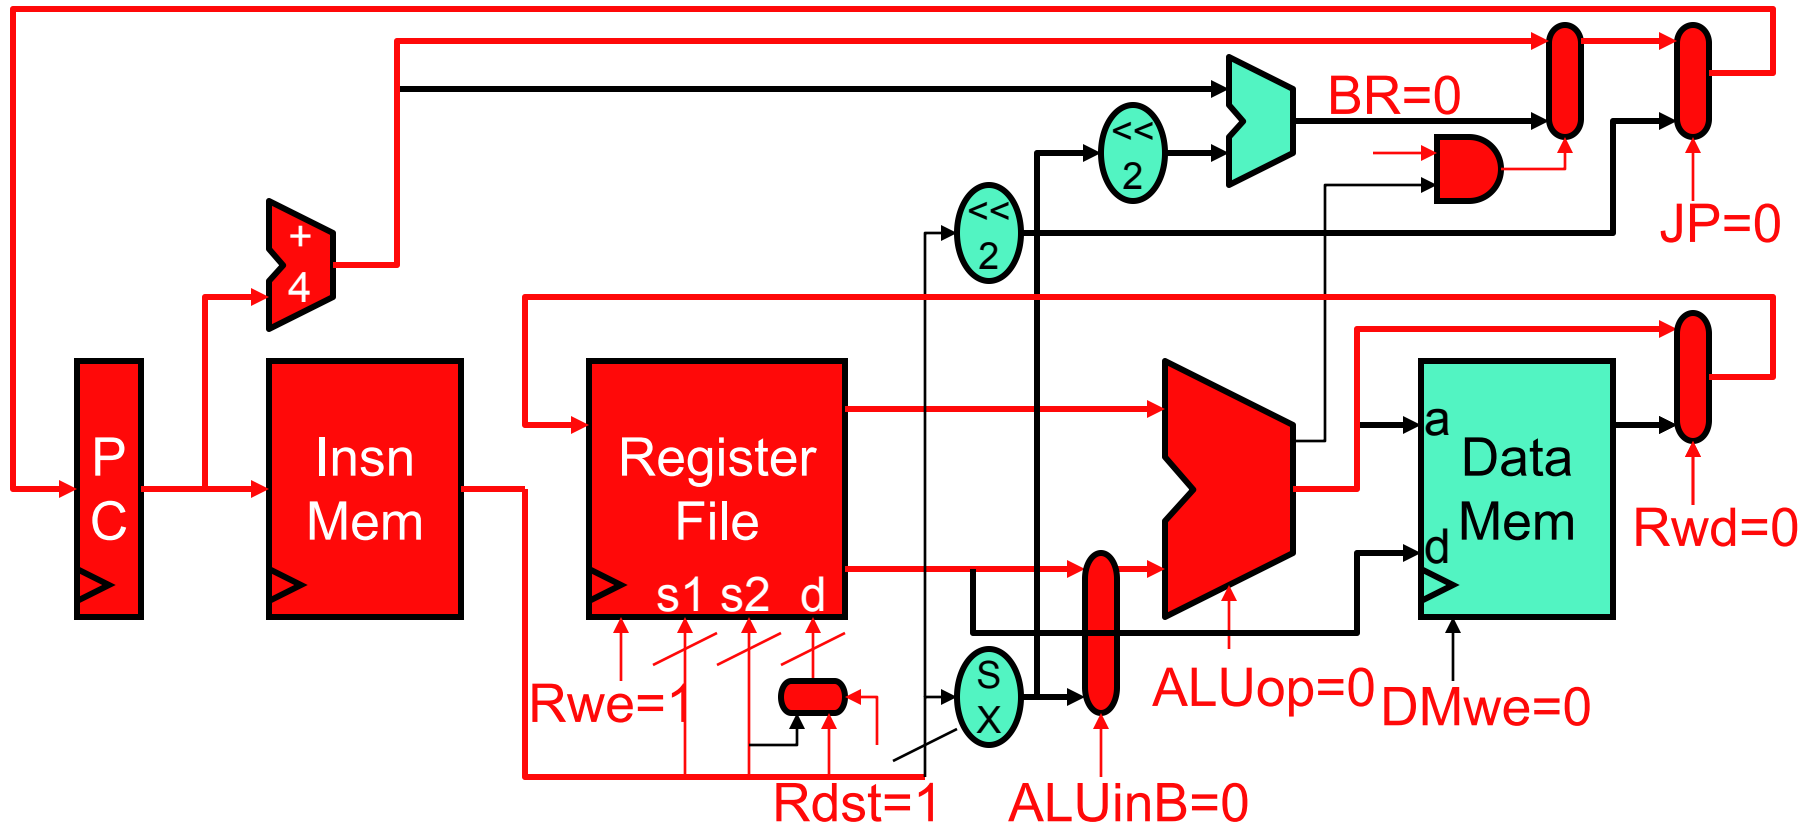


# What Is Control?



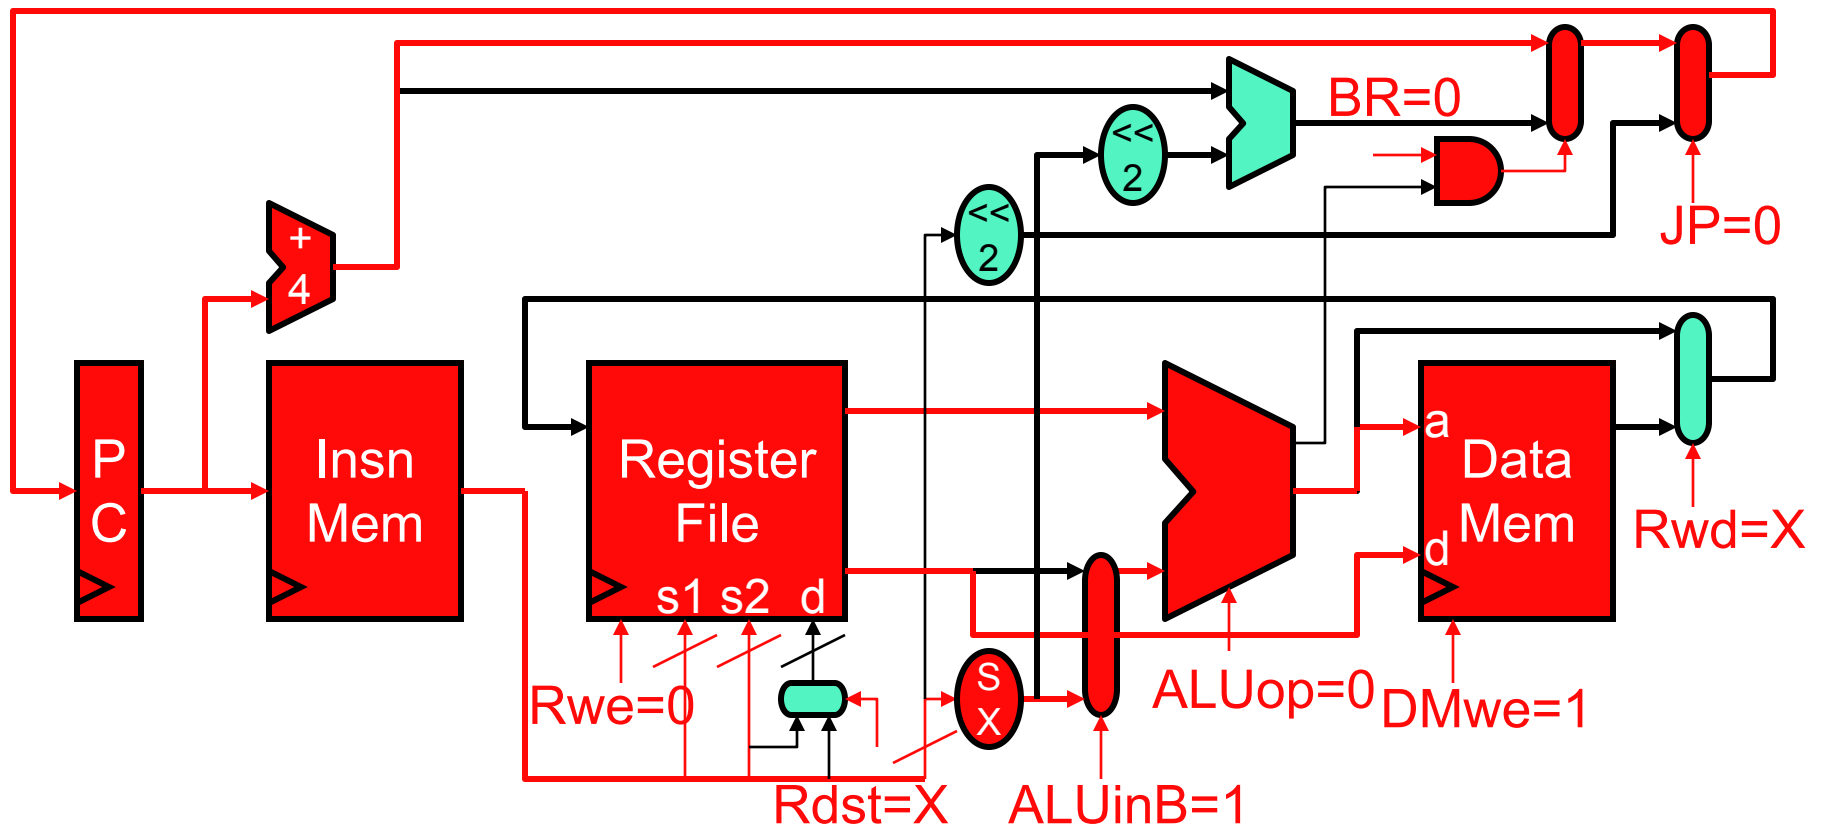
- 8 signals control flow of data through this datapath
  - MUX selectors, or register/memory write enable signals
  - A real datapath has 300-500 control signals

# Example: Control for **add**



**add \$4, \$3, \$4**

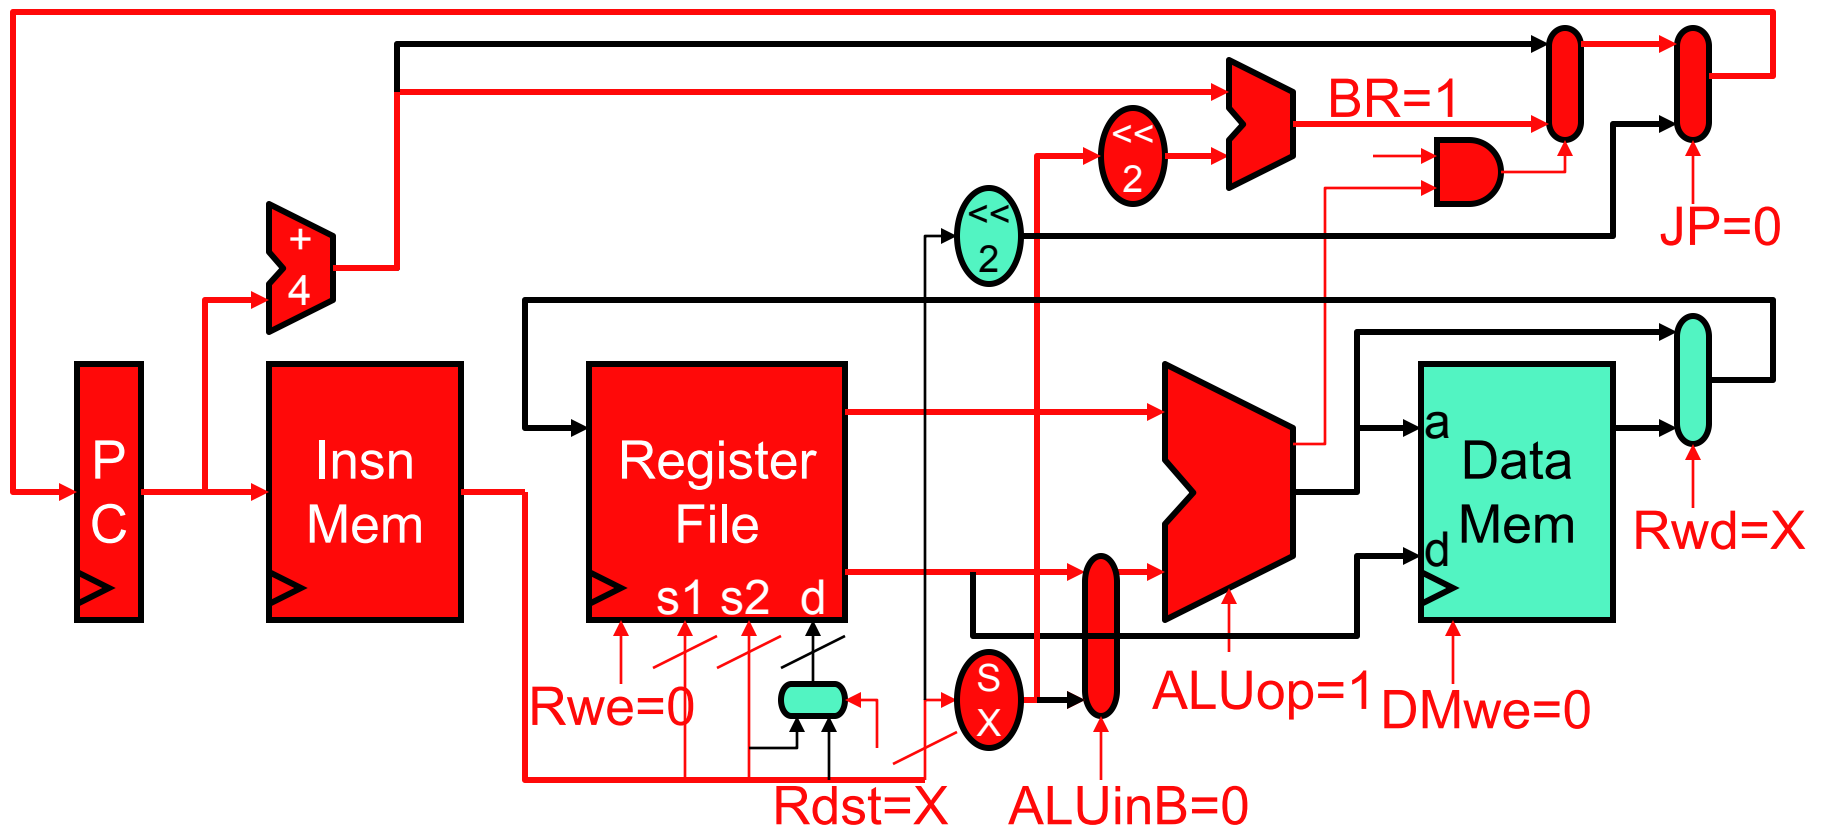
# Example: Control for **sw**



- Difference between **sw** and **add** is 5 signals
  - 3 if you don't count the X (don't care) signals

**sw** \$4, 0 (\$2)

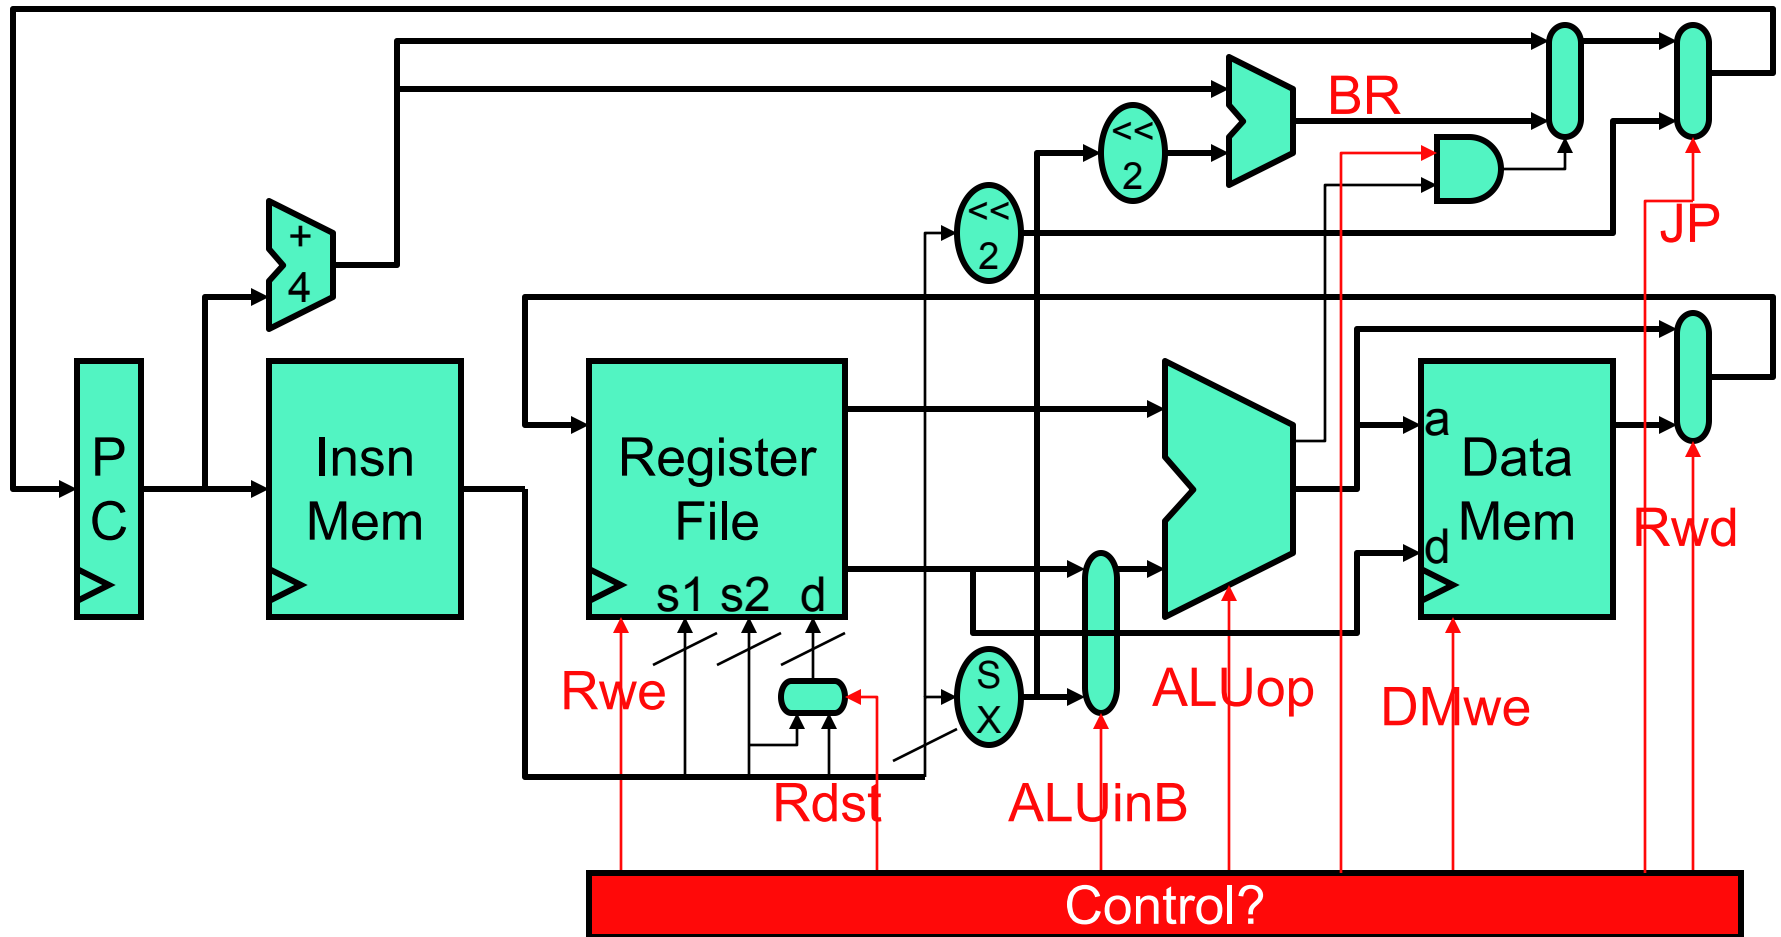
# Example: Control for **beq**



- Difference between **sw** and **beq** is only 4 signals

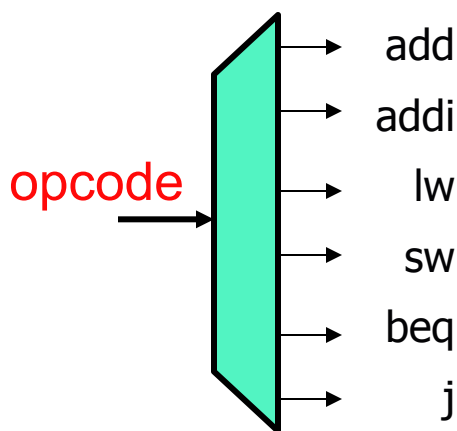
**beq** \$5, \$6, array\_sum\_loop

# How Is Control Implemented?



# Control Implementation: ROM


- **ROM (read only memory)**: like a RAM but unwritable
  - Bits in data words are control signals
  - Lines indexed by opcode
  - Example: ROM control for 6-insn MIPS datapath
  - X is “don’t care”



	BR	JP	ALUinB	ALUop	DMwe	Rwe	Rdst	Rwd
add	0	0	0	0	0	1	0	0
addi	0	0	1	0	0	1	1	0
lw	0	0	1	0	0	1	1	1
sw	0	0	1	0	1	0	X	X
beq	1	0	0	1	0	0	X	X
j	0	1	0	0	0	0	X	X


# Control Implementation: Logic

- Real machines have 100+ insns 300+ control signals
  - 30,000+ control bits (~4KB)
    - Not huge, but hard to make faster than datapath (important!)
- Alternative: **logic gates** or “random logic” (unstructured)
  - Exploits the observation: many signals have few 1s or few 0s
  - Example: random logic control for 6-insn MIPS datapath



# Control Logic in Verilog

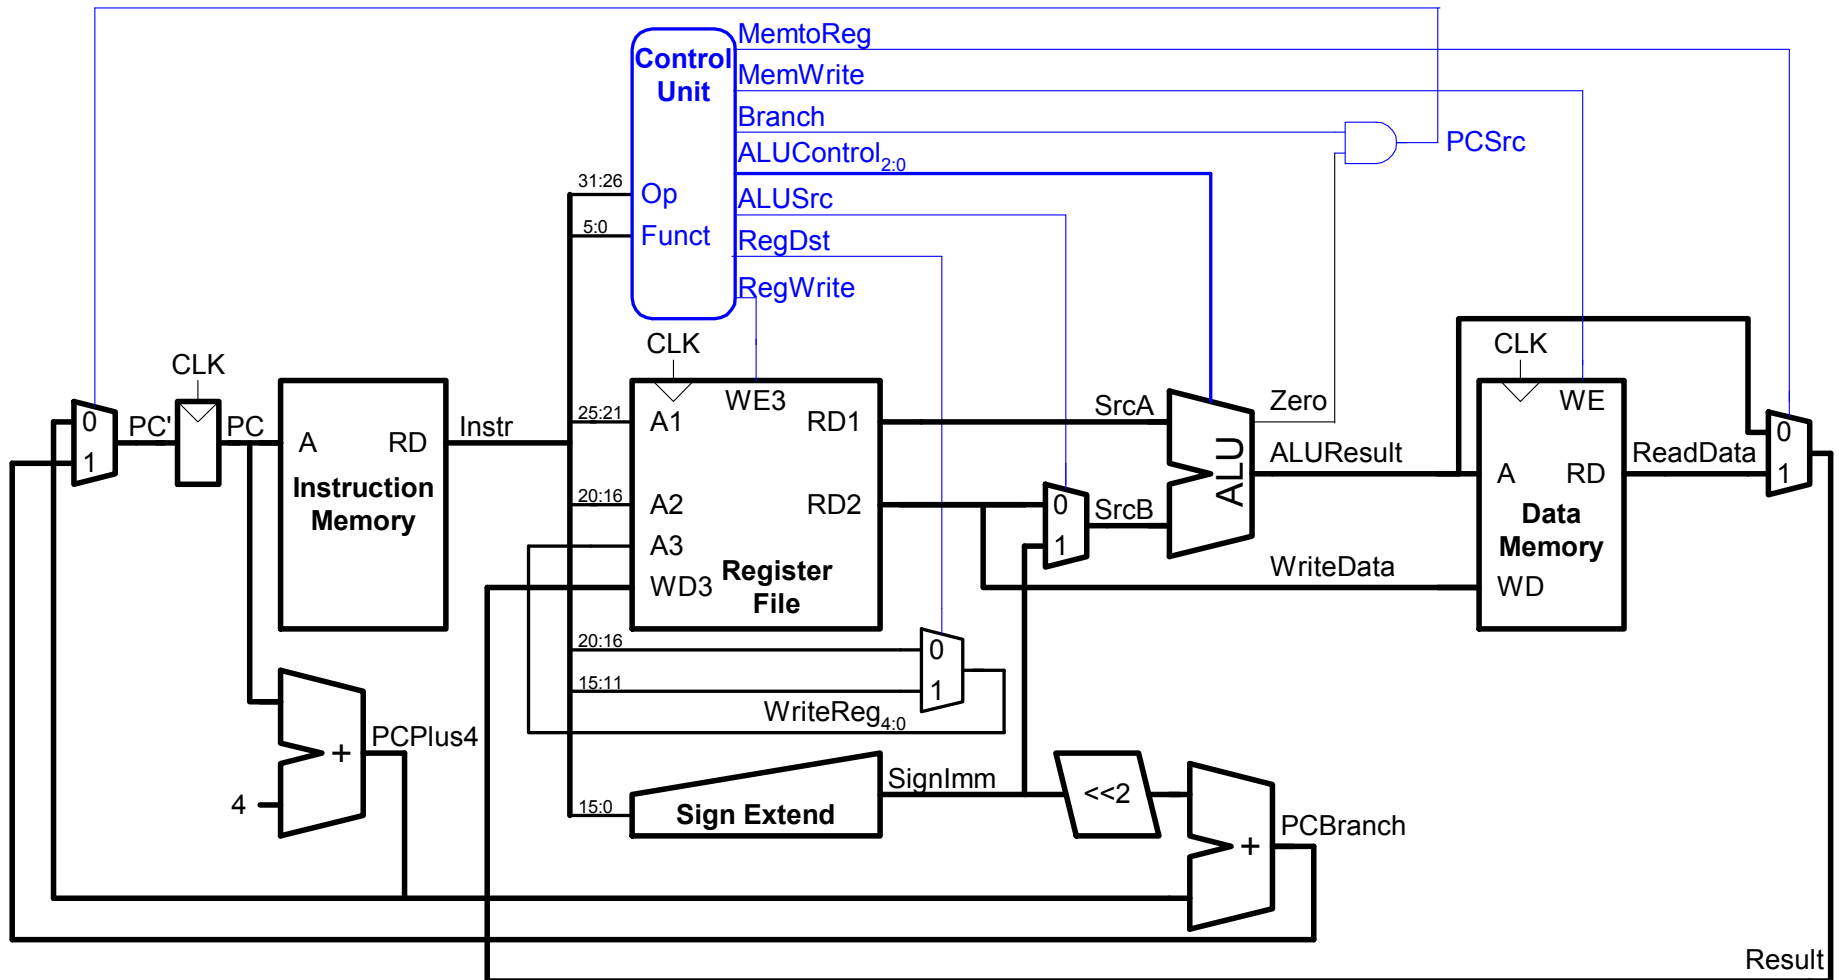
```
wire [31:0] insn;  
wire [5:0] func = insn[5:0]  
wire [5:0] opcode = insn[31:26];  
wire is_add = ((opcode == 6'h00) & (func == 6'h20));  
wire is_addi = (opcode == 6'h0F);  
wire is_lw = (opcode == 6'h23);  
wire is_sw = (opcode == 6'h2A);  
wire ALUinB = is_addi | is_lw | is_sw;  
wire Rwe = is_add | is_addi | is_lw;  
wire Rwd = is_lw;  
wire Rdst = ~is_add;  
wire DMwe = is_sw;
```





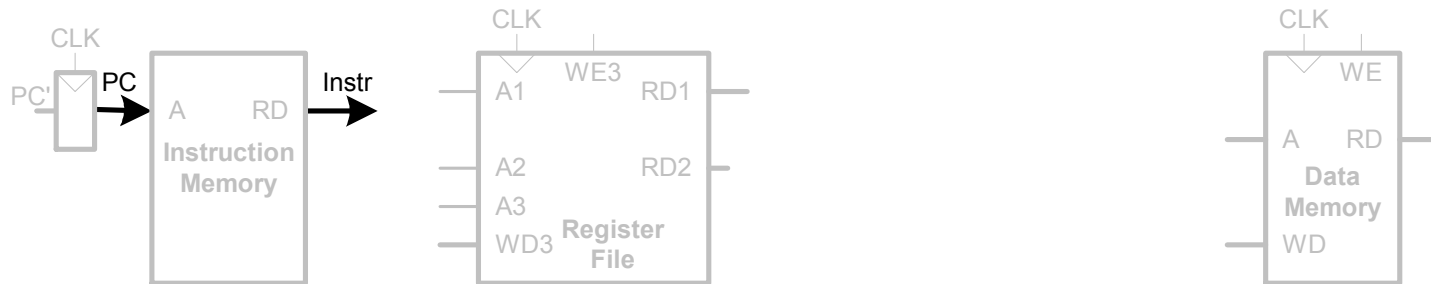
# Single-Cycle Performance

# Single-Cycle Datapath Performance



# Example: Single-Cycle Datapath: **lw** fetch

## ■ **STEP 1:** Fetch instruction



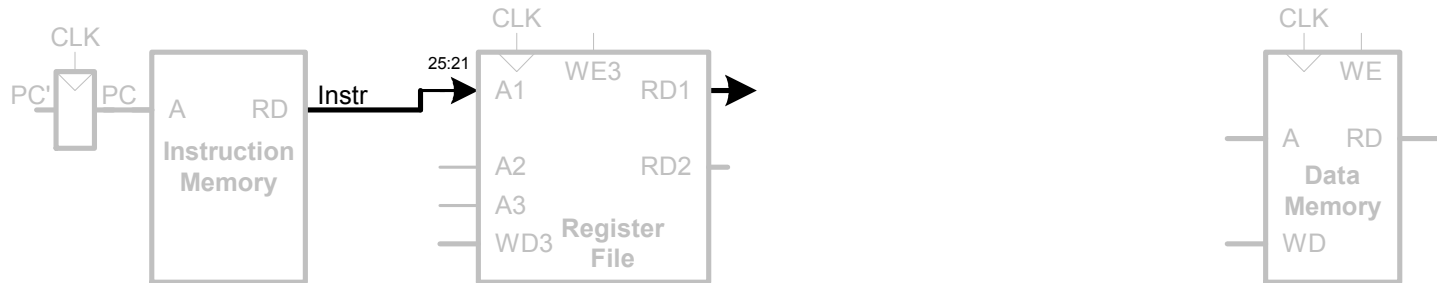
**lw** \$s3, 1(\$0) # read memory word 1 into \$s3

**I-Type**



# Single-Cycle Datapath: `lw` register read

## ■ **STEP 2:** Read source operands from register file



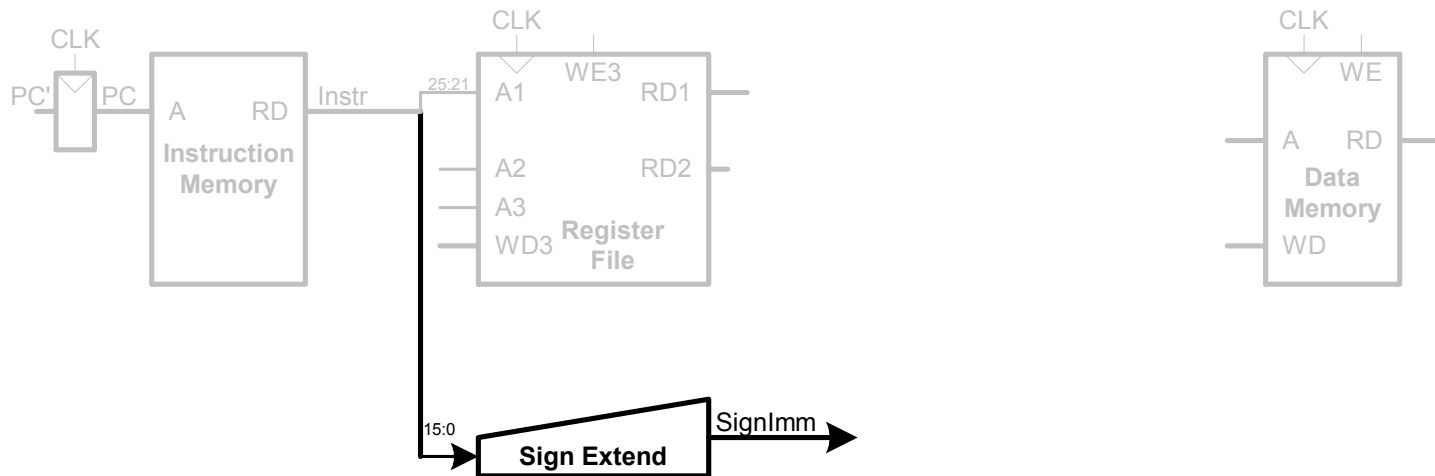
`lw $s3, 1($0) # read memory word 1 into $s3`

**I-Type**

op	rs	rt	imm
6 bits	5 bits	5 bits	16 bits

# Single-Cycle Datapath: `lw` immediate

## ■ **STEP 3:** Sign-extend the immediate



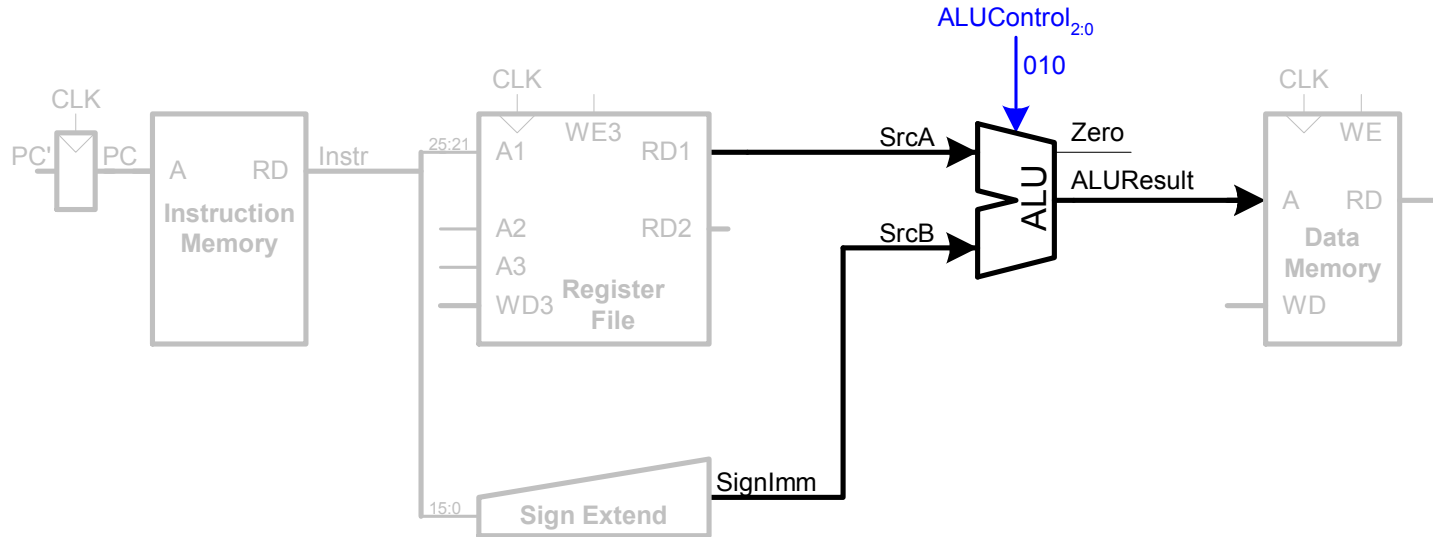
`lw $s3, 1($0) # read memory word 1 into $s3`

**I-Type**

op	rs	rt	imm
6 bits	5 bits	5 bits	16 bits

# Single-Cycle Datapath: `lw` address

## ■ **STEP 4:** Compute the memory address



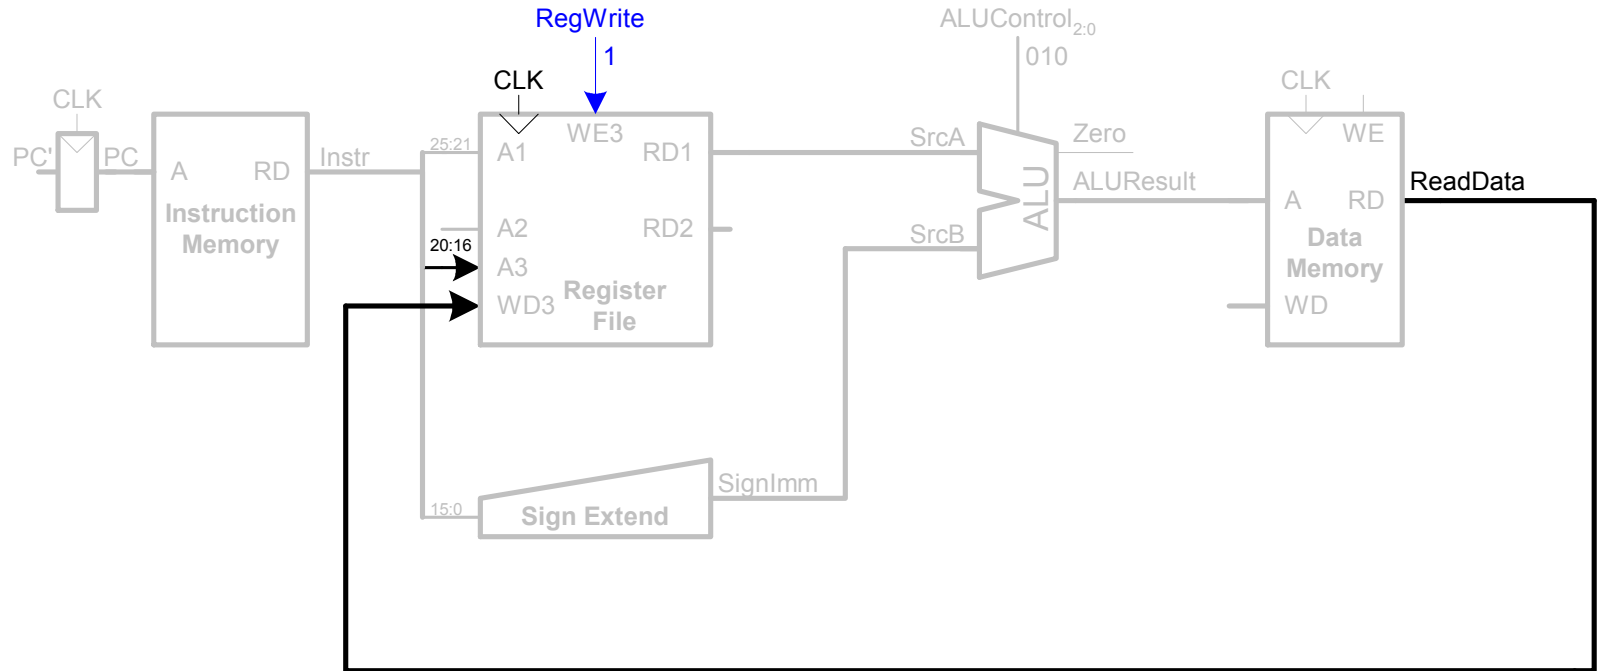
`lw $s3, 1($0) # read memory word 1 into $s3`

**I-Type**

op	rs	rt	imm
6 bits	5 bits	5 bits	16 bits

# Single-Cycle Datapath: **lw** memory read

## ■ **STEP 5:** Read from memory and write back to register file



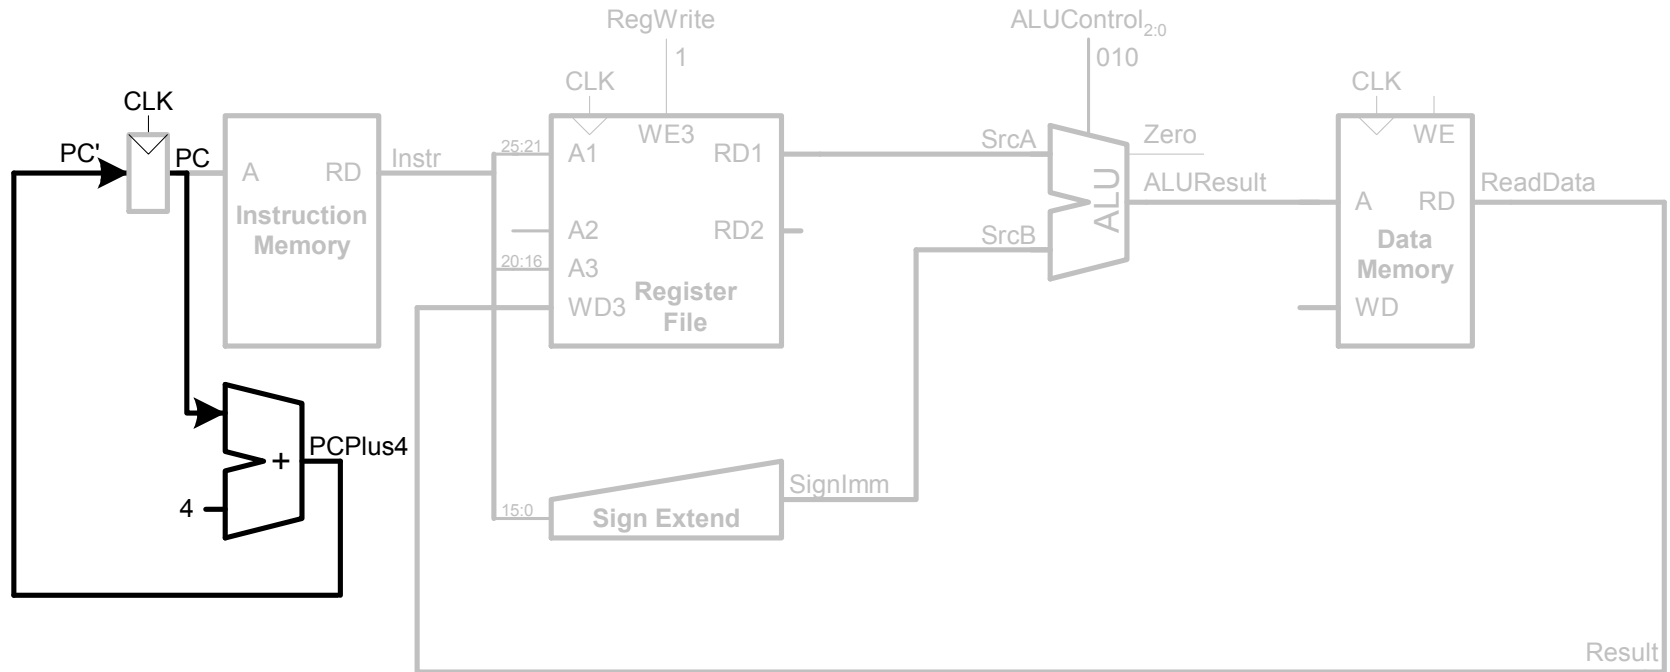
**lw \$s3, 1(\$0) # read memory word 1 into \$s3**

**I-Type**

op	rs	rt	imm
6 bits	5 bits	5 bits	16 bits

# Single-Cycle Datapath: 1w PC increment

## ■ **STEP 6:** Determine address of next instruction



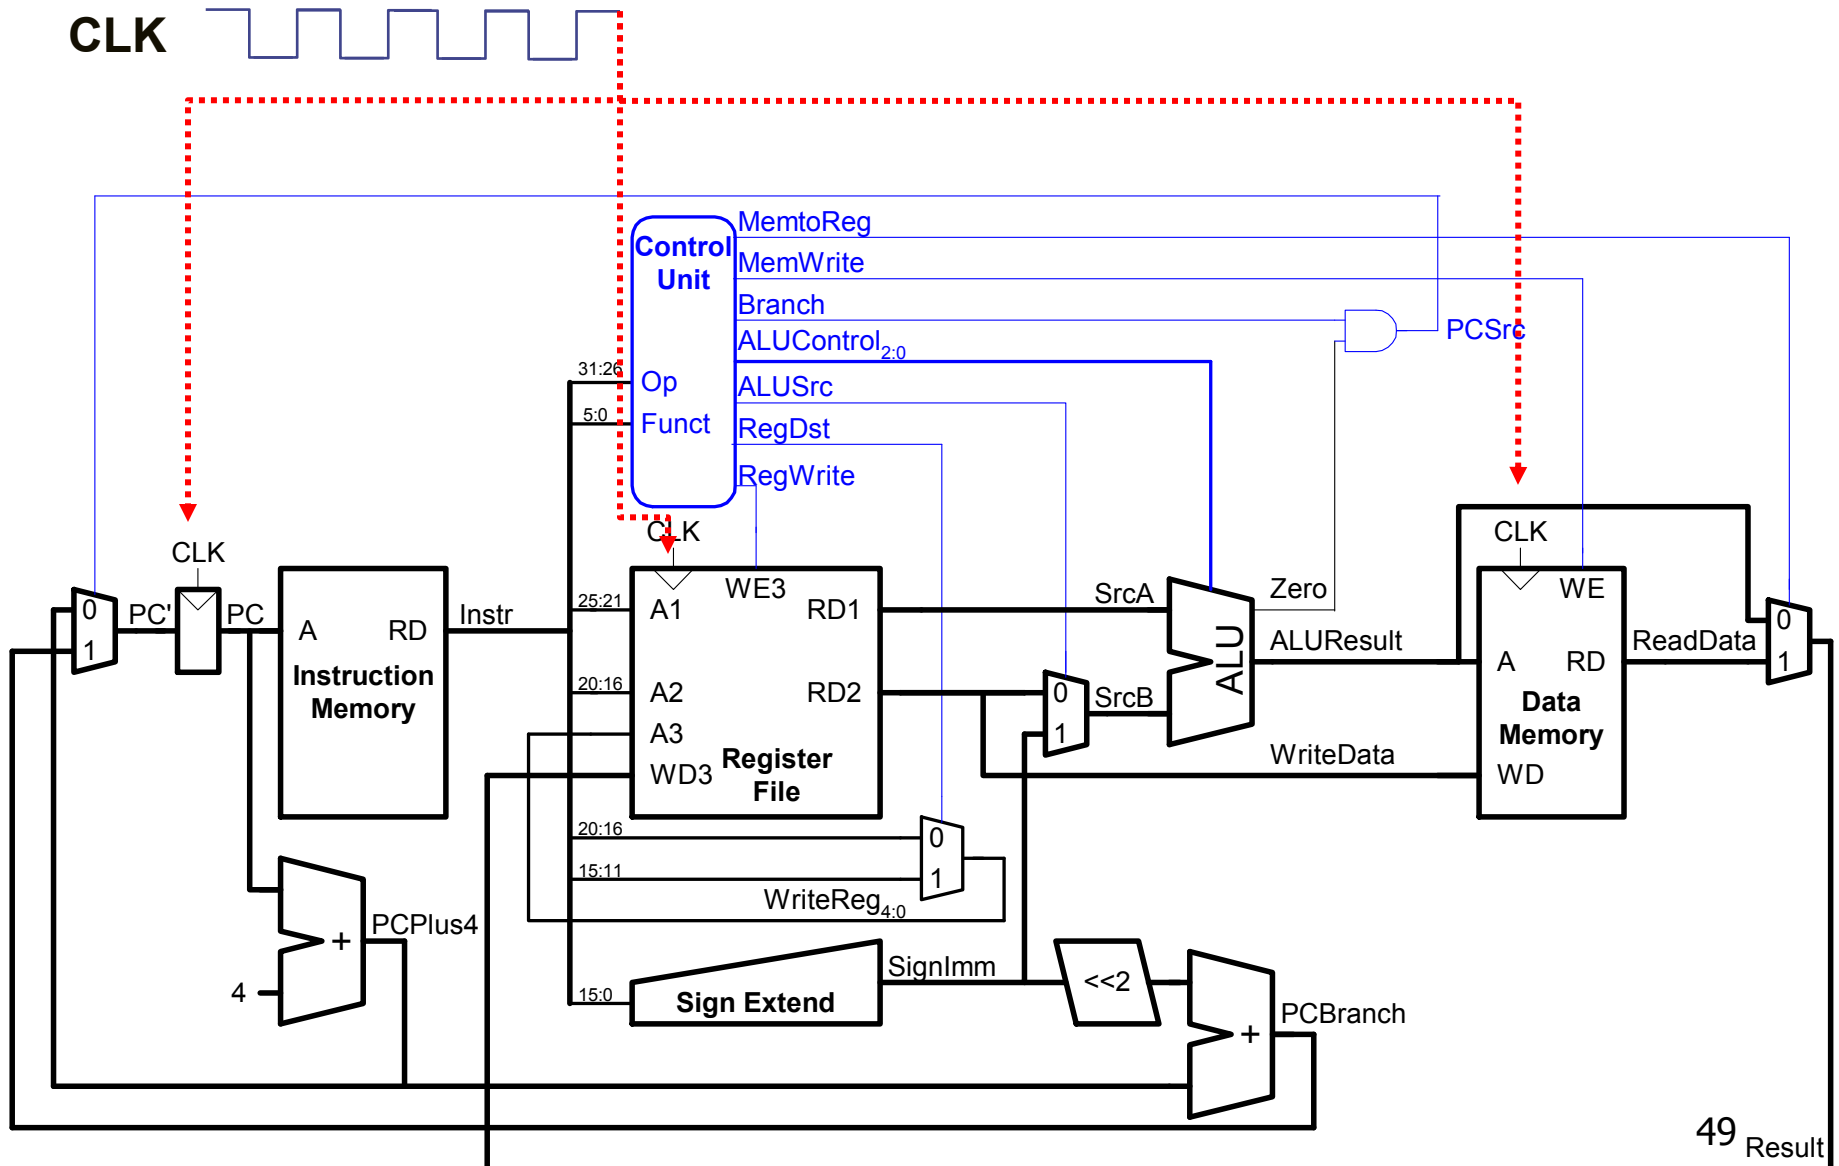
`lw $s3, 1($0) # read memory word 1 into $s3`

**I-Type**

op	rs	rt	imm
6 bits	5 bits	5 bits	16 bits



# Single-Cycle CPU



# Evaluating the Single-Cycle Microarchitecture

# Performance Analysis Basics

---

- Execution time of an instruction
  - $\{\text{CPI}\} \times \{\text{clock cycle time}\}$ 
    - CPI: Number of cycles it takes to execute an instruction
- Execution time of a program
  - Sum over all instructions  $[\{\text{CPI}\} \times \{\text{clock cycle time}\}]$
  - $\{\# \text{ of instructions}\} \times \{\text{Average CPI}\} \times \{\text{clock cycle time}\}$

# A Single-Cycle Microarchitecture: Analysis

---

- Every instruction takes 1 cycle to execute
  - CPI (Cycles per instruction) is strictly 1
- How long each instruction takes is determined by how long the slowest instruction takes to execute
  - Even though many instructions do not need that long to execute
- Clock cycle time of the microarchitecture is determined by how long it takes to complete the slowest instruction
  - Critical path of the design is determined by the processing time of the slowest instruction

# What is the Slowest Instruction to Process?

---

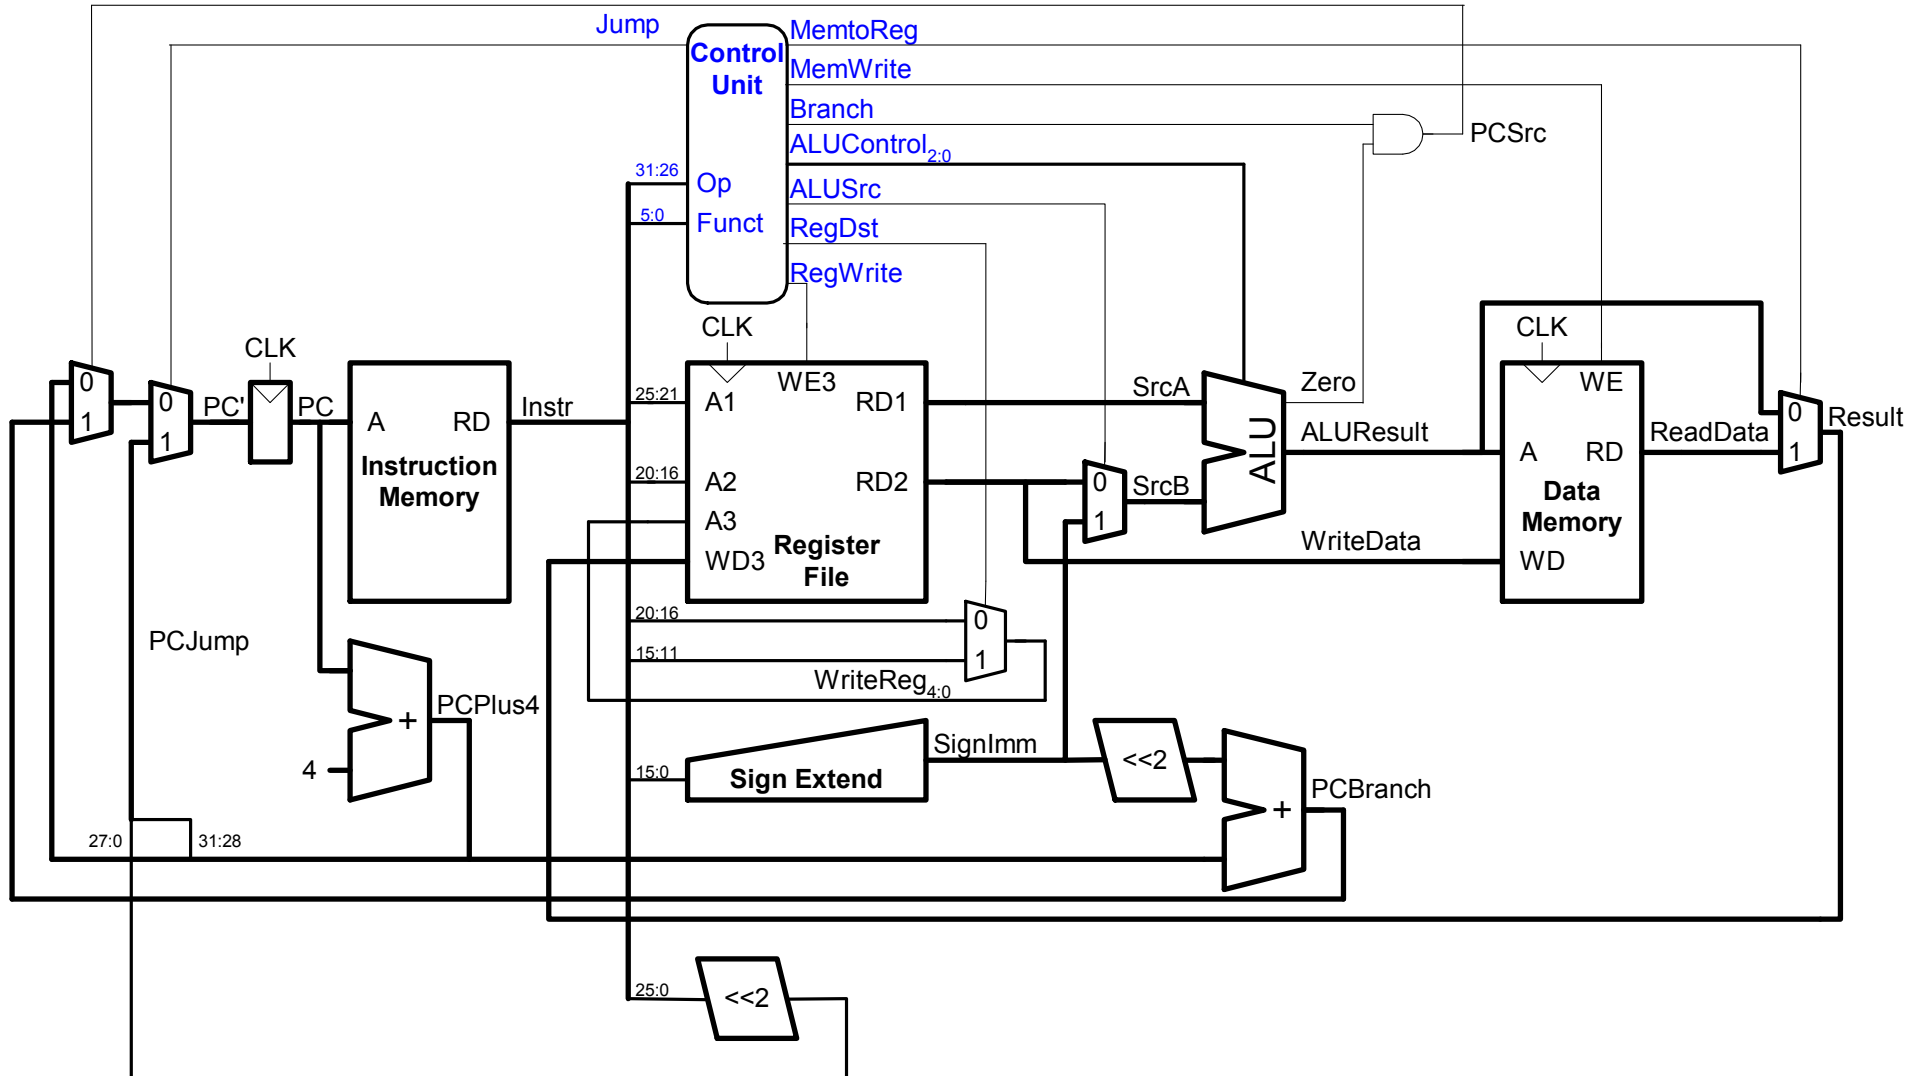
- Let's go back to the basics
  
- All six phases of the instruction processing cycle take a *single machine clock cycle* to complete
  - Fetch
  - Decode
  - Evaluate Address
  - Fetch Operands
  - Execute
  - Store Result
- 1. Instruction fetch (IF)**
  - 2. Instruction decode and register operand fetch (ID/RF)**
  - 3. Execute/Evaluate memory address (EX/AG)**
  - 4. Memory operand fetch (MEM)**
  - 5. Store/writeback result (WB)**
- Do each of the above phases take the same time (latency) for all instructions?

# Example Single-Cycle Datapath Analysis

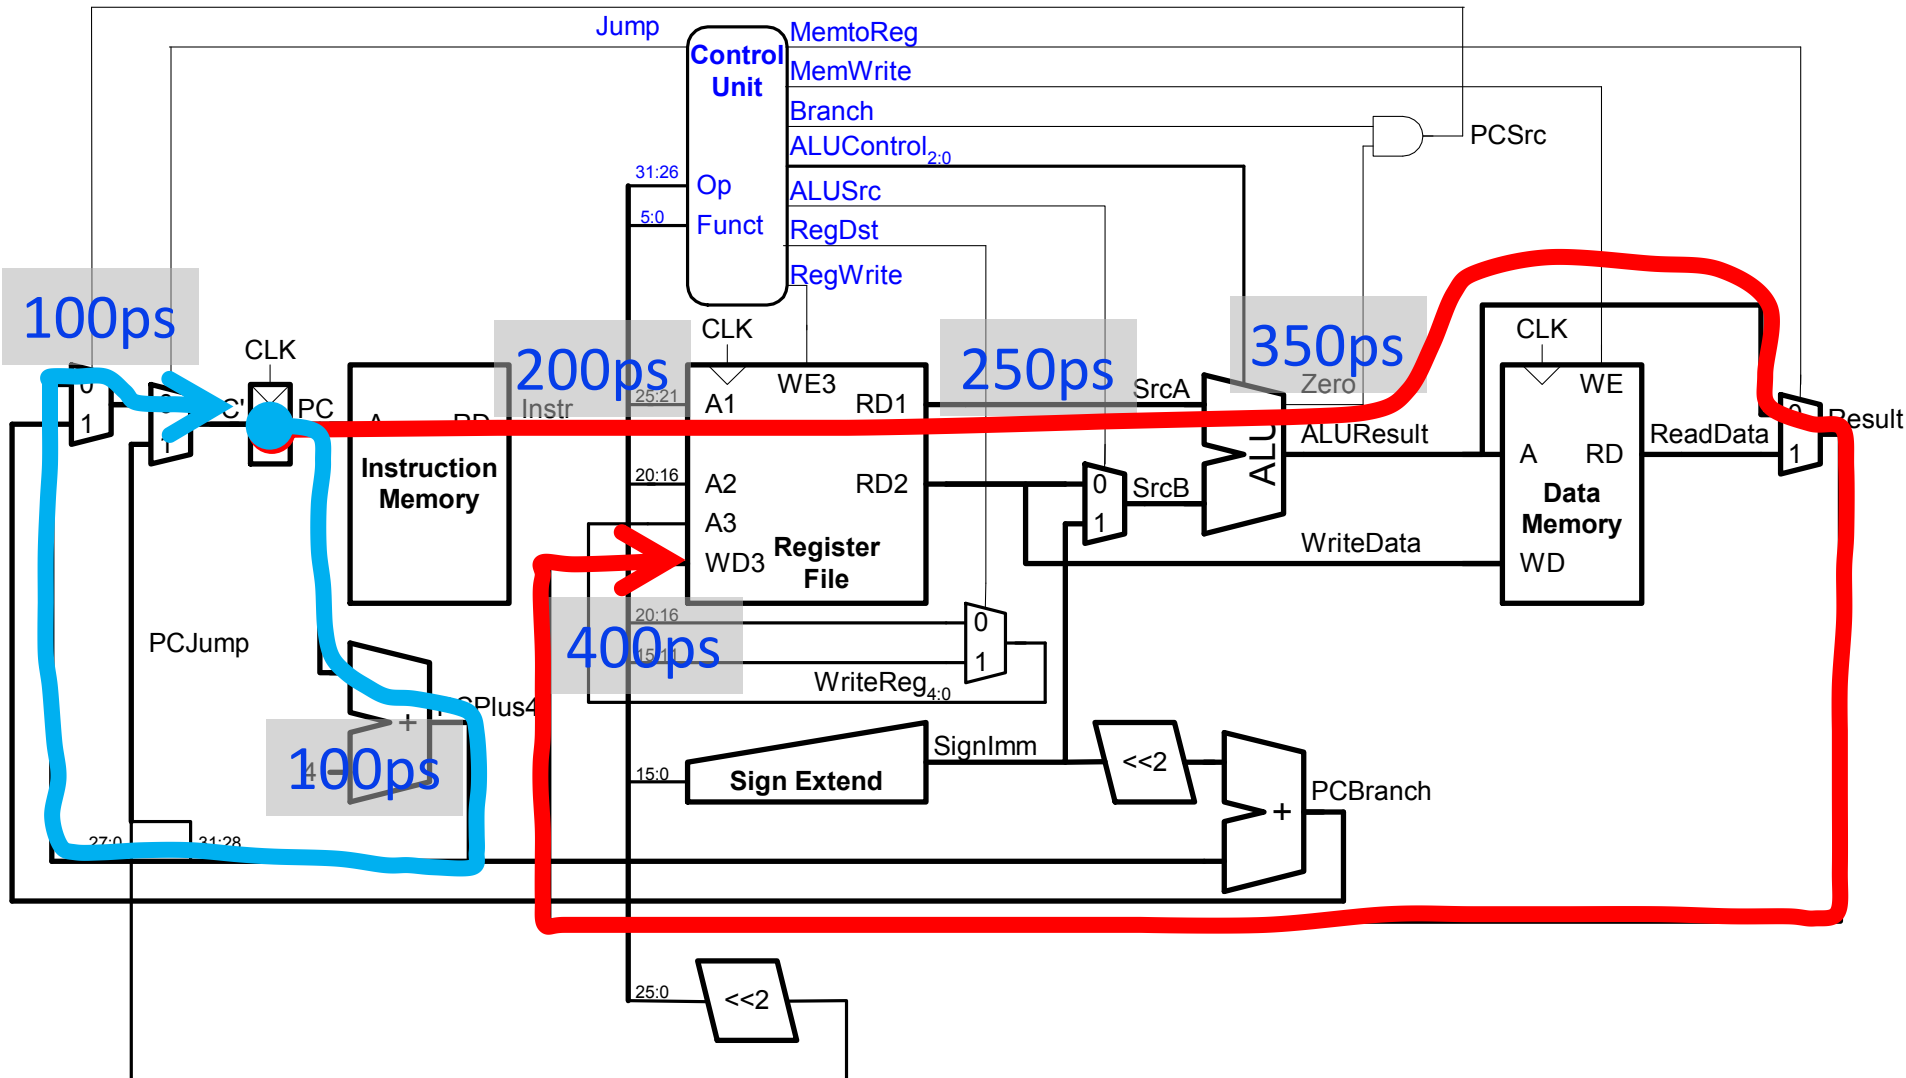
- Assume (for the design in the previous slide)
  - memory units (read or write): 200 ps
  - ALU and adders: 100 ps
  - register file (read or write): 50 ps
  - other combinational logic: 0 ps

steps	IF	ID	EX	MEM	WB	Delay
resources	mem	RF	ALU	mem	RF	
R-type	200	50	100		50	400
I-type	200	50	100		50	400
LW	200	50	100	200	50	600
SW	200	50	100	200		550
Branch	200	50	100			350
Jump	200					250

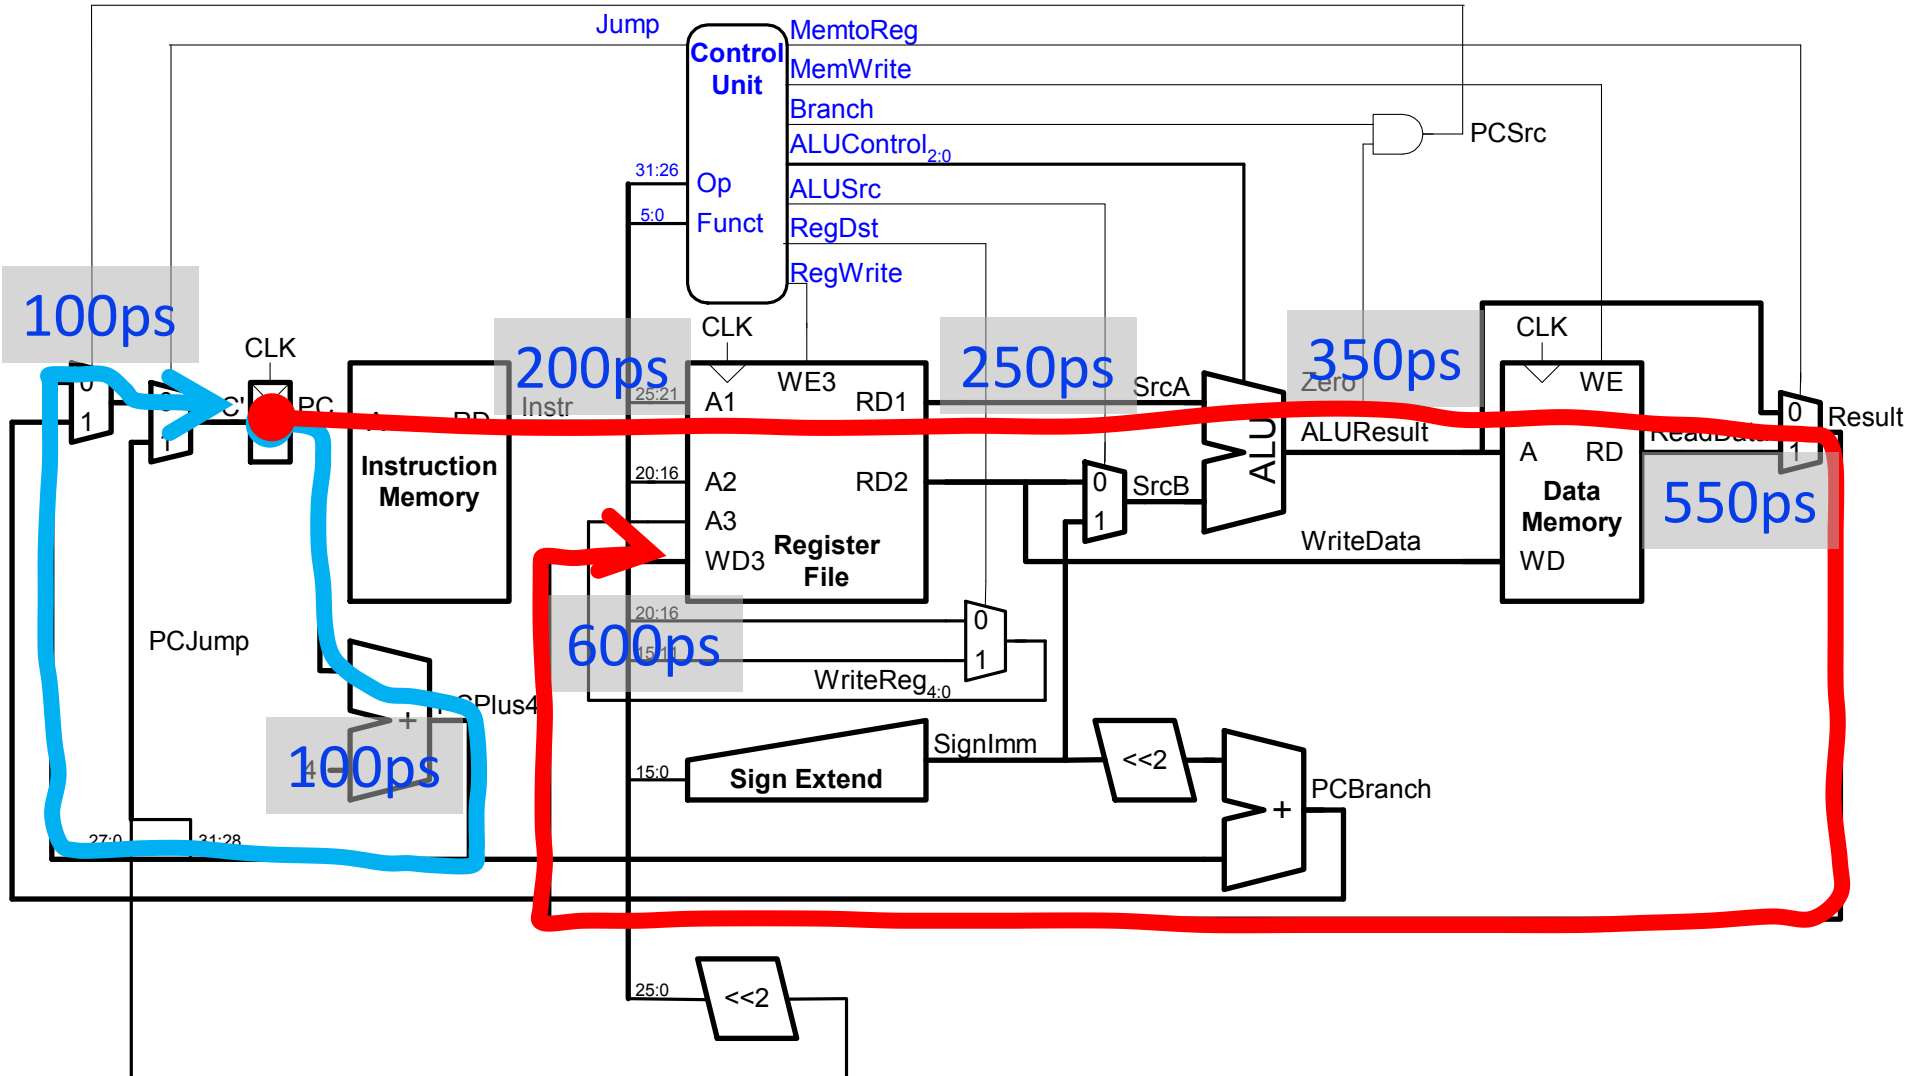
# Let's Find the Critical Path



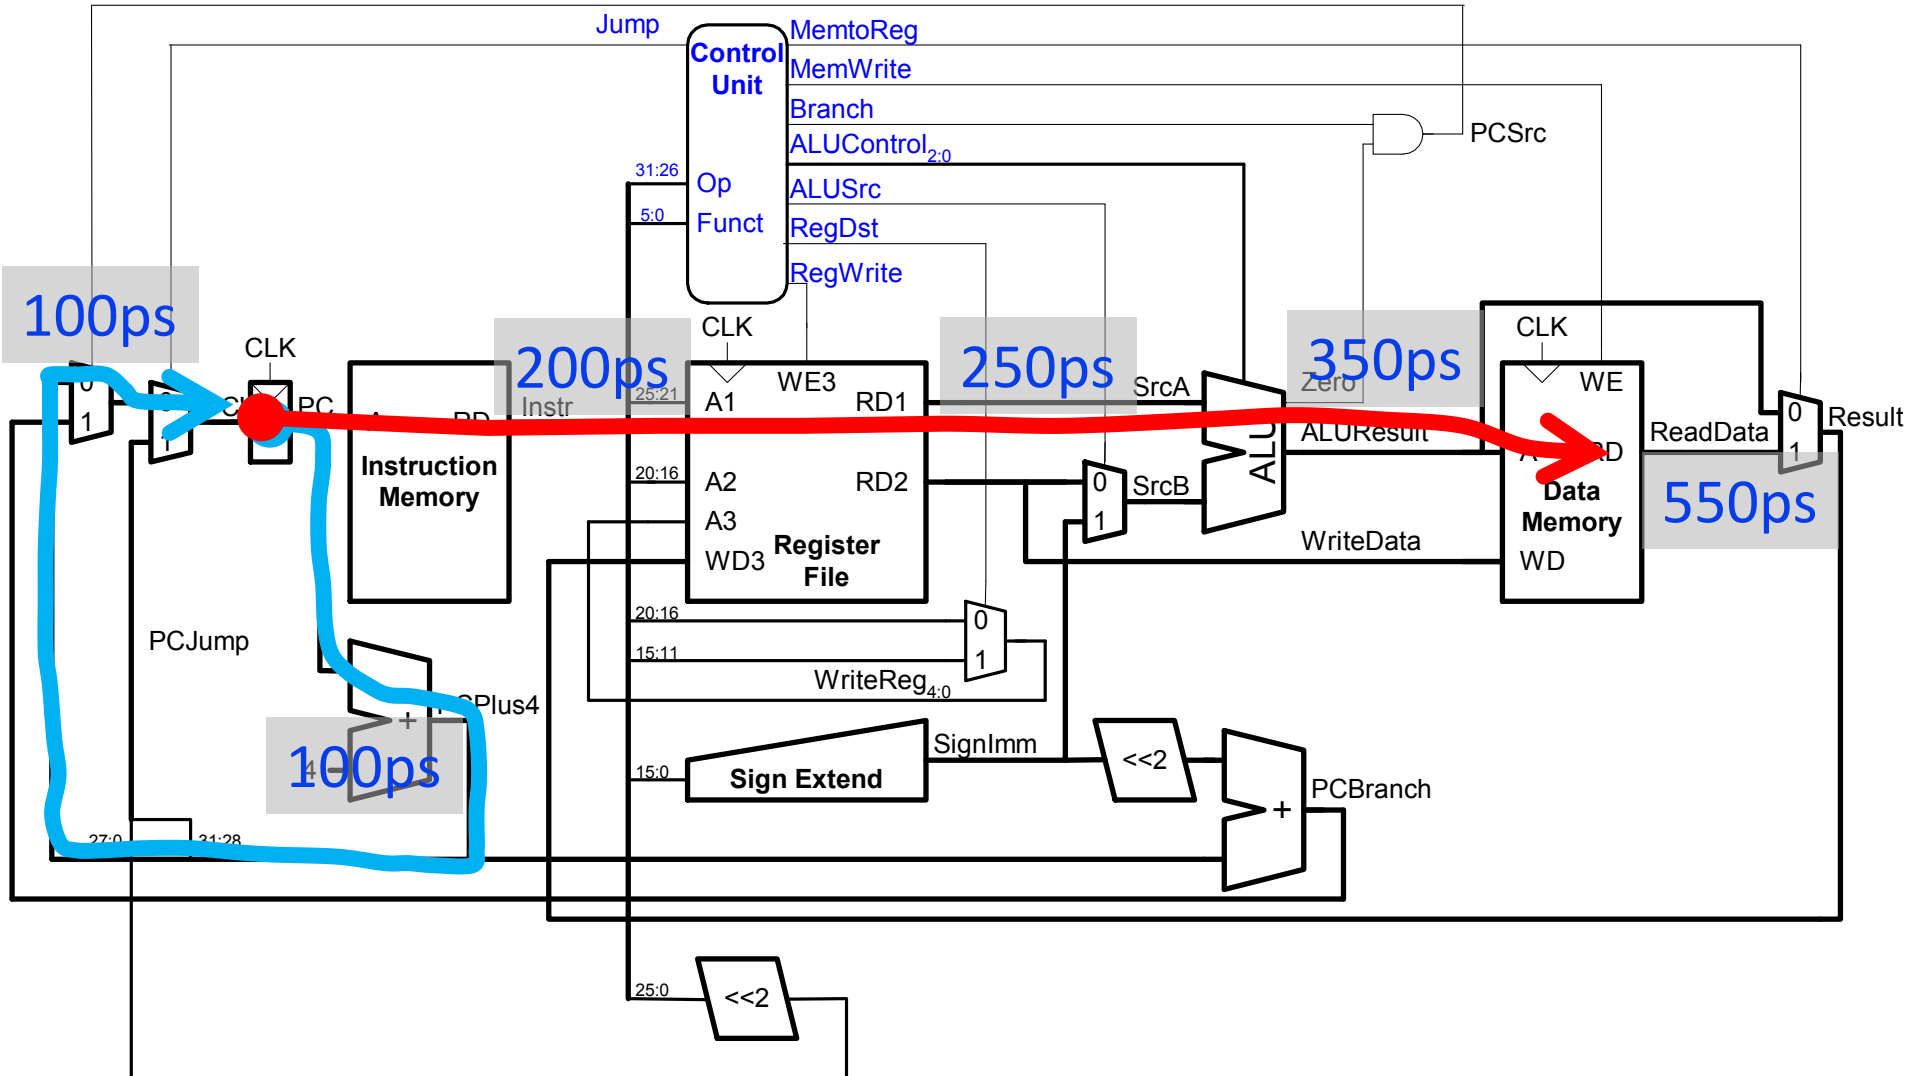
# R-Type and I-Type ALU



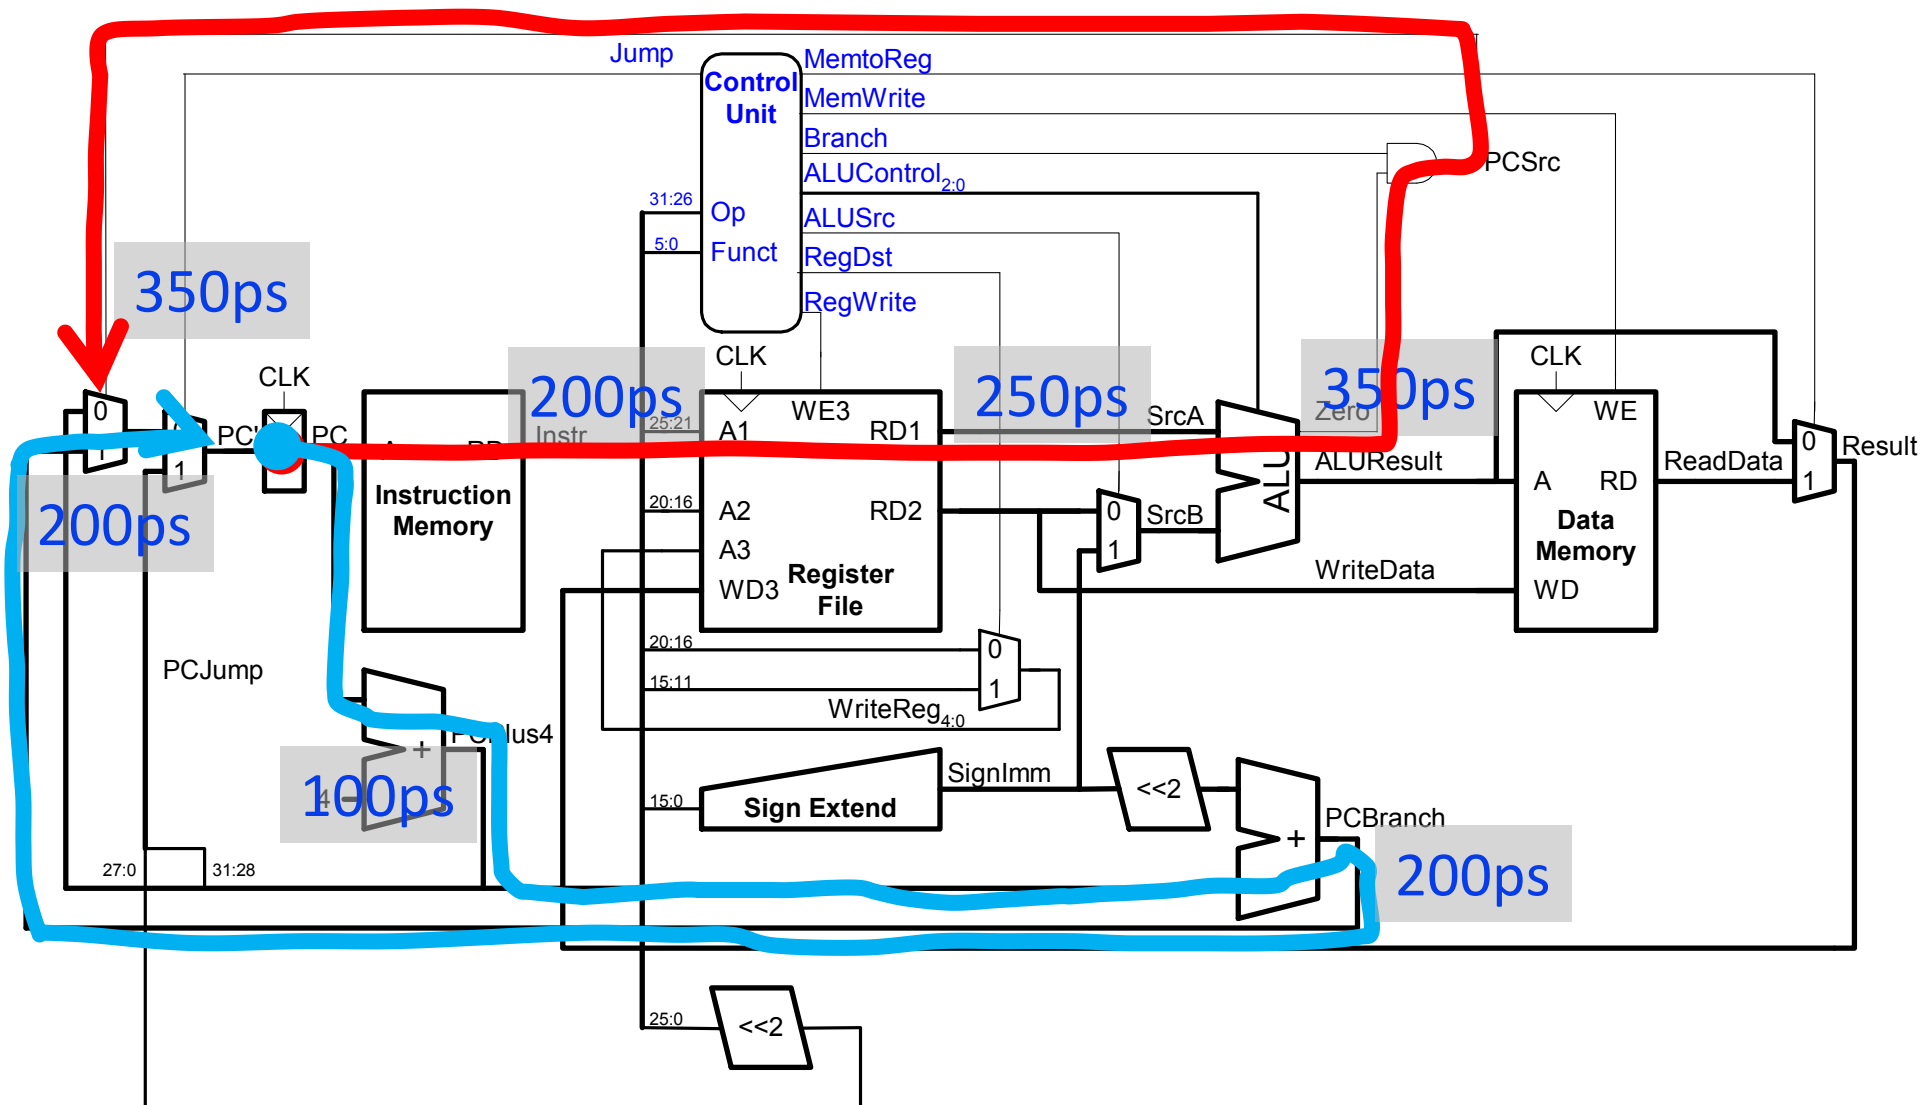




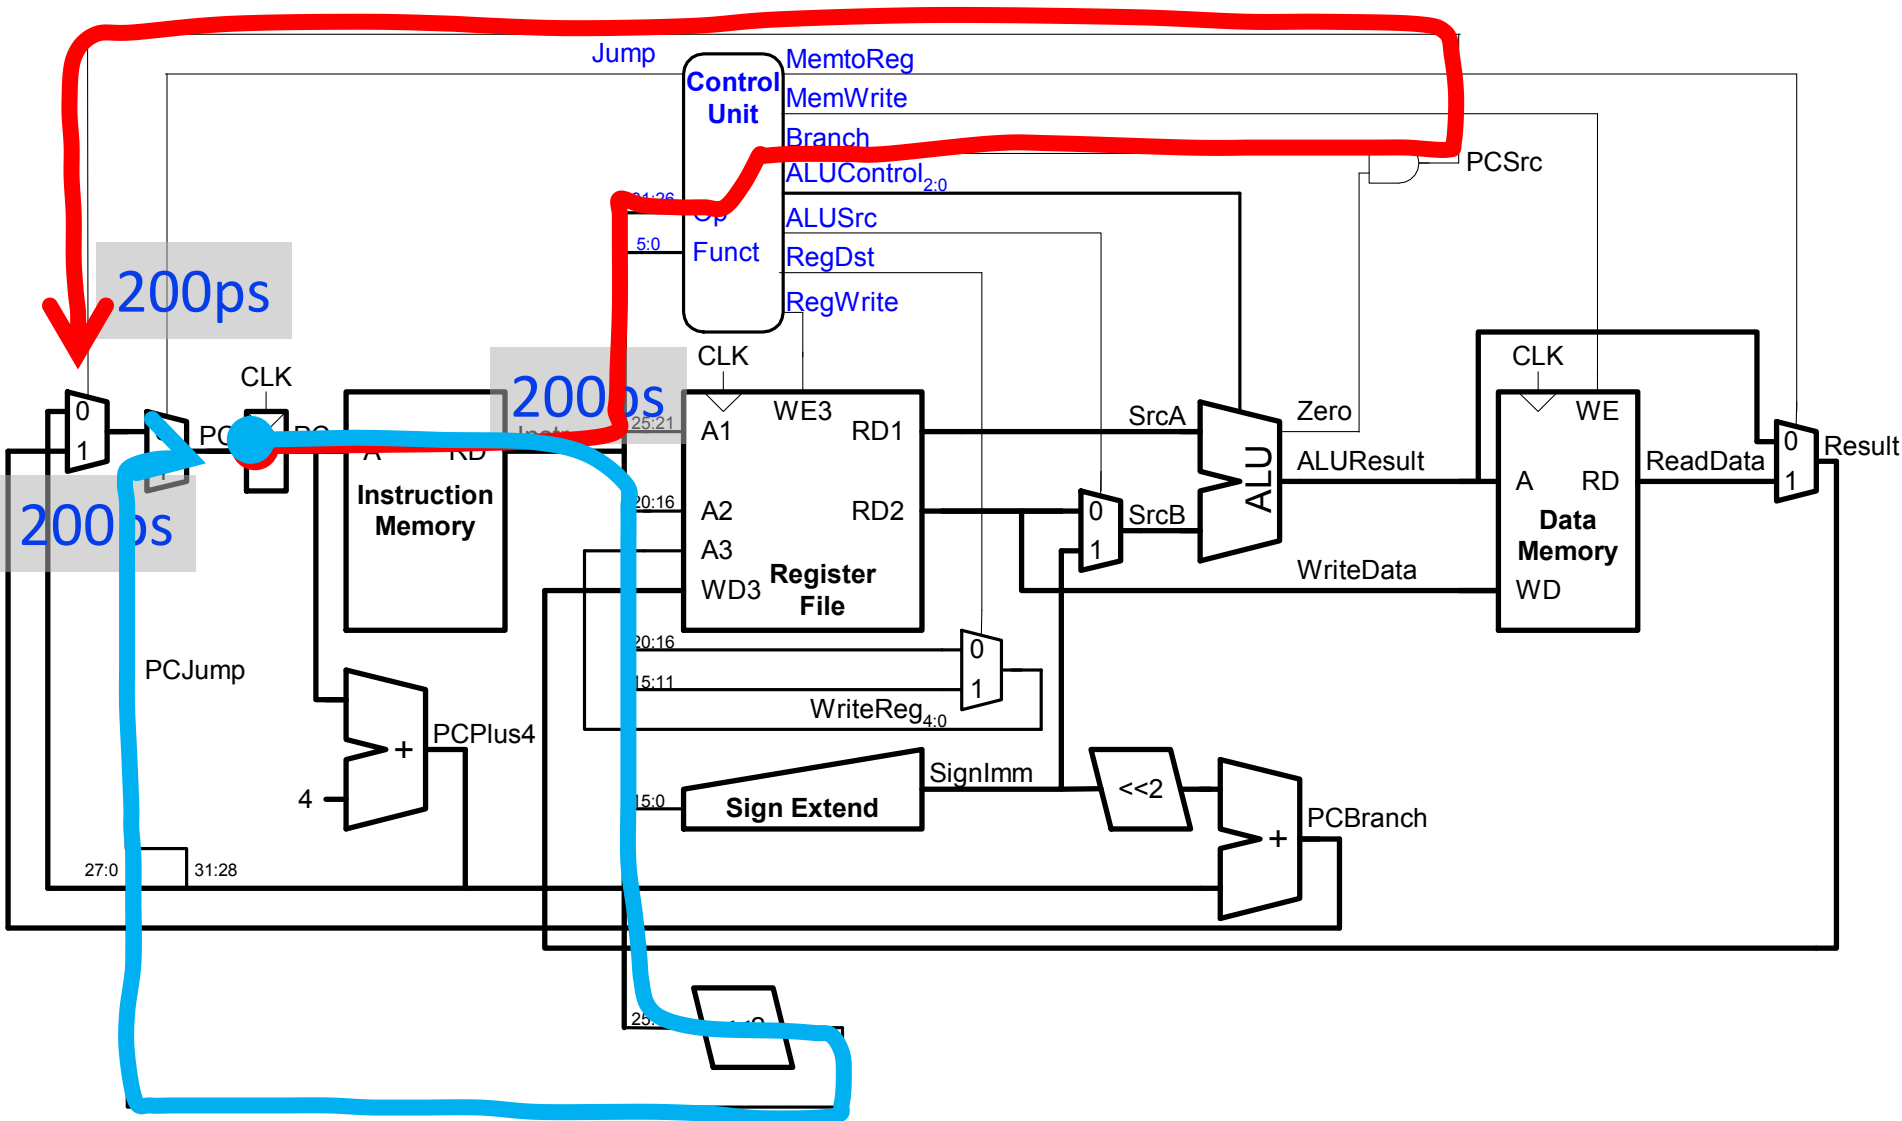
# SW



# Branch Taken



# Jump




# Single Cycle Arc: Complexity

---

- Contrived
  - All instructions run as slow as the slowest instruction
- Inefficient
  - All instructions run as slow as the slowest instruction
  - Must provide worst-case combinational resources in parallel as required by any instruction
  - Need to replicate a resource if it is needed more than once by an instruction during different parts of the instruction processing cycle
- Not easy to optimize/improve performance
  - Optimizing the common case does not work (e.g. common instructions)
  - Need to optimize the worst case all the time

# Summary

---



- Overview of ISAs
- Datapath storage elements
- MIPS Datapath
- MIPS Control