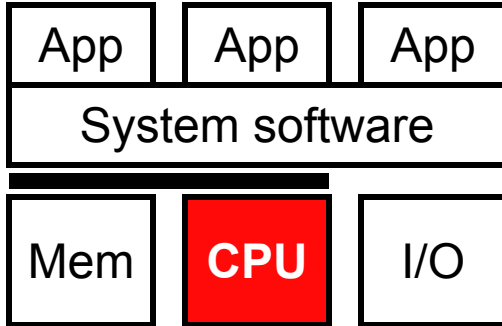


# Computer Architecture

## Lec 6: Branch Prediction

# This Lecture

---



- Control Dependences
- Branch Prediction

# Control Dependences

# Control Dependence

---

- Question: What should the fetch PC be in the next cycle?
- Answer: The address of the next instruction
  - ❑ All instructions are control dependent on previous ones. Why?
- If the fetched instruction is a non-control-flow instruction:
  - ❑ Next Fetch PC is the address of the next-sequential instruction
  - ❑ Easy to determine if we know the size of the fetched instruction
- If the instruction that is fetched is a control-flow instruction:
  - ❑ How do we determine the next Fetch PC?
- In fact, how do we even know whether or not the fetched instruction is a control-flow instruction?

# Branch Types

---

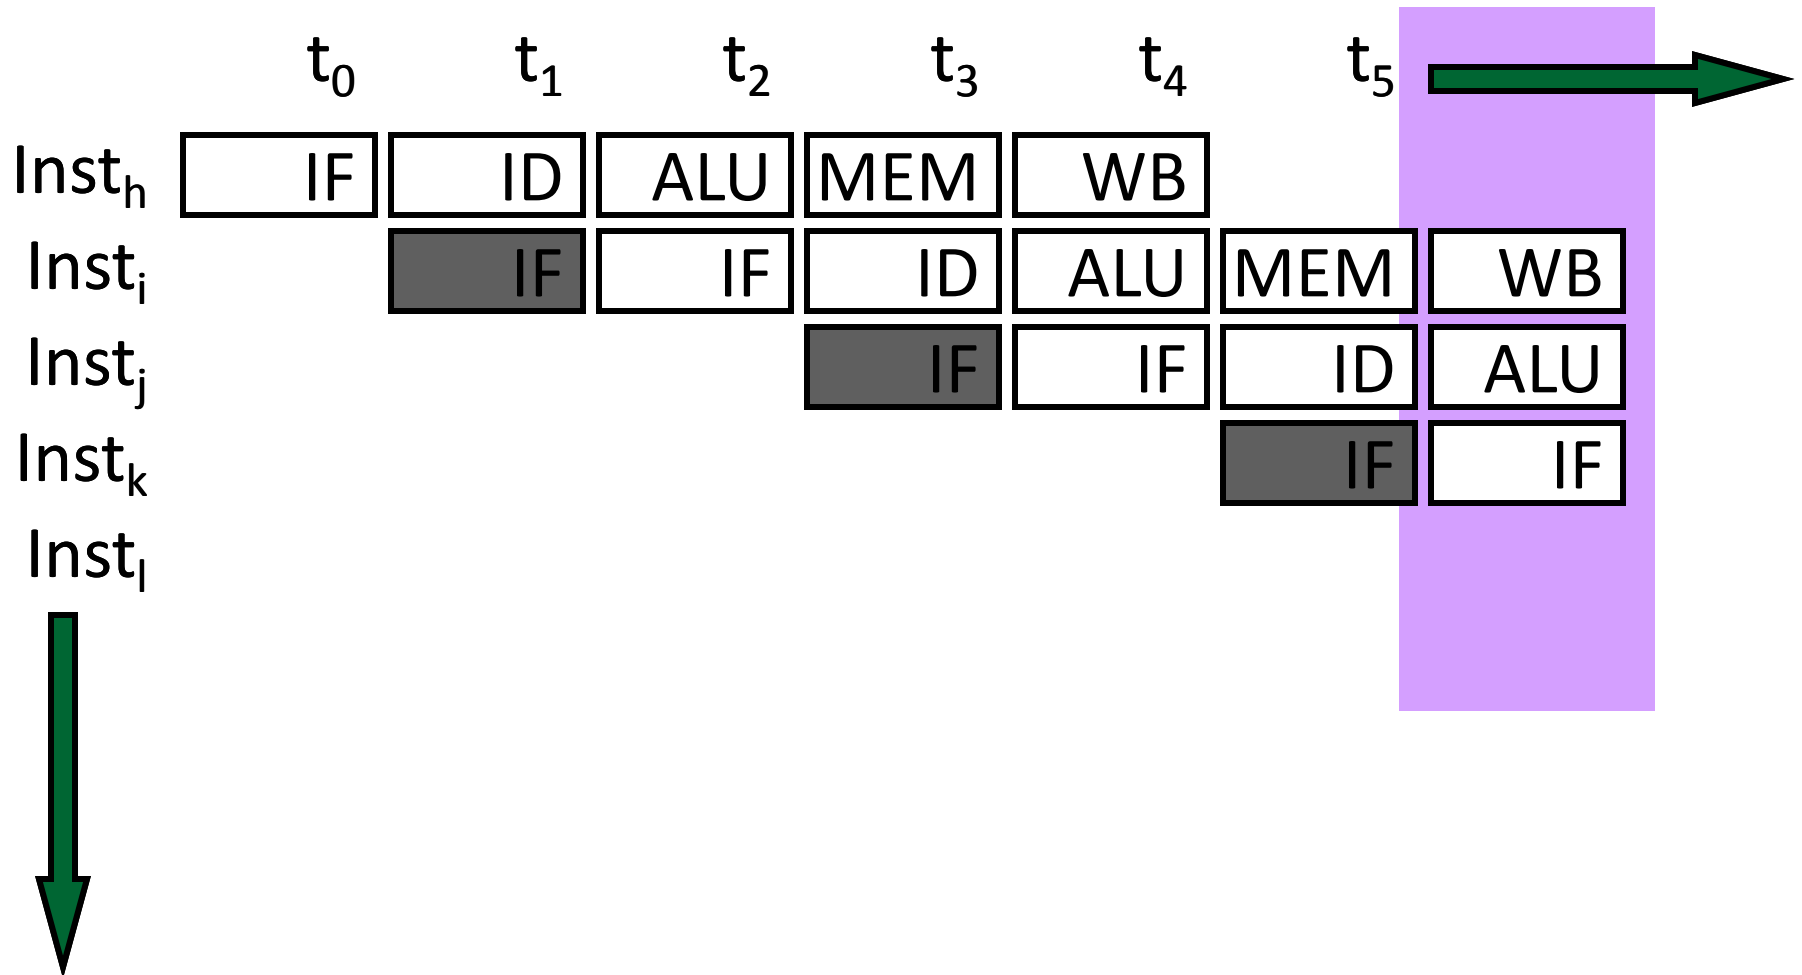
Type	Direction at fetch time	Number of possible next fetch addresses?	When is next fetch address resolved?
Conditional (beq)	Unknown	2	Execution (register dependent)
Unconditional (j)	Always taken	1	Decode (PC + offset)
Call (jal X)	Always taken	1	Decode (PC + offset)
Return (jr \$ra)	Always taken	Many	Execution (register dependent)
Indirect (jr \$r1)	Always taken	Many	Execution (register dependent)

# How to Handle Control Dependences

---

- Critical to keep the pipeline full with correct sequence of dynamic instructions.
- Potential solutions if the instruction is a control-flow instruction:
  - **Stall** the pipeline until we know the next fetch address
  - Guess the next fetch address (**branch prediction**)
  - Employ delayed branching (**branch delay slot**)
  - Do something else (**fine-grained multithreading**)
  - Eliminate control-flow instructions (**predicated execution**)
  - Fetch from both possible paths (if you know the addresses of both possible paths) (**multipath execution**)

# Stall Fetch Until Next PC is Known: Good Idea?



This is the case with non-control-flow and unconditional br instructions!

# The Branch Problem

---

- Control flow instructions (branches) are frequent
  - 15-25% of all instructions
- Problem: Next fetch address after a control-flow instruction is not determined after  $N$  cycles in a pipelined processor
  - $N$  cycles: (minimum) branch resolution latency
- If we are fetching  $W$  instructions per cycle (i.e., if the pipeline is  $W$  wide)
  - A branch misprediction leads to  $N \times W$  wasted instruction slots



# Importance of The Branch Problem

---

- Assume  $N = 20$  (20 pipe stages),  $W = 5$  (5 wide fetch)
  - Assume: 1 out of 5 instructions is a branch
- How long does it take to fetch 500 instructions?
  - ❑ 100% accuracy
    - 100 cycles (all instructions fetched on the correct path)
    - No wasted work;  $IPC = 500/100$
  - ❑ 99% accuracy
    - 100 (correct path) +  $20 * 1$  (wrong path) = 120 cycles
    - 20% extra instructions fetched;  $IPC = 500/120$
  - ❑ 90% accuracy
    - 100 (correct path) +  $20 * 10$  (wrong path) = 300 cycles
    - 200% extra instructions fetched;  $IPC = 500/300$
  - ❑ 60% accuracy
    - 100 (correct path) +  $20 * 40$  (wrong path) = 900 cycles
    - 800% extra instructions fetched;  $IPC = 500/900$

# Average Run-Length between Branches

---

Average dynamic instruction mix of SPEC CPU 2017 [Limaye and Adegbiya, ISPASS'18]:

	SPECint	SPECfp
Branches	19 %	11 %
Loads	24 %	26 %
Stores	10 %	7 %
Other	47 %	56 %

SPECint17: *perlbench, gcc, mcf, omnetpp, xalancbmk, x264, deepsjeng, leela, exchange2, xz*

SPECfp17: *bwaves, cactus, lbm, wrf, pop2, imagick, nab, fotonik3d, roms*

***What is the average run length between branches?***

**Roughly 5-10 instructions**

# Branch Prediction

# How to Handle Control Dependences

---

- Critical to keep the pipeline full with correct sequence of dynamic instructions.
- Potential solutions if the instruction is a control-flow instruction:
  - **Stall** the pipeline until we know the next fetch address
  - Guess the next fetch address (**branch prediction**)
  - Employ delayed branching (**branch delay slot**)
  - Do something else (**fine-grained multithreading**)
  - Eliminate control-flow instructions (**predicated execution**)
  - Fetch from both possible paths (if you know the addresses of both possible paths) (**multipath execution**)

# Branch Speculation

Without Prediction:

	1	2	3	4	5	6	7	8	9
addi r3←r1,1	F	D	X	M	W				
bnez r3,targ		F	D	X	M	W			
stall			--	--	--	--	--		
stall				--	--	--	--	--	
st r6←[r7+4]					F	D	X	M	W
mul r10←r8,r9									
targ:add r4←r4,r5									

Correct Prediction:

	1	2	3	4	5	6	7	8	9
addi r3←r1,1	F	D	X	M	W				
bnez r3,targ		F	D	X	M	W			
st r6←[r7+4]			F	D	X	M	W		
mul r10←r8,r9				F	D	X	M	W	

speculative

**Correct Prediction Saves 2 Cycles!**

# Mis-prediction Penalty

	1	2	3	4	5	6	7	8	9
Recovery:	F	D	X	M	W				
addi r3←r1,1									
bnez r3,targ		F	D	X	M	W			
<del>st r6 [r7+4]</del>			F	D	--	--	--		
<del>mul r10←r8,r9</del>				F	--	--	--	--	
targ:add r4←r4,r5					F	D	X	M	W

- **Mis-speculation recovery**: what to do on wrong guess
  - Not too painful in a short, in-order pipeline
  - Branch resolves in X
  - + Younger insns (in F, D) haven't changed permanent state
  - **Flush** insns currently in D and X (i.e., replace with **nops**)

# Performance Analysis

---

- correct guess  $\Rightarrow$  no penalty ~86% of the time
- incorrect guess  $\Rightarrow$  2 bubbles
- Assume
  - no data dependency related stalls
  - 20% control flow instructions
  - 70% of control flow instructions are taken
  - $\text{CPI} = [ 1 + (0.20 * 0.7) * 2 ] =$   
 $= [ 1 + 0.14 * 2 ] = 1.28$


probability of  
a wrong guess

penalty for  
a wrong guess

Can we reduce either of the two penalty terms?

# Branch Prediction Components

---




- Step #1: is it a branch?
  - Easy after decode...
- Step #2: is the branch taken or not taken?
  - Direction predictor (applies to conditional branches only)
  - Predicts taken/not-taken
- Step #3: if the branch is taken, where does it go?
  - Easy after decode...



# When to Perform Branch Prediction?


- Option #1: During Decode
  - Look at instruction opcode to determine branch instructions
  - Can calculate next PC from instruction (for PC-relative branches)
    - One cycle “mis-fetch” penalty **even if branch predictor is correct**



- Option #2: During Fetch?
  - How do we do that?

Is It a Branch?


# Revisiting Branch Prediction Components



- Step #1: is it a branch?
  - Easy after decode... during fetch: **predictor**
- Step #2: is the branch taken or not taken?
  - Direction predictor (later)
- Step #3: if the branch is taken, where does it go?
  - Branch target predictor (BTB)
  - Supplies target PC if branch is taken

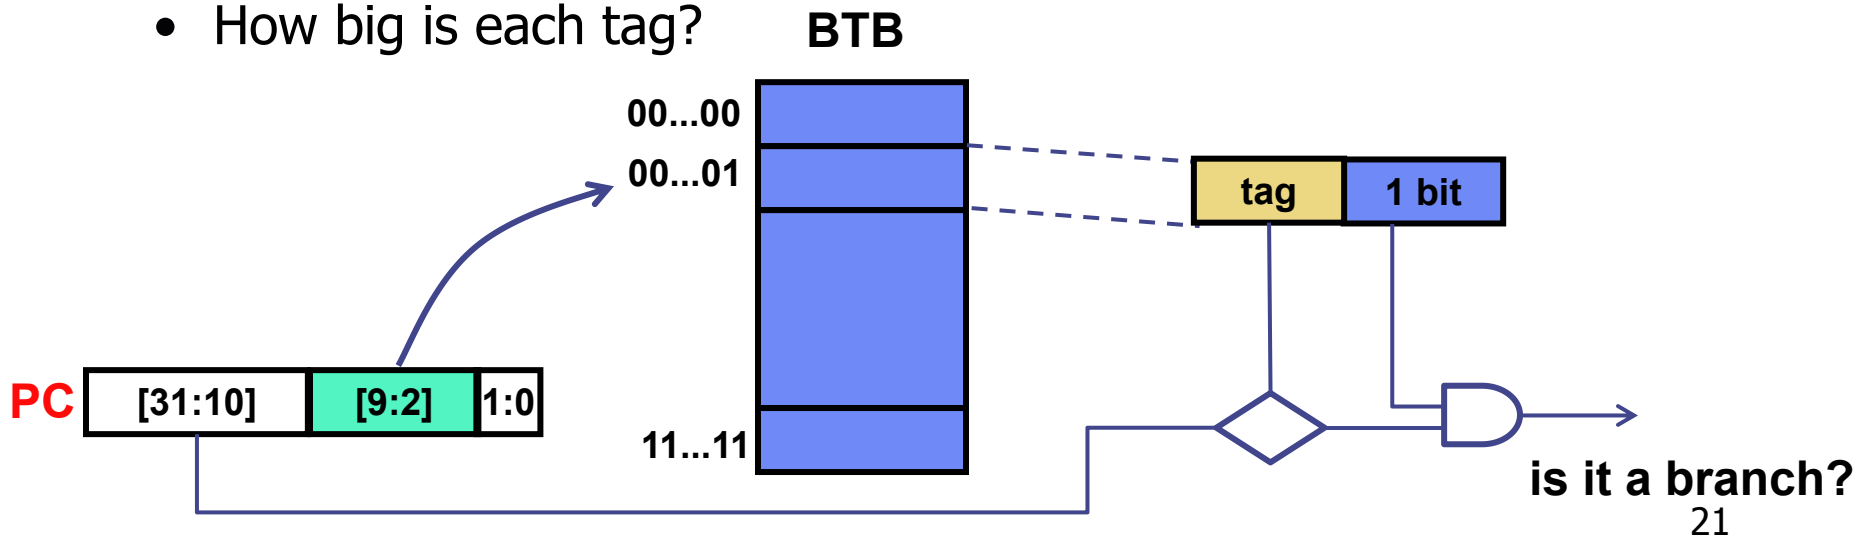
# Branch Target Buffer

- **Branch target buffer (BTB):**
  - Record a list of branches we have seen
    - + code doesn't change
  - PC indexes table of bits
    - each entry is 1 bit: is there a branch here?
  - What about aliasing?
    - Two PCs with the same lower bits?




# Branch Target Buffer

- BTB entries are too coarse-grained
- + Record only branches that were taken at least once
  - a never-taken branch might as well be a NOP
    - doesn't help enough
- better idea: **Tag each BTB entry**
  - remember **some** things precisely, rather than everything imprecisely
  - record a subset of actual taken branches
  - $\text{is\_a\_branch} = (\text{BTB}[\text{PC}].\text{branch} \ \&\& \ \text{BTB}[\text{PC}].\text{tag} == \text{PC})$
  - How big is each tag?



# Branch Target Prediction


# Revisiting Branch Prediction Components



- Step #1: is it a branch?
  - Easy after decode... during fetch: predictor
- Step #2: is the branch taken or not taken?
  - Direction predictor
- Step #3: if the branch is taken, where does it go?
  - **Branch target predictor (BTB)**
  - Supplies target PC if branch is taken

# Branch Target Buffer, Again


- **Branch target buffer (BTB):**
  - “guess” the future PC based on past behavior
  - “Last time the branch X was taken, it went to address Y”
    - “So, in the future, if address X is fetched, fetch address Y next”
  - Essentially: branch will go to same place it went last time
  - PC indexes table of **target addresses**
    - use tags to precisely remember a subset of branch targets
  - What about aliasing?
    - Two PCs with the same lower bits? (just a prediction)






# Branch Target Buffer (continued)

- At Fetch, how do we know we have a branch? We don't...
  - ...**all insns access BTB in parallel with Imem Fetch**
- **BTB predicts which insn are branches, and targets**
  - **tag each entry with its corresponding PC**
  - Update BTB on every taken branch insn, record target PC:
    - $\text{BTB}[\text{PC}].\text{tag} = \text{PC}$ ,  $\text{BTB}[\text{PC}].\text{target} = \text{target of branch}$
  - All insns access at Fetch *in parallel* with Imem
    - Check for tag match, indicates insn at that PC is a branch
      - otherwise, assume insn is **not** a branch
    - $\text{Predicted PC} = (\text{BTB}[\text{PC}].\text{tag} == \text{PC}) ? \text{BTB}[\text{PC}].\text{target} : \text{PC}+4$



# Branch Direction Prediction

# Revisiting Branch Prediction Components



- Step #1: is it a branch?
  - Easy after decode... during fetch: predictor
- Step #2: is the branch taken or not taken?
  - **Direction predictor**
- Step #3: if the branch is taken, where does it go?
  - Branch target predictor (BTB)
  - Supplies target PC if branch is taken

# Simple Branch Direction Prediction

---

- Compile time (static)
  - Always not taken
  - Always taken
  - BTFN (Backward taken, forward not taken)
  - Profile based (likely direction)
  - Program analysis based (likely direction)
- Run time (dynamic)
  - Last time prediction (single-bit)
  - Two-bit counter based prediction
  - Two-level prediction (global vs. local)
  - Hybrid
  - Advanced algorithms (e.g., using perceptrons)

# Static Branch Prediction (I)

---

- Always not-taken
  - Simple to implement: no need for BTB, no direction prediction
  - Low accuracy: ~30-40% (for conditional branches)
- Always taken
  - No direction prediction
  - Better accuracy: ~60-70% (for conditional branches)
    - Backward branches (i.e. loop branches) are usually taken
    - Backward branch: target address lower than branch PC
- Backward taken, forward not taken (BTFN)
  - Predict backward (loop) branches as taken, others not-taken
- Profile-based
  - Idea: Compiler determines likely direction for each branch using a profile run. Encodes that direction as a hint bit in the branch instruction format.

# Static Branch Prediction (II)

---

- **Program-based (or, program analysis based)**
  - Idea: Use heuristics based on program analysis to determine statically-predicted direction
  - Example opcode heuristic: Predict BLEZ as NT (negative integers used as error values in many programs)
  - Example loop heuristic: Predict a branch guarding a loop execution as taken (i.e., execute the loop)
  - Pointer and FP comparisons: Predict not equal
- + Does not require profiling
- Heuristics might be not representative or good
- Requires compiler analysis and ISA support (ditto for other static methods)
- Ball and Larus, "Branch prediction for free," PLDI 1993.
  - 20% misprediction rate

# Static Branch Prediction (III)

---

- **Programmer-based**
    - Idea: Programmer provides the statically-predicted direction
    - Via *pragmas* in the programming language that qualify a branch as likely-taken versus likely-not-taken
- + Does not require profiling or program analysis
- + Programmer may know some branches and their program better than other analysis techniques
- Requires programming language, compiler, ISA support
- Burdens the programmer?

# Pragmas

---

- Idea: **Keywords that enable a programmer to convey hints to lower levels of the transformation hierarchy**
- `if (likely(x)) { ... }`
- `if (unlikely(error)) { ... }`
- Many other hints and optimizations can be enabled with pragmas
  - E.g., whether a loop can be parallelized
  - **#pragma omp parallel**
  - **Description**
    - The omp parallel directive explicitly instructs the compiler to parallelize the chosen segment of code.



# Static Branch Prediction

---

- All previous techniques can be combined
  - Profile based
  - Program based
  - Programmer based
- What is the common disadvantage of all three techniques?
  - Cannot adapt to dynamic changes in branch behavior
    - This can be mitigated by a dynamic compiler, but not at a fine granularity (and a dynamic compiler has its overheads...)
    - What is a Dynamic Compiler?
      - A compiler that generates code at runtime
      - Java JIT (just in time) compiler, Microsoft CLR (common lang. runtime)

# More Sophisticated Direction Prediction

---

- Compile time (static)
  - Always not taken
  - Always taken
  - BTFN (Backward taken, forward not taken)
  - Profile based (likely direction)
  - Program analysis based (likely direction)
- Run time (dynamic)
  - Last time prediction (single-bit)
  - Two-bit counter based prediction
  - Two-level prediction (global vs. local)
  - Hybrid
  - Advanced algorithms (e.g., using perceptrons)

# Dynamic Branch Prediction

---


- Idea: Predict branches based on dynamic information (collected at run-time)
- Advantages
  - + Prediction based on history of the execution of branches
    - + It can adapt to dynamic changes in branch behavior
  - + No need for static profiling: input set representativeness problem goes away
- Disadvantages
  - More complex (requires additional hardware)

# Last Time Predictor

- Last time predictor

- Single bit per branch (stored in BTB)
- 1 bit for prediction (0 = N, 1 = T)
- Indicates which direction branch went last time it executed

TTTTTTTTTTNNNNNNNNNN → 90% accuracy



# Last Time Predictor

- Problem: **inner loop branch** below

```
for (i=0;i<100;i++)
  for (j=0;j<3;j++)
    // whatever
```

  - Two “built-in” mis-predictions per inner loop iteration
  - Branch predictor “changes its mind too quickly”

Time	State	Prediction	Outcome	Result?
1	N	N	T	Wrong
2	T	T	T	Correct
3	T	T	T	Correct
4	T	T	N	Wrong
5	N	N	T	Wrong
6	T	T	T	Correct
7	T	T	T	Correct
8	T	T	N	Wrong
9	N	N	T	Wrong
10	T	T	T	Correct
11	T	T	T	Correct
12	T	T	N	Wrong


# Improving the Last Time Predictor

---

- Problem: A last-time predictor changes its prediction from  $T \rightarrow NT$  or  $NT \rightarrow T$  too quickly
  - even though the branch may be mostly taken or mostly not taken
- Solution Idea: Add hysteresis to the predictor so that prediction does not change on a single different outcome
  - Use two bits to track the history of predictions for a branch instead of a single bit
  - Can have 2 states for T or NT instead of 1 state for each
- Smith, "A Study of Branch Prediction Strategies," ISCA 1981.

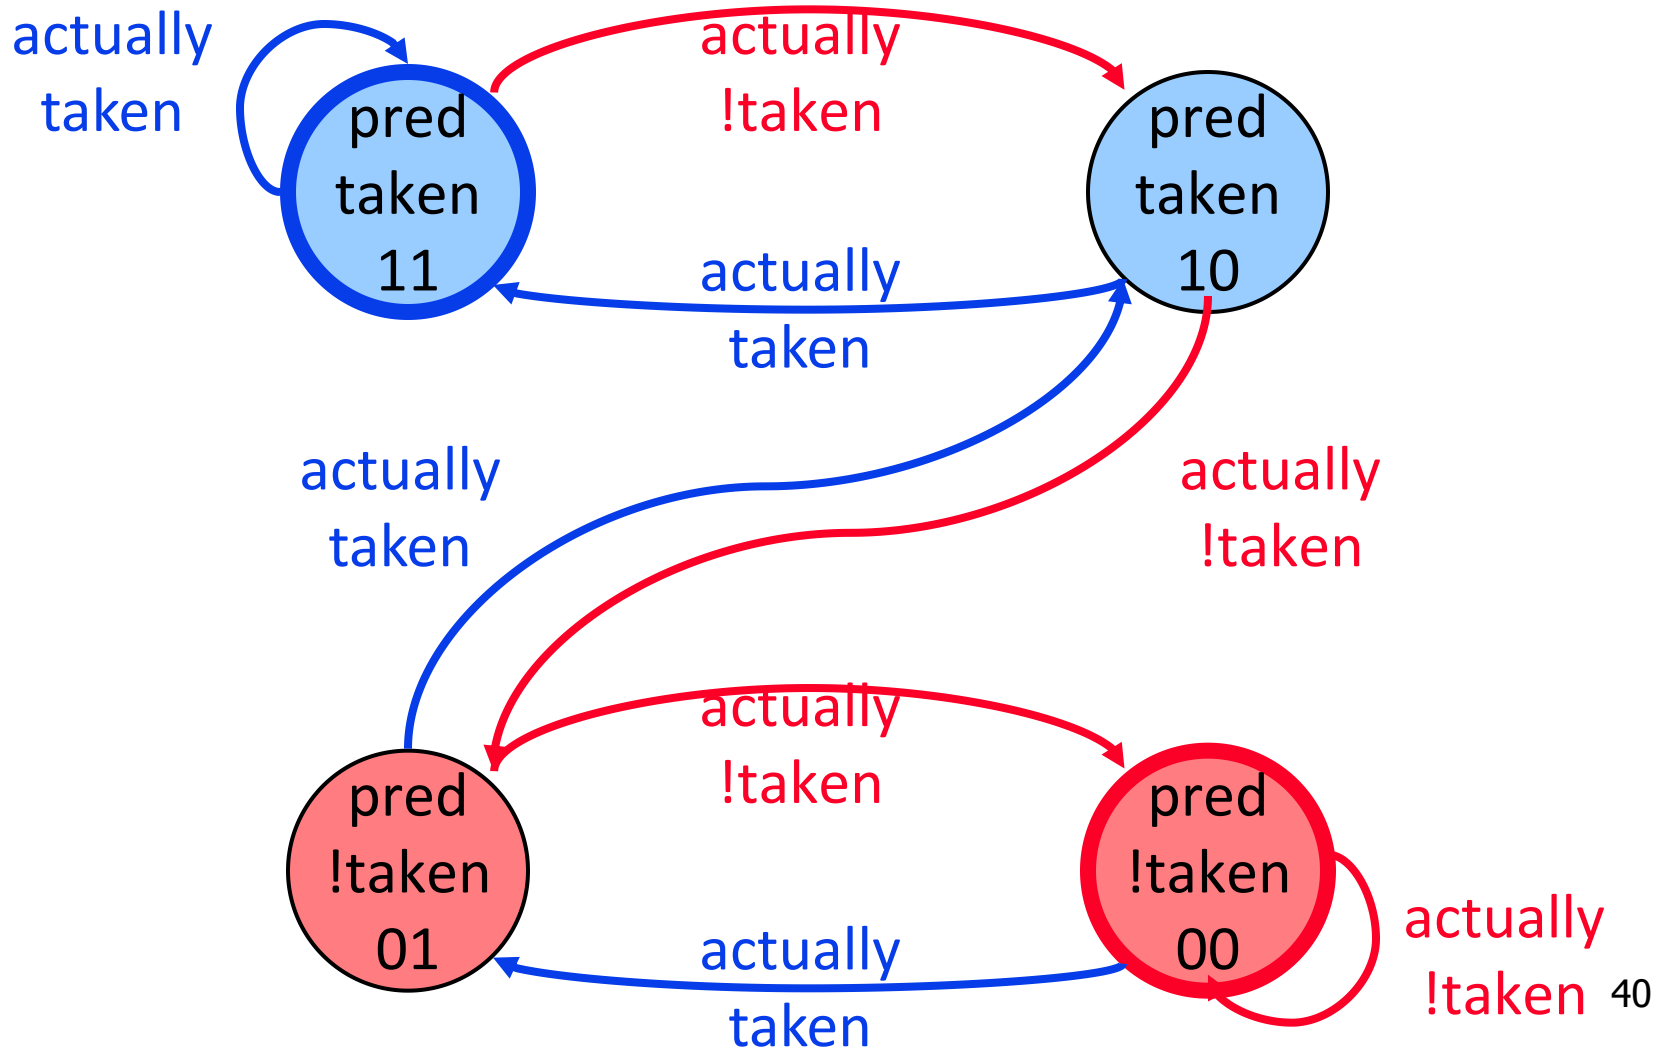
# Two-Bit Counter Based Prediction

- Each branch associated with a two-bit counter
- One more bit provides hysteresis
- 2 bits: (00->N, 01->n, 10->t, 11->T)



# State Machine for 2-bit Saturating Counter

- Counter using *saturating arithmetic*
  - Arithmetic with maximum and minimum values





# Two-Bit Counter Based Prediction

- A strong prediction does not change with one single different outcome
- Accuracy for a loop with N iterations =  $(N-1)/N$

TNTNTNTNTNTNTNTNTNTN →  
50% accuracy

(assuming counter initialized to weakly taken)

- + Better prediction accuracy
- More hardware cost (but counter can be part of a BTB entry)

Time	State	Prediction	Outcome	Result?
1	N	N	T	Wrong
2	n	N	T	Wrong
3	t	T	T	Correct
4	T	T	N	Wrong
5	t	T	T	Correct
6	T	T	T	Correct
7	T	T	T	Correct
8	T	T	N	Wrong
9	t	T	T	Correct
10	T	T	T	Correct
11	T	T	T	Correct
12	T	T	N	Wrong

# Is This Good Enough?

---

- ~85-90% accuracy for **many** programs with 2-bit counter based prediction (also called **bimodal prediction**)
- Is this good enough?
- How big is the branch problem?

# Importance of The Branch Problem

---

- Assume  $N = 20$  (20 pipe stages),  $W = 5$  (5 wide fetch)
- Assume: 1 out of 5 instructions is a branch
- Assume: Each 5 instruction-block ends with a branch
- How long does it take to fetch 500 instructions?
  - ❑ 100% accuracy
    - 100 cycles (all instructions fetched on the correct path)
    - No wasted work;  $IPC = 500/100$
  - ❑ 90% accuracy
    - 100 (correct path) +  $20 * 10$  (wrong path) = 300 cycles
    - 200% extra instructions fetched;  $IPC = 500/300$
  - ❑ 85% accuracy
    - 100 (correct path) +  $20 * 15$  (wrong path) = 400 cycles
    - 300% extra instructions fetched;  $IPC = 500/400$
  - ❑ 80% accuracy
    - 100 (correct path) +  $20 * 20$  (wrong path) = 500 cycles
    - 400% extra instructions fetched;  $IPC = 500/500$

# Can We Do Better: Two-Level Prediction

---

- Last-time and 2BC predictors exploit “last-time” predictability
- Realization 1: A branch’s outcome can be correlated with other branches’ outcomes
  - Global branch correlation
- Realization 2: A branch’s outcome can be correlated with past outcomes of the same branch (other than the outcome of the branch “last-time” it was executed)
  - Local branch correlation

# Global Branch Correlation (I)

---

- Recently executed branch outcomes in the execution path are correlated with the outcome of the next branch

- If first branch not taken, second also not taken

```
if (cond1)
...
if (cond1 AND cond2)
```

- If first branch taken, second definitely not taken

```
branch Y: if (cond1) a = 2;
...
branch X: if (a == 0)
```

# Global Branch Correlation (II)

---

- If Y and Z both taken, then X also taken
- If Y or Z not taken, then X also not taken

```
branch Y: if (cond1)
...
branch Z: if (cond2)
...
branch X: if (cond1 AND
cond2)
```


# Capturing Global Branch Correlation

---

- Idea: Associate branch outcomes with “global T/NT history” of all branches
- Make a prediction based on the outcome of the branch the last time the same global branch history was encountered
- Implementation:
  - Keep track of the “global T/NT history” of all branches in a register → Global History Register (GHR)
  - Use GHR to index into a table that recorded the outcome that was seen for each GHR value in the recent past → Pattern History Table (table of 2-bit counters)

# Two Level Global Branch Prediction

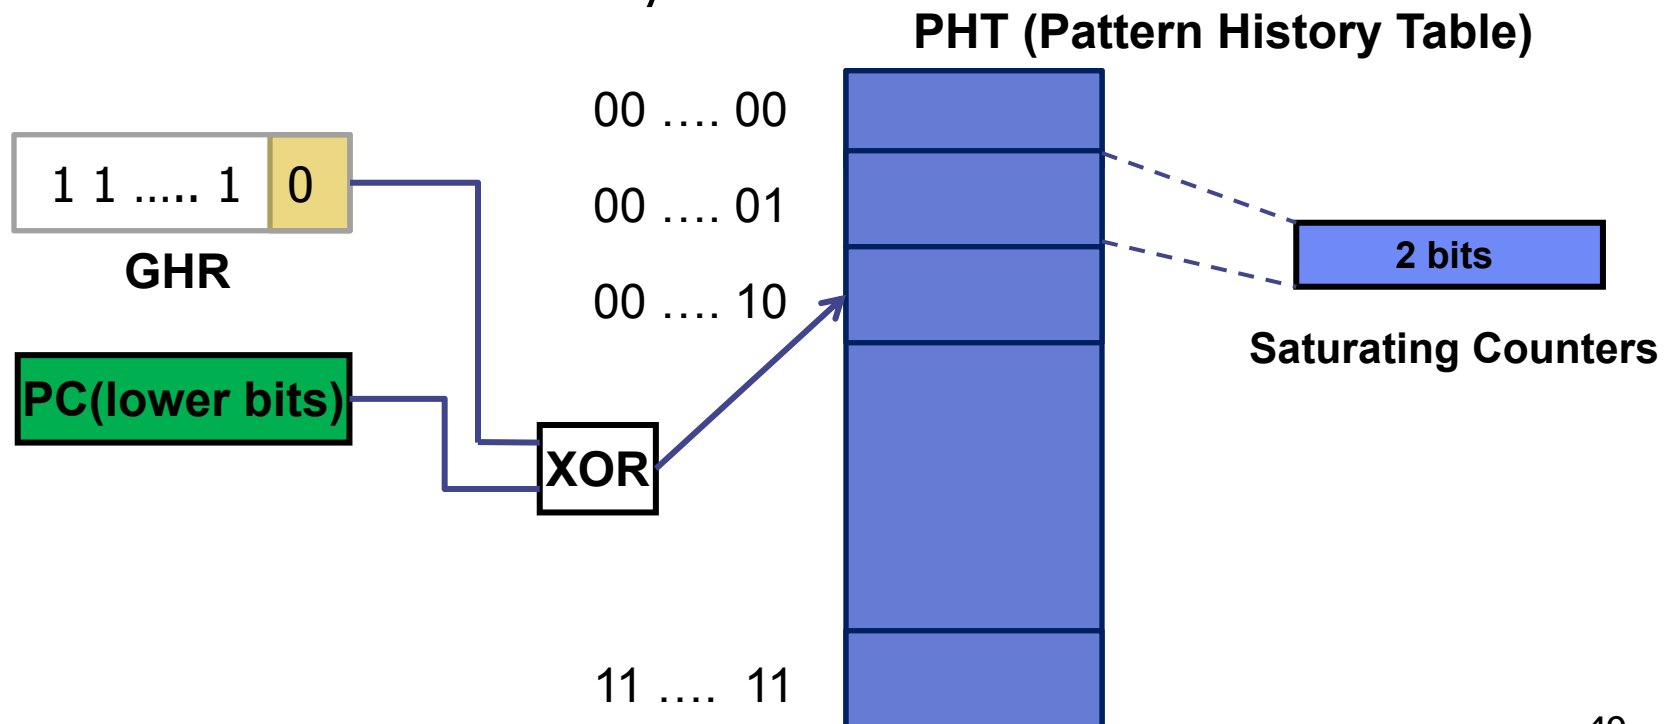
- First level: **Global branch history register** (N bits)
  - The direction of last N branches (1 bit per branch)
  - 1->tacken, 0->not taken
- Second level: **Table of saturating counters for each history entry**
  - The direction the branch took the last time the same history was seen






# Improving Global Predictor Accuracy (I)

- Idea: Add more context information to the global predictor to take into account which branch is being predicted
  - **Gshare predictor**: GHR hashed with the Branch PC
    - + More context information used for prediction
    - + Better utilization of the two-bit counter array
    - Increases access latency



# Improving Global Predictor Accuracy (II)

- Idea: Main separate PHT for different branches
  - use PC to determine the PHT
  - use GHR to determine the entry of the PHT
- + More context information used for prediction
- Increases hardware complexity



# Intel Pentium Pro Branch Predictor

---

- Two level global branch predictor
- 4-bit global history register
- Multiple pattern history tables (of 2 bit counters)
  - Which pattern history table to use is determined by lower order bits of the branch address

# Can We Do Better: Two-Level Prediction

---

- Last-time and 2BC predictors exploit “last-time” predictability
- Realization 1: A branch’s outcome can be correlated with other branches’ outcomes
  - Global branch correlation
- Realization 2: A branch’s outcome can be correlated with past outcomes of the same branch (other than the outcome of the branch “last-time” it was executed)
  - Local branch correlation

# Local Branch Correlation

---

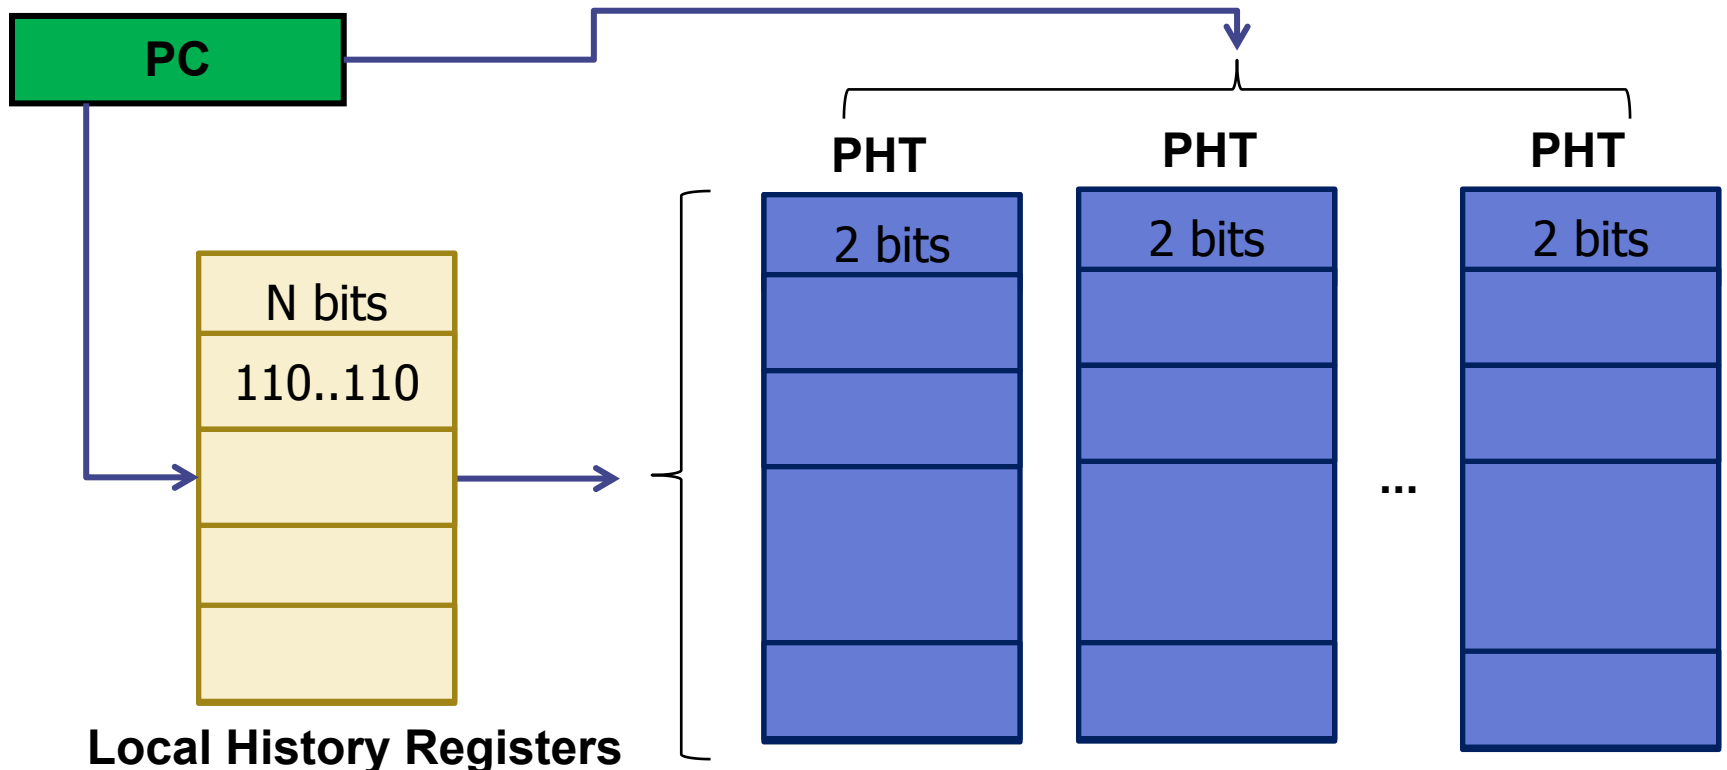
```
for(int i=1; i<4; i++)
```

Pattern: **TTTNTTTNTTTNTTTNTTTNTTTN**

Clearly, if we know the direction this branch had gone on the previous three executions, then we could always be able to predict the next branch direction

# Two Level Local Branch Prediction

- First level: **A set of local history registers** (N bits each)
  - Select the history register based on the PC of the branch
- Second level: **Table of saturating counters for each history entry**
  - PC determines the PHT of this branch
  - the local history register determines the PHT entry



# Can We Do Even Better?

---

- Predictability of branches varies
- Some branches are more predictable using local history
- Some using global
- For others, a simple two-bit counter is enough
- Yet for others, a single bit is enough
- Observation: There is heterogeneity in predictability behavior of branches
  - No one-size fits all branch prediction algorithm for all branches
- Idea: Exploit that heterogeneity by designing heterogeneous branch predictors

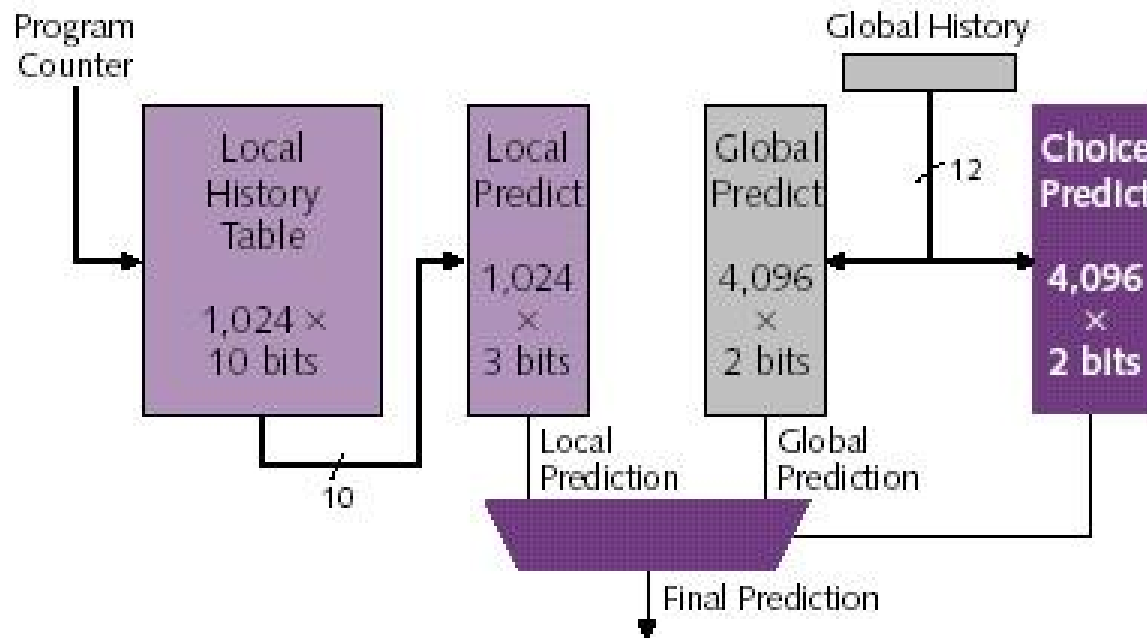
# Hybrid Branch Predictors

---

- Idea: Use more than one type of predictor (i.e., multiple algorithms) and select the “best” prediction
    - E.g., hybrid of 2-bit counters and global predictor
  - Advantages:
    - + Better accuracy: different predictors are better for different branches
    - + Reduced warmup time (faster-warmup predictor used until the slower-warmup predictor warms up)
  - Disadvantages:
    - Need “meta-predictor” or “selector”
    - Longer access latency
- McFarling, “Combining Branch Predictors,” DEC WRL Tech Report, 1993.



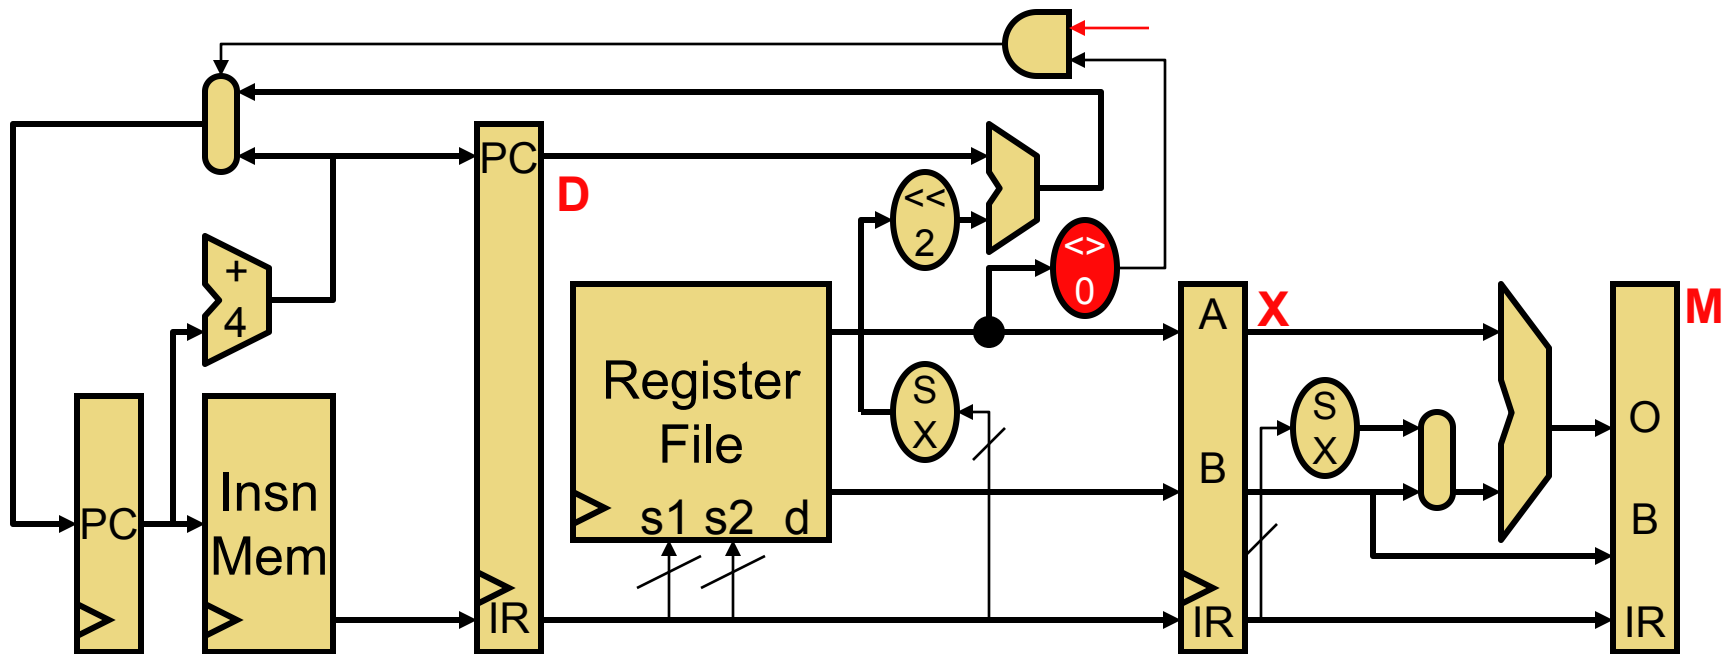
# Alpha 21264 Tournament Predictor



- Minimum branch penalty: 7 cycles
- Typical branch penalty: 11+ cycles
- 48K bits of target addresses stored in I-cache
- Predictor tables are reset on a context switch
- Kessler, "[The Alpha 21264 Microprocessor](#)," IEEE Micro 1999.

# Reducing Branch Penalty


# Reducing Penalty: Fast Branches



- **Fast branch:** can decide at D, not X
  - Test must be comparison to zero or equality, no time for ALU
  - + New taken branch penalty is 1
  - Additional insns (**s1t**) for more complex tests, must bypass to D too

# Reducing Penalty: Fast Branches

- **Fast branch**: targets control-hazard penalty
  - Basically, branch insns that can resolve at D, not X
    - Test must be comparison to zero or equality, **no time for ALU**
- + New taken branch penalty is 1
- Additional comparison insns (e.g., `cmplt`, `slt`) for complex tests
- Must bypass into decode stage now, too




# Fast Branch Performance

---

- Assume: Branch: 20%, 75% of branches are taken
  - $\text{CPI} = 1 + 20\% * 75\% * 1 = 1 + 0.20 * 0.75 * 1 = 1.15$
  - **15% slowdown**
- But wait, fast branches assume only simple comparisons
  - Fine for MIPS
  - But not fine for ISAs with “branch if \$1 > \$2” operations
- In such cases, say 25% of branches require an extra insn
  - $\text{CPI} = 1 + (20\% * 75\% * 1) + 20\% * 25\% * 1(\text{extra insn}) = 1.2$

# Putting It All Together

- BTB & branch direction predictor during fetch



- If branch prediction correct, no taken branch penalty

# Branch Prediction Performance

---

- Dynamic branch prediction
  - 20% of instruction branches
  - Simple predictor: branches predicted with 75% accuracy
    - $\text{CPI} = 1 + (20\% * 25\% * 2) = 1.1$
  - More advanced predictor: 95% accuracy
    - $\text{CPI} = 1 + (20\% * 5\% * 2) = 1.02$
- Branch mis-predictions still a big problem though
  - Pipelines are long: typical mis-prediction penalty is 10+ cycles
  - For cores that do more per cycle, predictions more costly (later)

# Branch Prediction in the Future


---

- Further improvement is difficult
  - 95%~98%
  - Branch predictor is getting larger, Alpha21464 has 48KB
- Improvement is significant
  - 5% precision improvement can reduce 50 pipeline flush per 10000 insts
- New applications and new architecture need new branch prediction approach
  - power aware
  - SMT
  - neural network
- Branch Prediction Competetion
  - supported by industry



# Summary

---



- Control hazards
- Branch prediction