

Computer Architecture

Lec 13: Data-Level Parallelism & Accelerators

This Lecture

- Data Level Parallelism
 - “medium-grained” parallelism between ILP and TLP
 - Still one flow of execution (unlike TLP)
 - Compiler/programmer must explicitly express it (unlike ILP)
- GPUs
 - Embrace data parallelism via “SIMT” execution model
 - Becoming more programmable all the time
- TPUs
 - Neural network accelerator
 - Fast matrix multiply machine

Data-Level Parallelism

How to Compute SAXPY Quickly?

- Performing the **same** operations on **many** data items
 - Example: SAXPY

```
for (I = 0; I < 1024; I++) {           L1: ldf [X+r1]->f1    // I is in r1
    Z[I] = A*X[I] + Y[I];             mulf f0,f1->f2    // A is in f0
}                                         ldf [Y+r1]->f3
                                         addf f2,f3->f4
                                         stf f4->[Z+r1]
                                         addi r1,4->r1
                                         blti r1,4096,L1
```

- Instruction-level parallelism (ILP) - fine grained
 - Loop unrolling with static scheduling –or– dynamic scheduling
 - Wide-issue superscalar (non-)scaling limits benefits
- Thread-level parallelism (TLP) - coarse grained
 - Multicore
- Can we do some “medium grained” parallelism?

Data-Level Parallelism

- **Data-level parallelism (DLP)**
 - Single operation repeated on multiple data elements
 - SIMD (**S**ingle-**I**nstruction, **M**ultiple-**D**ata)
 - Less general than ILP: parallel insns are all same operation
 - Exploit with **vectors**
- Old idea: Cray-1 supercomputer from late 1970s
 - Eight 64-entry x 64-bit floating point “vector registers”
 - 4096 bits (0.5KB) in each register! 4KB for vector register file
 - Special vector instructions to perform vector operations
 - Load vector, store vector (wide memory operation)
 - Vector+Vector or Vector+Scalar
 - addition, subtraction, multiply, etc.
 - In Cray-1, each instruction specifies 64 operations!
 - ALUs were expensive, so one operation per cycle (not parallel)

Example Vector ISA Extensions (SIMD)

- Extend ISA with vector storage ...
 - **Vector register**: fixed-size array of FP/int elements
 - **Vector length**: For example: 4, 8, 16, 64, ...
- ... and example operations for vector length of 4
 - Load vector: `ldf.v [X+r1]->v1`
`ldf [X+r1+0]->v10`
`ldf [X+r1+1]->v11`
`ldf [X+r1+2]->v12`
`ldf [X+r1+3]->v13`
 - Add two vectors: `addf.vv v1,v2->v3`
`addf v1i,v2i->v3i (where i is 0,1,2,3)`
 - Add vector to scalar: `addf-vs v1,f2,v3`
`addf v1i,f2->v3i (where i is 0,1,2,3)`
- Today's vectors: short (128-512 bits), but fully parallel

Example Use of Vectors – 4-wide

```
ldf [X+r1]->f1  
mulf f0,f1->f2  
ldf [Y+r1]->f3  
addf f2,f3->f4  
stf f4->[Z+r1]  
  
addi r1,4->r1  
blti r1,4096,L1
```

7x1024 instructions

```
ldf.v [X+r1]->v1  
mulf.vs v1,f0->v2  
ldf.v [Y+r1]->v3  
addf.vv v2,v3->v4  
stf.v v4,[Z+r1]  
  
addi r1,16->r1  
blti r1,4096,L1
```

7x256 instructions

(4x fewer instructions)

- Operations
 - Load vector: **ldf.v [X+r1]->v1**
 - Multiply vector to scalar: **mulf.vs v1,f0->v2**
 - Add two vectors: **addf.vv v1,v2->v3**
 - Store vector: **stf.v v1->[X+r1]**
- Performance?
 - Best case: 4x speedup
 - But, vector instructions don't always have single-cycle throughput
 - Execution width (implementation) vs vector width (ISA)

Vector Datapath & Implementation

- Vector insn. are just like normal insn... only “wider”
 - Single instruction fetch (no extra N^2 checks)
 - Wide register read & write (not multiple ports)
 - Wide execute: replicate floating point unit (same as superscalar)
 - Wide bypass (avoid N^2 bypass problem)
 - Wide cache read & write (single cache tag check)
- Execution width (implementation) vs vector width (ISA)
 - Example: Pentium 4 and “Core 1” executes vector ops at half width
 - “Core 2” executes them at full width
- Because they are just instructions...
 - ...superscalar execution of vector instructions
 - Multiple n-wide vector instructions per cycle

Vector Insn Sets for Different ISAs

- x86
 - Intel and AMD: MMX, SSE, SSE2, SSE3, SSE4, AVX, AVX2
 - currently: AVX 512 offers 512b vectors
- PowerPC
 - AltiVEC/VMX: 128b
- ARM
 - NEON: 128b
 - Scalable Vector Extensions (SVE): up to 2048b
 - implementation is narrower than this!
 - makes vector code portable

GPU

Part I: Graphics + GPU history

What GPUs were originally designed to do: 3D rendering

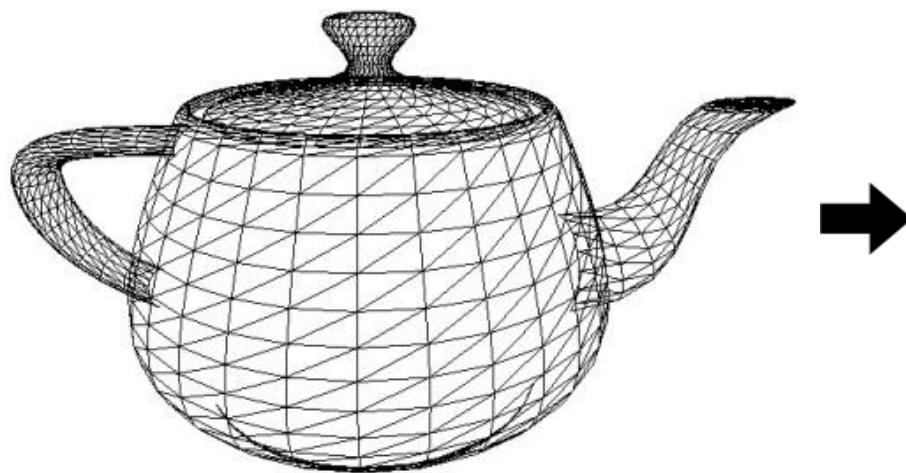


Image credit: Henrik Wann Jensen

Input: description of a scene:

3D surface geometry (e.g., triangle mesh)

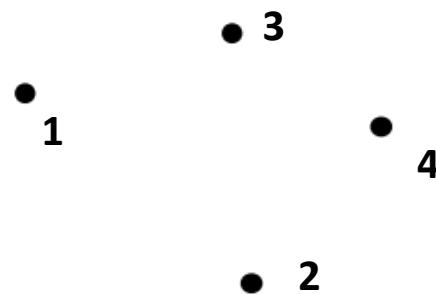
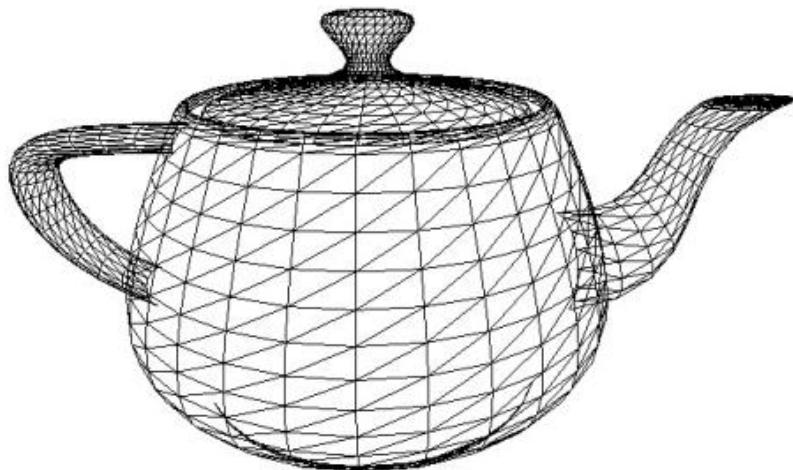
surface materials, lights, camera, etc.

Output: image of the scene

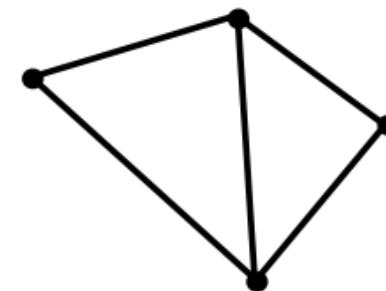
Simple definition of rendering task: computing how each triangle in 3D mesh contributes to appearance of each pixel in the image?

Real-time graphics primitives (entities)

Represent surface as a 3D triangle mesh

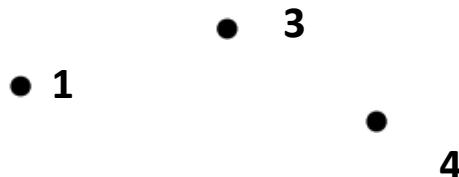


Vertices
(points in space)

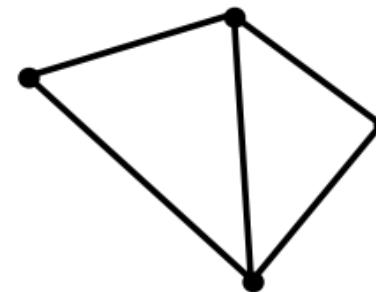


Primitives
(e.g., triangles, points, lines)

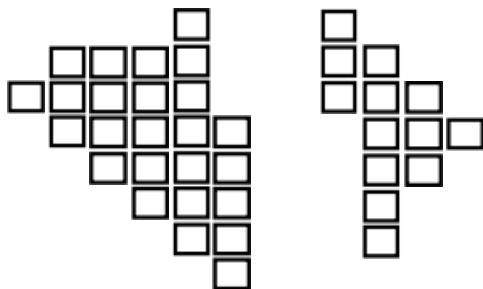
Real-time graphics primitives (entities)



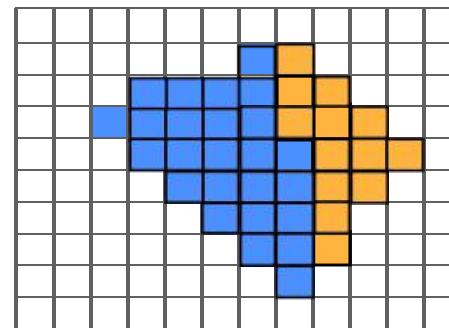
Vertices
(points in space)



Primitives
(e.g., triangles, points, lines)



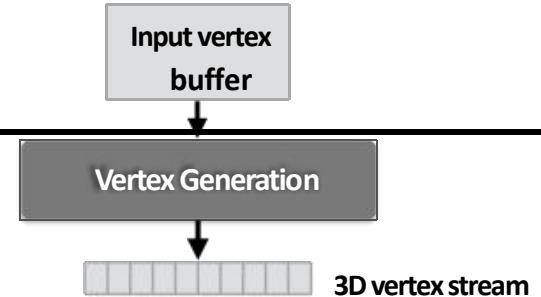
Fragments



Pixels (in an image)

Rendering a picture

Input: a list of vertices in 3D space (and their connectivity into primitives)



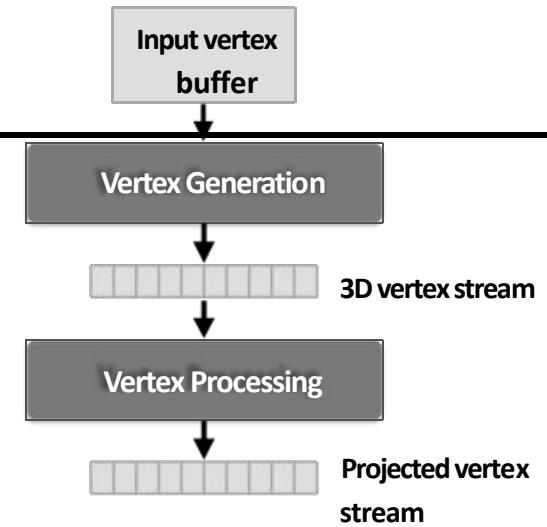
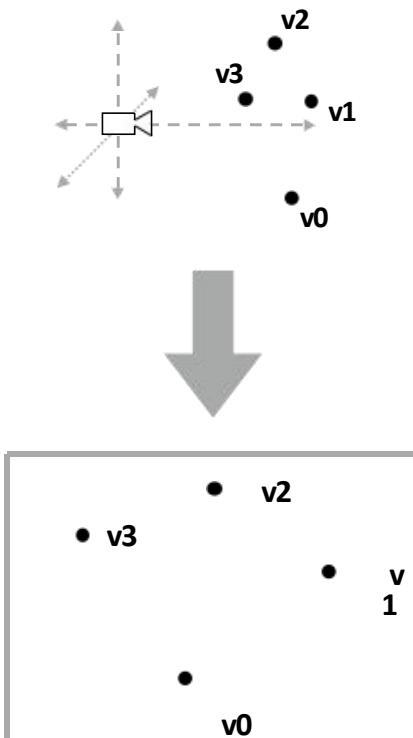
Example: every three vertices defines a triangle

```
list_of_positions =  
{  v1x, v1y, v1x,  
  v0x, v0y, v0z,  
  v2x, v2y, v2z,  
  v3x, v3y, v3x }  
};
```

triangle 0 = {v0, v1, v2}
triangle 1 = {v1, v2, v3}

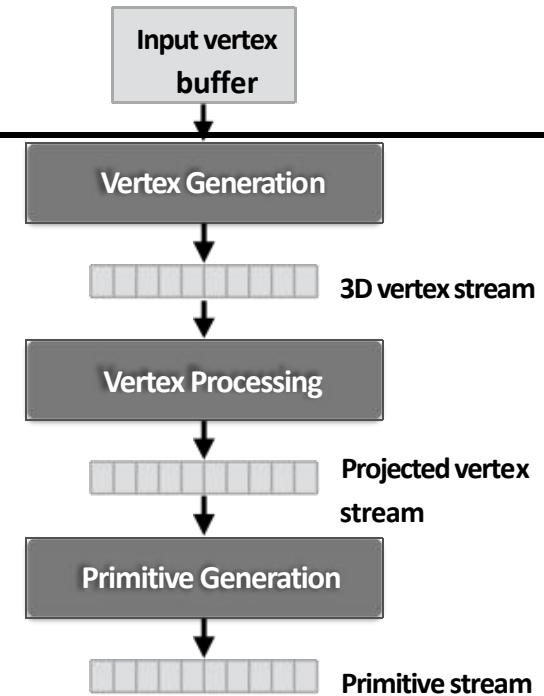
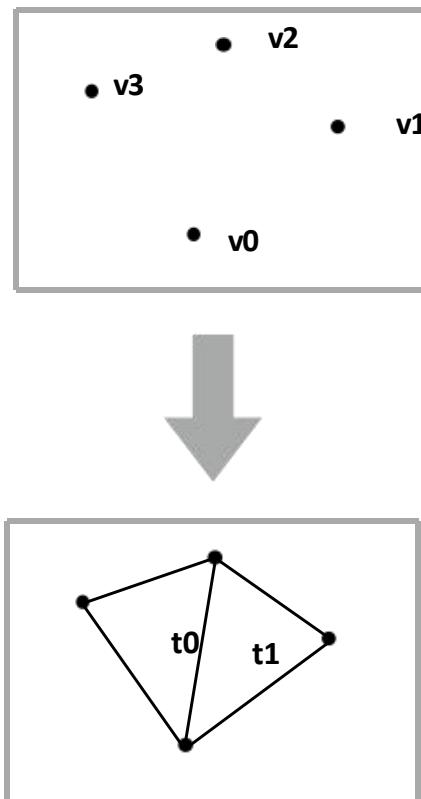
Rendering a picture

Step 1: given a scene camera position,
compute where the vertices lie on screen



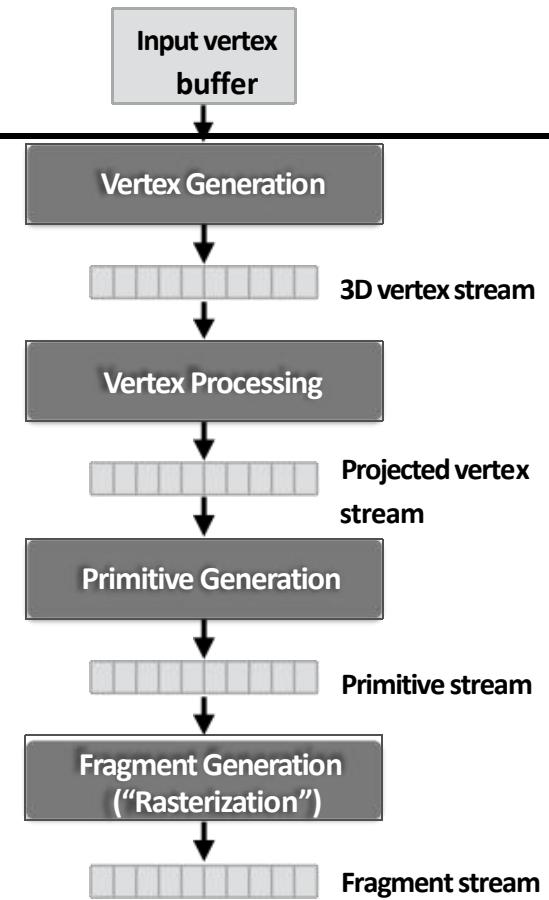
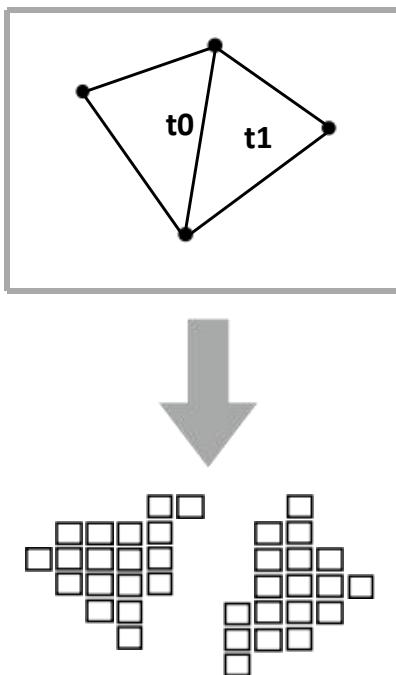
Rendering a picture

Step 2: group vertices into primitives



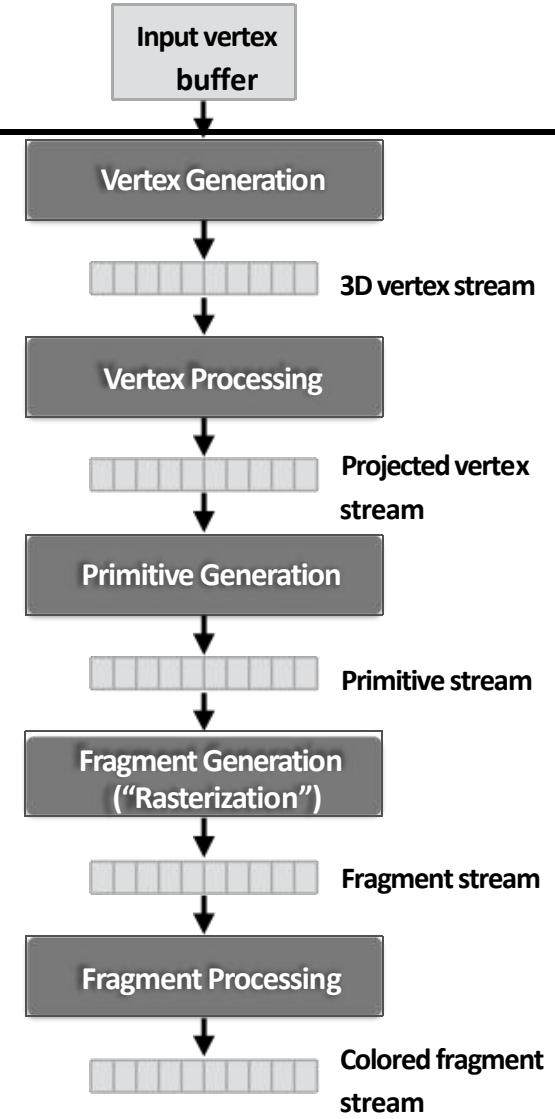
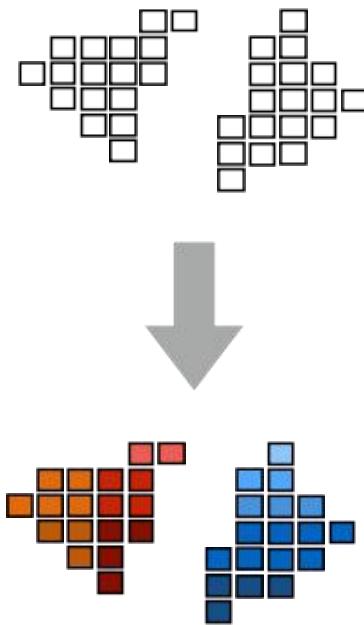
Rendering a picture

Step 3: generate one fragment for each pixel a primitive overlaps



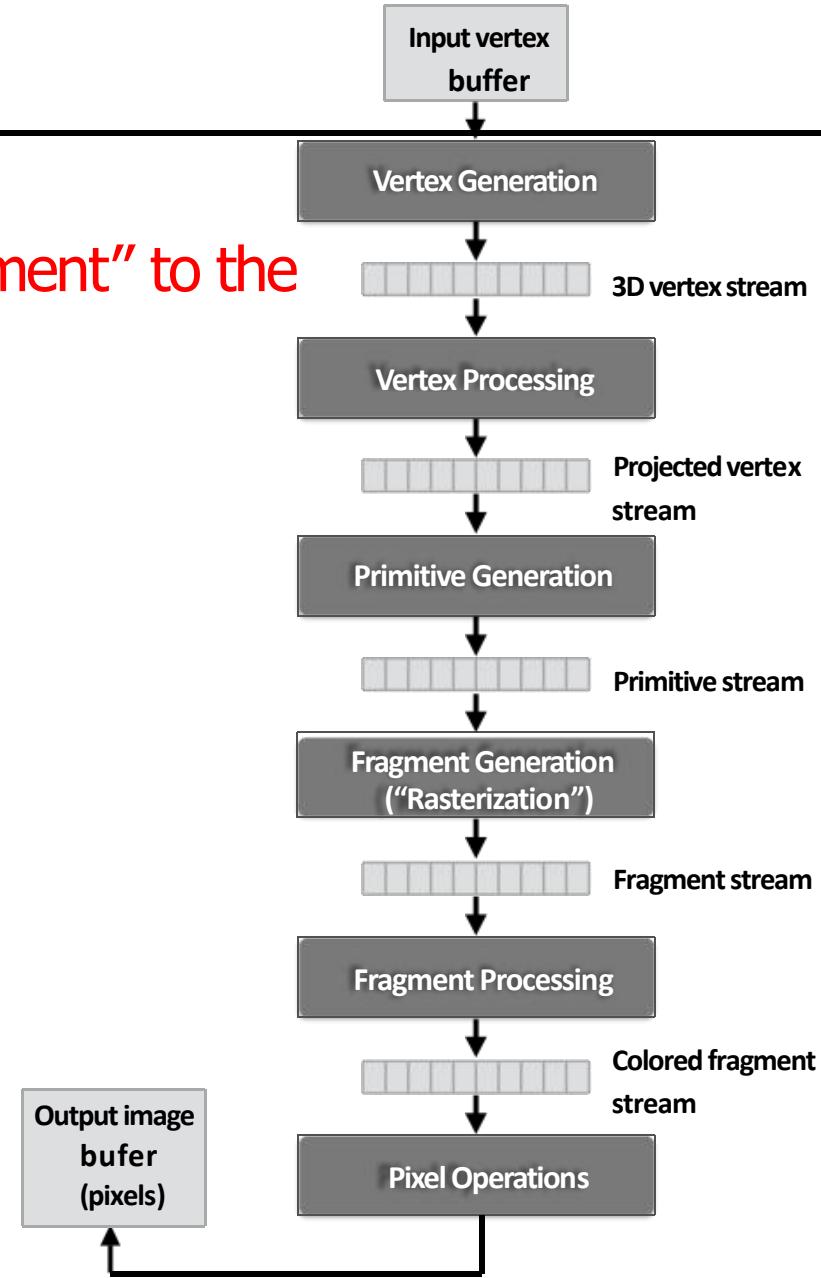
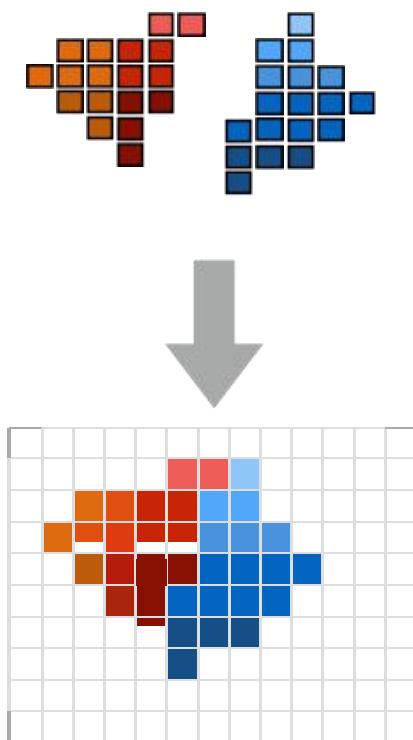
Rendering a picture

Step 4: compute color of primitive for each fragment (based on scene lighting and primitive material properties)



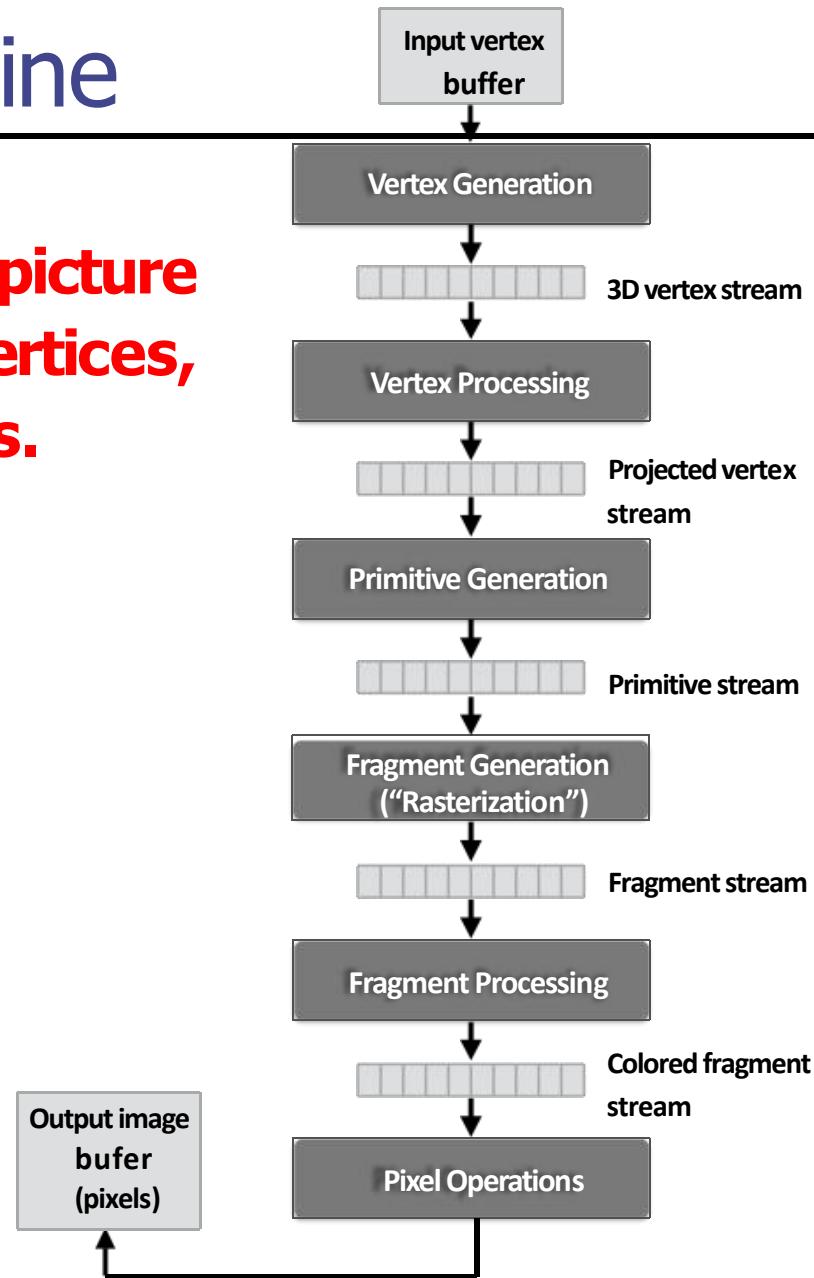
Rendering a picture

Step 5: put color of the “closest fragment” to the camera in the output image



Real-time graphics pipeline

Abstracts process of rendering a picture as a sequence of operations on vertices, primitives, fragments, and pixels.

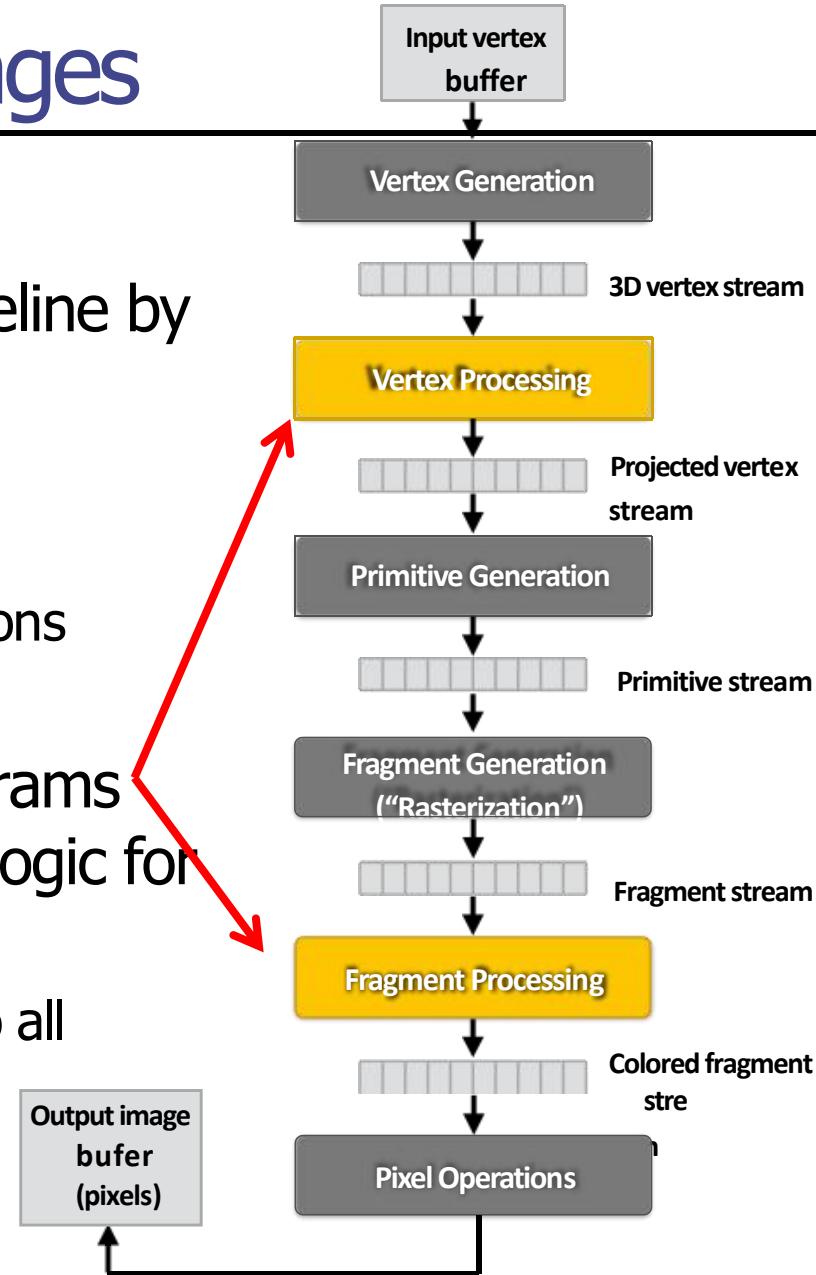


Early graphics programming (OpenGL API)

- Graphics programming APIs provided programmer mechanisms to set parameters of scene lights and materials
 - **glLight(light_id, parameter_id, parameter_value)**
 - Examples of light parameters: color, position, direction
 - **glMaterial(face, parameter_id, parameter_value)**
 - Examples of material parameters: color, shininess

Graphics shading languages

- Allow application to extend the functionality of the graphics pipeline by specifying materials and lights programmatically!
 - Support diversity in materials
 - Support diversity in lighting conditions
- Programmer provides mini-programs (“shaders”) that define pipeline logic for certain stages
 - Pipeline maps shader function onto all elements of input stream



Example fragment shader program

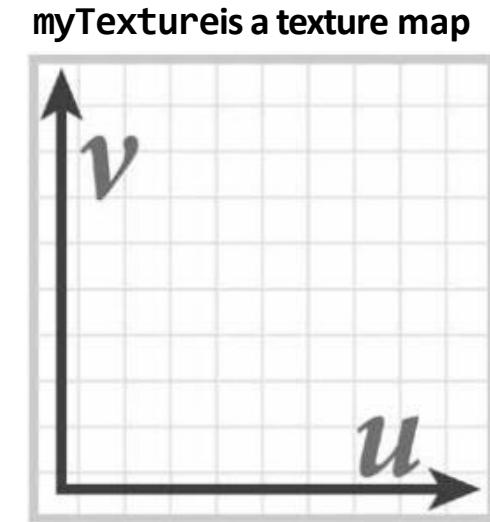
Run once per fragment (per pixel covered by a triangle)

OpenGL shading language (GLSL) shader program:
defines behavior of **fragment processing stage**

```
uniform sampler2D myTexture;
uniform float3 lightDir;
varying vec3 norm;
varying vec2 uv;

void myFragmentShader()
{
    vec3 kd = texture2D(myTexture, uv);
    kd *= clamp(dot(lightDir, norm), 0.0, 1.0);
    return vec4(kd, 1.0);
}
```

read-only global variables
per-fragment inputs



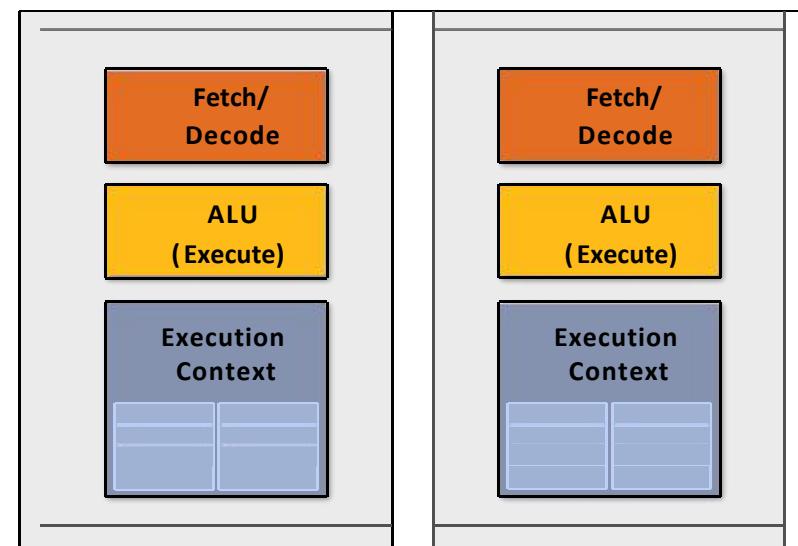
“fragment shader”
(a.k.a kernel function mapped
onto input fragment stream)

per-fragment output: RGBA surface color at pixel

Part II: GPU compute mode

Review: how to run code on a CPU

- Lets say a user wants to run a program on a multi-core CPU...
 - OS loads program text into memory
 - OS selects CPU execution context
 - OS interrupts processor, prepares execution context (sets contents of registers, program counter, etc. to prepare execution context)
- Go!
- Processor begins executing instructions from within the environment maintained in the execution context

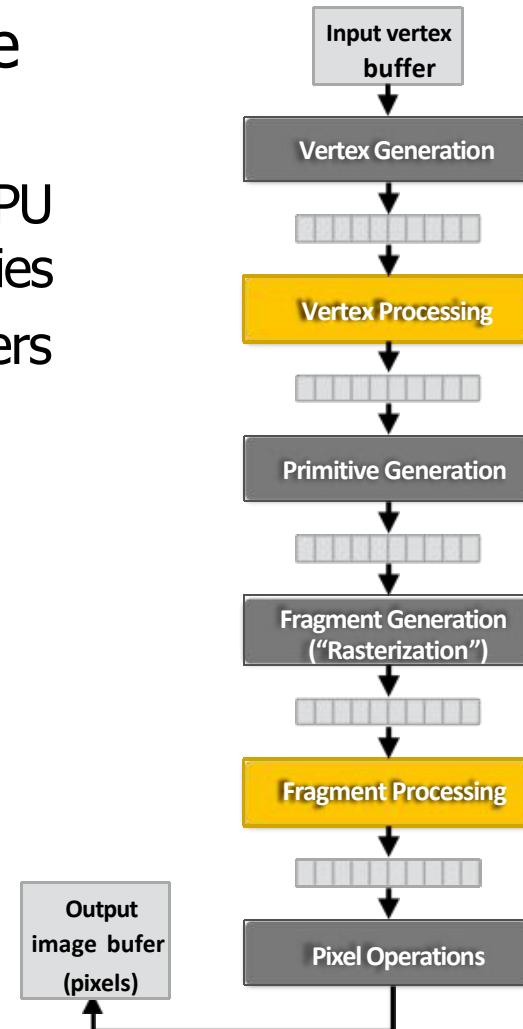


Multi-core CPU

How to run code on a GPU (prior to 2007)

- Lets say a user wants to draw a picture using a GPU...
 - Application (via graphics driver) provides GPU vertex and fragment shader program binaries
 - Application sets graphics pipeline parameters (e.g., output image size)
 - Application provides hardware a buffer of vertices
 - Go! `launch(myKernel, N)`

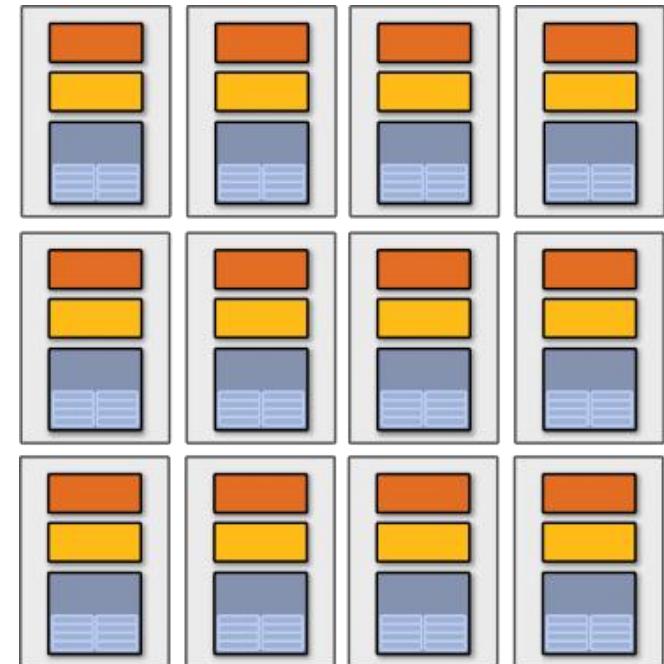
This was the only interface to GPU hardware. GPU hardware could only execute graphics pipeline computations.



NVIDIA Tesla architecture (2007)

First alternative, non-graphics-specific ("compute mode") interface to GPU hardware

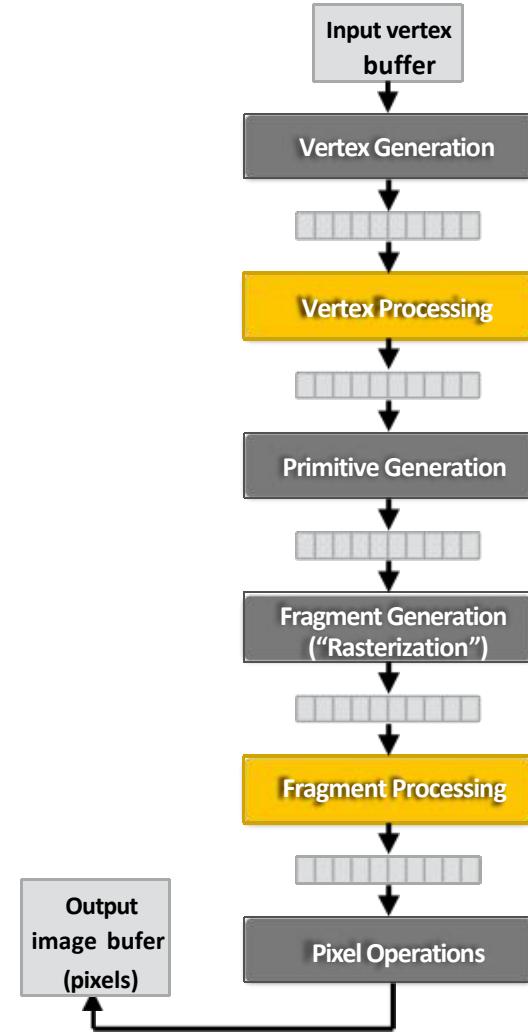
- Lets say a user wants to run a non-graphics program on the GPU's programmable cores...
 - Application can allocate buffers in GPU memory and copy data to/from buffers
 - Application (via graphics driver) provides GPU a single kernel program binary
 - Application tells GPU to run the kernel in an SIMD fashion ("run N instances")
 - Go! (`drawPrimitives(vertex_buffer)`)



Aside: interestingly, this is a far simpler operation than `drawPrimitives()`

CUDA programming language

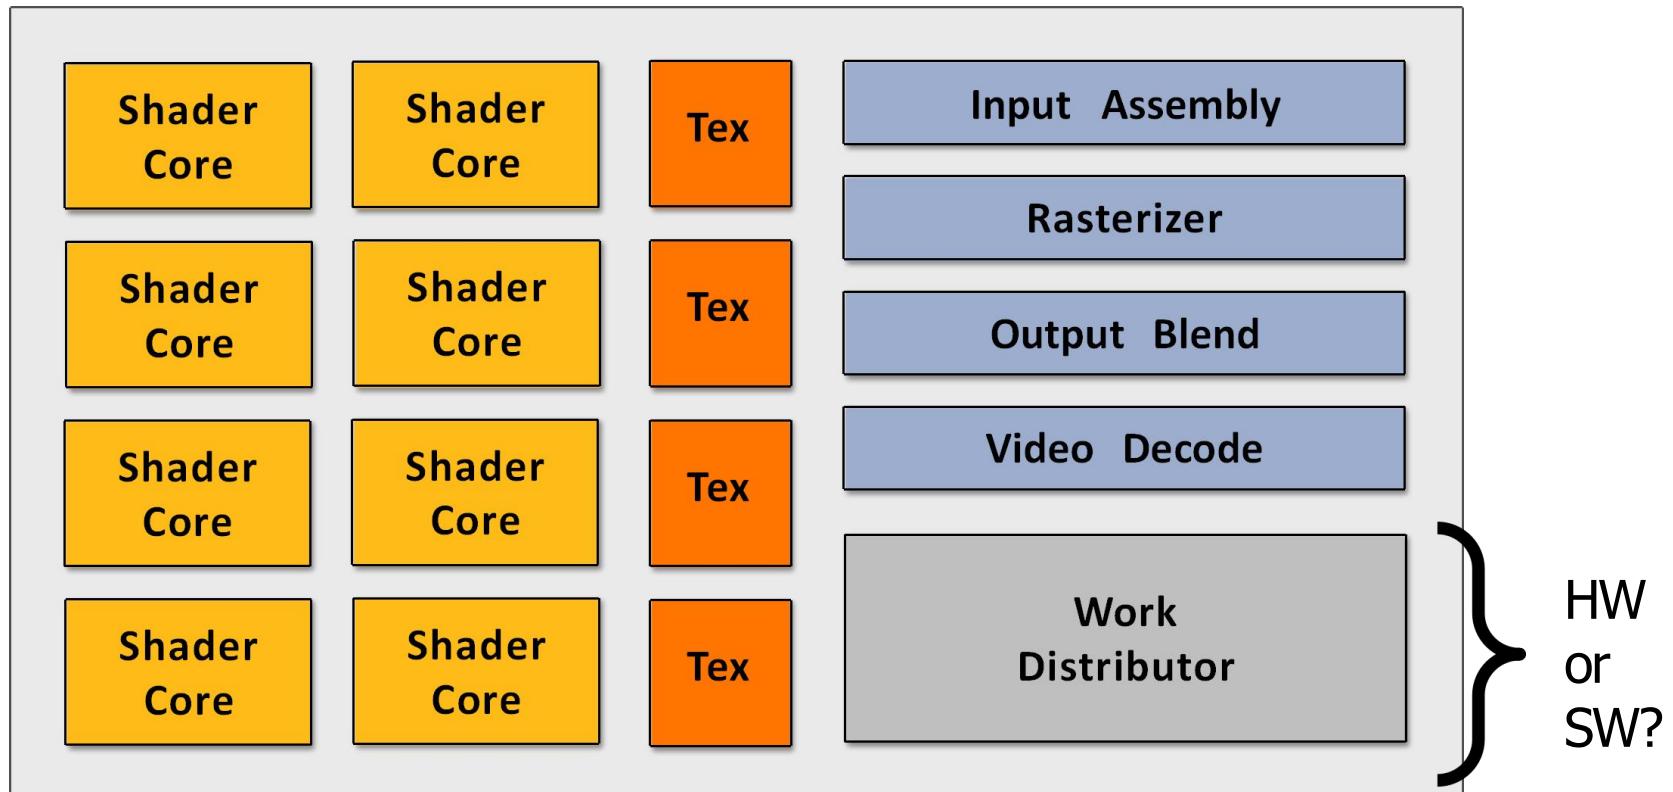
- Introduced in 2007 with NVIDIA Tesla architecture
- “C-like” language to express programs that run on GPUs using the compute-mode hardware interface
- Relatively low-level: CUDA’s abstractions closely match the capabilities/performance characteristics of modern GPUs (design goal: maintain low abstraction distance)
- Note: OpenCL is an open standards version of CUDA
 - CUDA only runs on NVIDIA GPUs
 - OpenCL runs on CPUs and GPUs from many vendors
 - Almost everything I say about CUDA also holds for OpenCL
 - CUDA is better documented



Part III: GPU architecture

What's in a GPU?

A GPU is a heterogeneous chip multi-processor (highly tuned for graphics)



A diffuse relectance shader

```
sampler mySamp;  
Texture2D<float3> myTex;  
float3 lightDir;  
  
float4 diffuseShader(float3 norm, float2 uv) {  
  
    float3 kd;  
  
    kd = myTex.Sample(mySamp, uv);  
  
    kd *= clamp( dot(lightDir, norm), 0.0, 1.0);  
  
    return float4(kd, 1.0);  
}
```

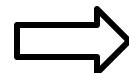
Shader programming model:

Fragments are processed
independently,
but there is no explicit parallel
programming

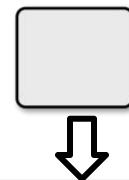
Compile shader

1 unshaded fragment input record

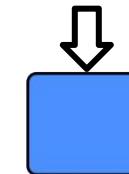
```
sampler mySamp;  
Texture2D<float3> myTex;  
float3 lightDir;  
  
float4 diffuseShader(float3 norm, float2 uv) {  
  
    float3 kd;  
  
    kd = myTex.Sample(mySamp, uv);  
  
    kd *= clamp( dot(lightDir,norm),0.0, 1.0);  
  
    return float4(kd, 1.0);  
}
```



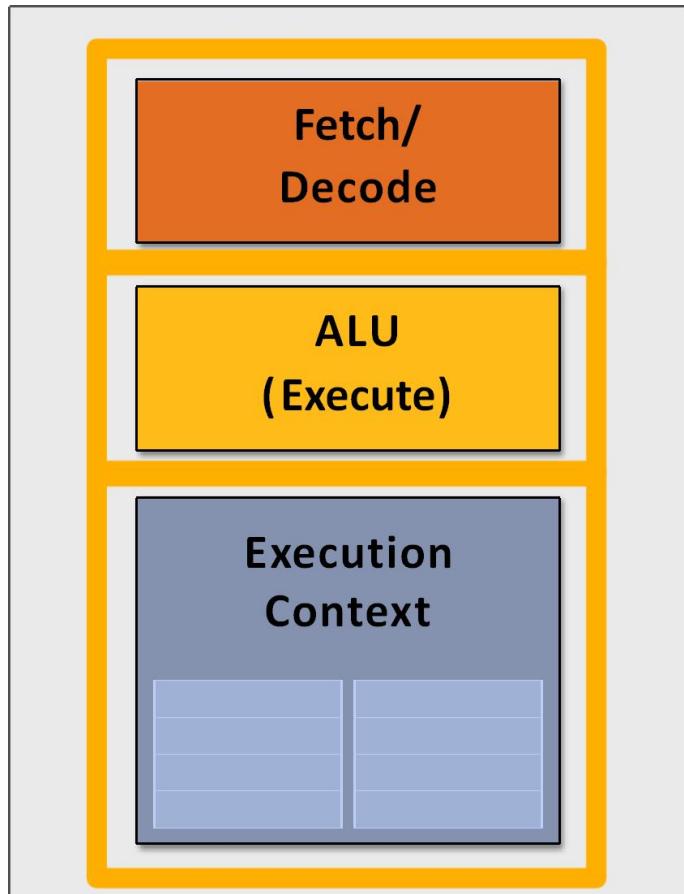
```
<diffuseShader>:  
sample r0, v4, t0, s0  
mul r3, v0, cb0[0]  
madd r3, v1, cb0[1], r3  
madd r3, v2, cb0[2], r3  
clmp r3, r3, l(0.0), l(1.0)  
mul o0, r0, r3  
mul o1, r1, r3  
mul o2, r2, r3  
mov o3, l(1.0)
```



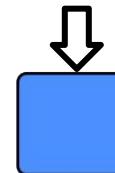
1 shaded fragment output record



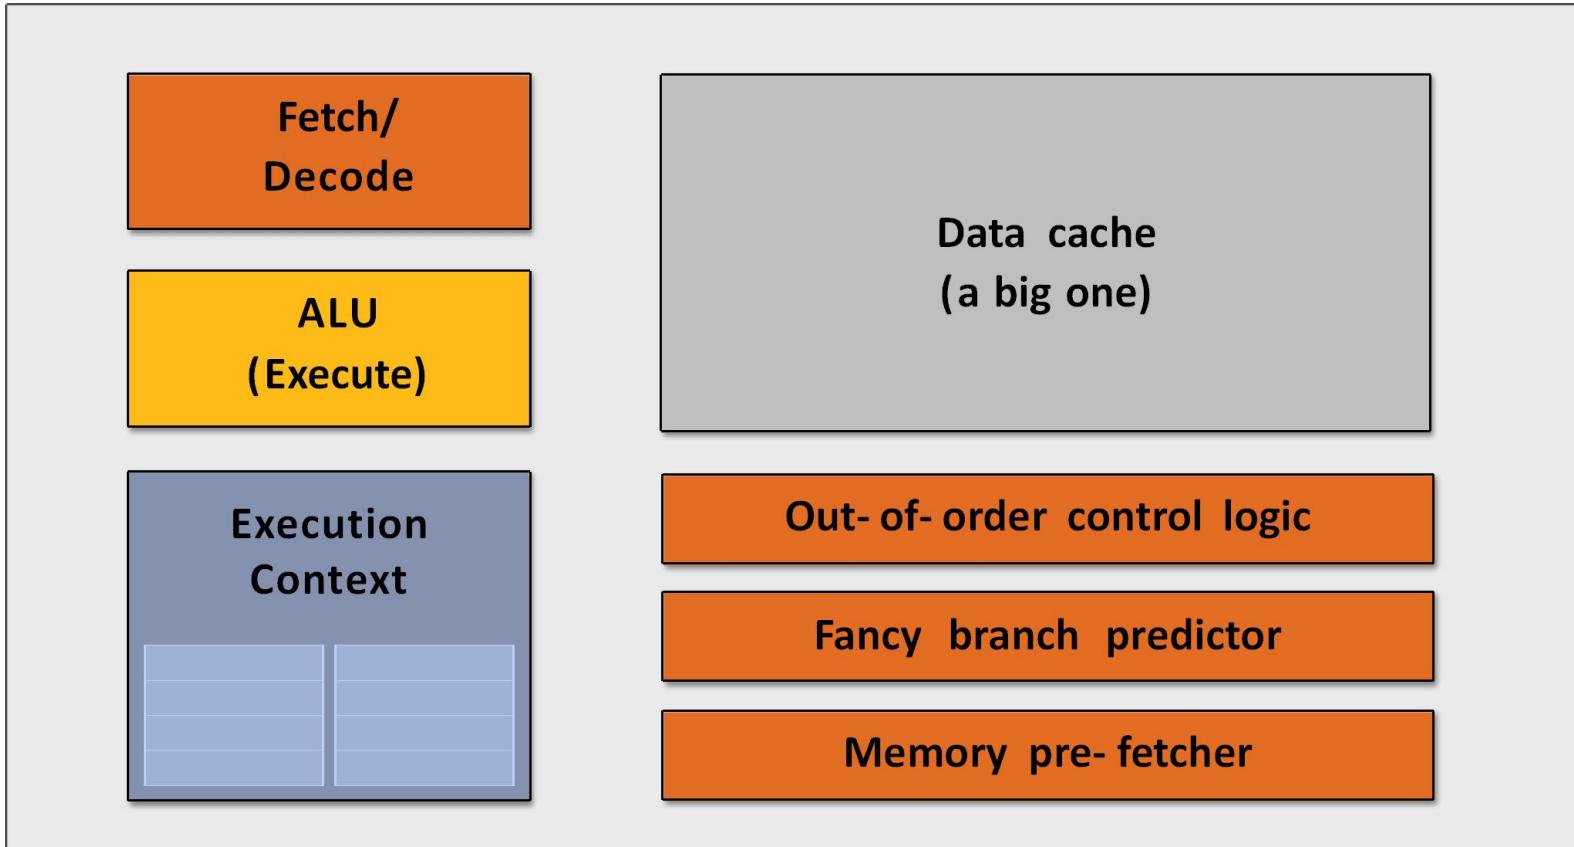
Execute shader



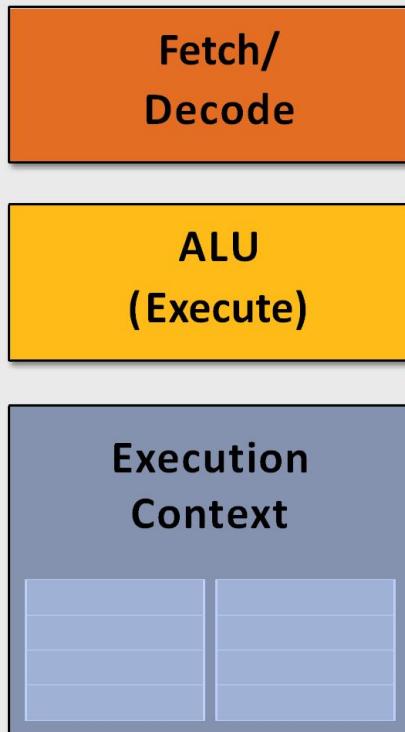
```
<diffuseShader>:  
sample r0, v4, t0, s0  
mul r3, v0, cb0[0]  
madd r3, v1, cb0[1], r3  
madd r3, v2, cb0[2], r3  
clmp r3, r3, 1(0.0), 1(1.0)  
mul o0, r0, r3  
mul o1, r1, r3  
mul o2, r2, r3  
mov o3, 1(1.0)
```



“CPU-style” cores



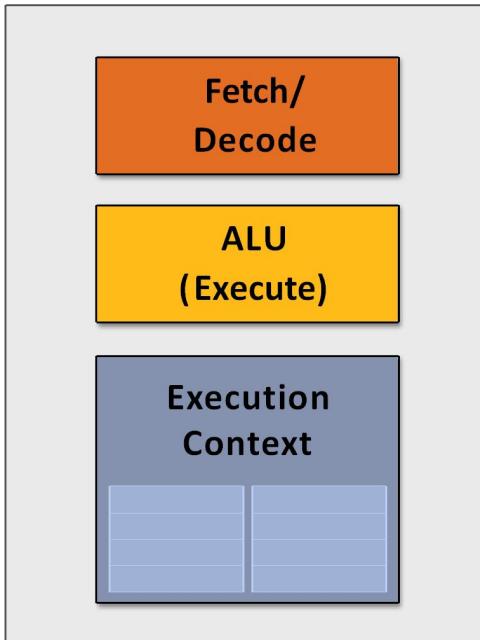
Slimming down



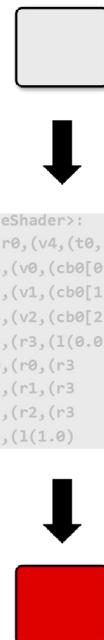
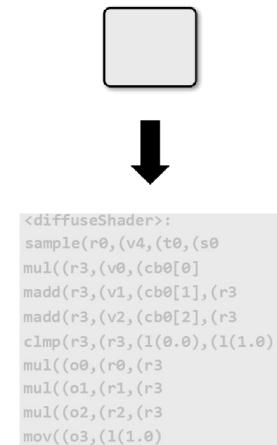
Idea #1:
Remove components that help a single instruction stream run fast

Two cores (two fragments in parallel)

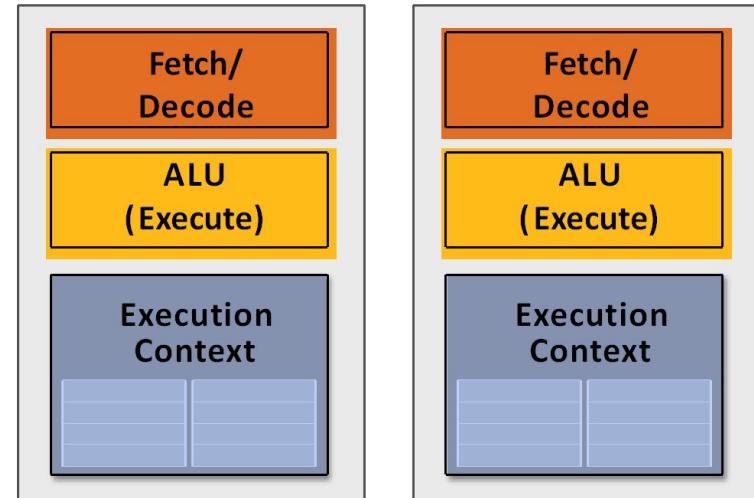
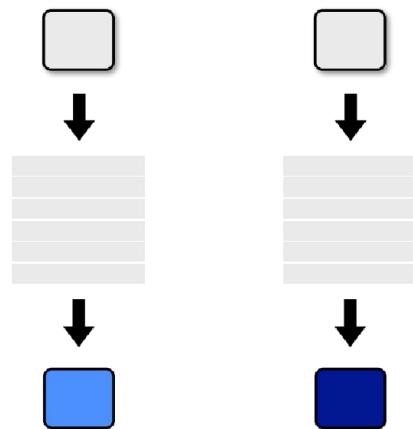
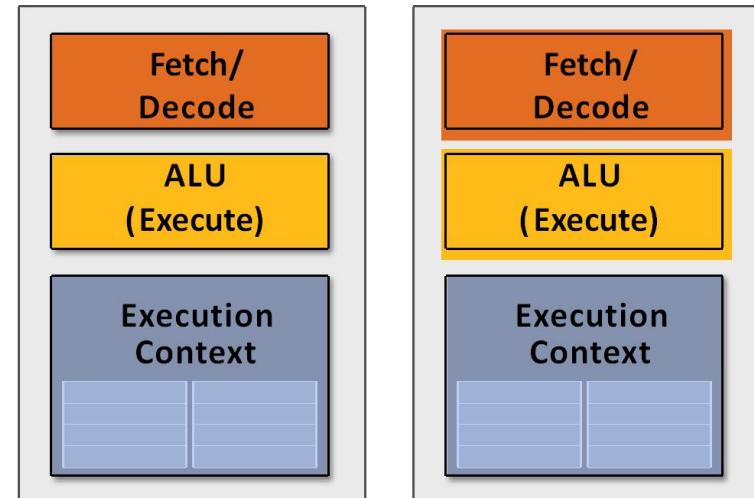
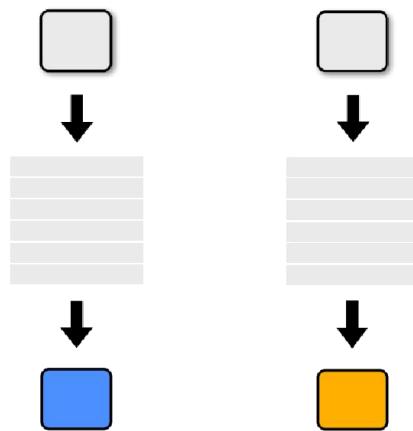
fragment 1



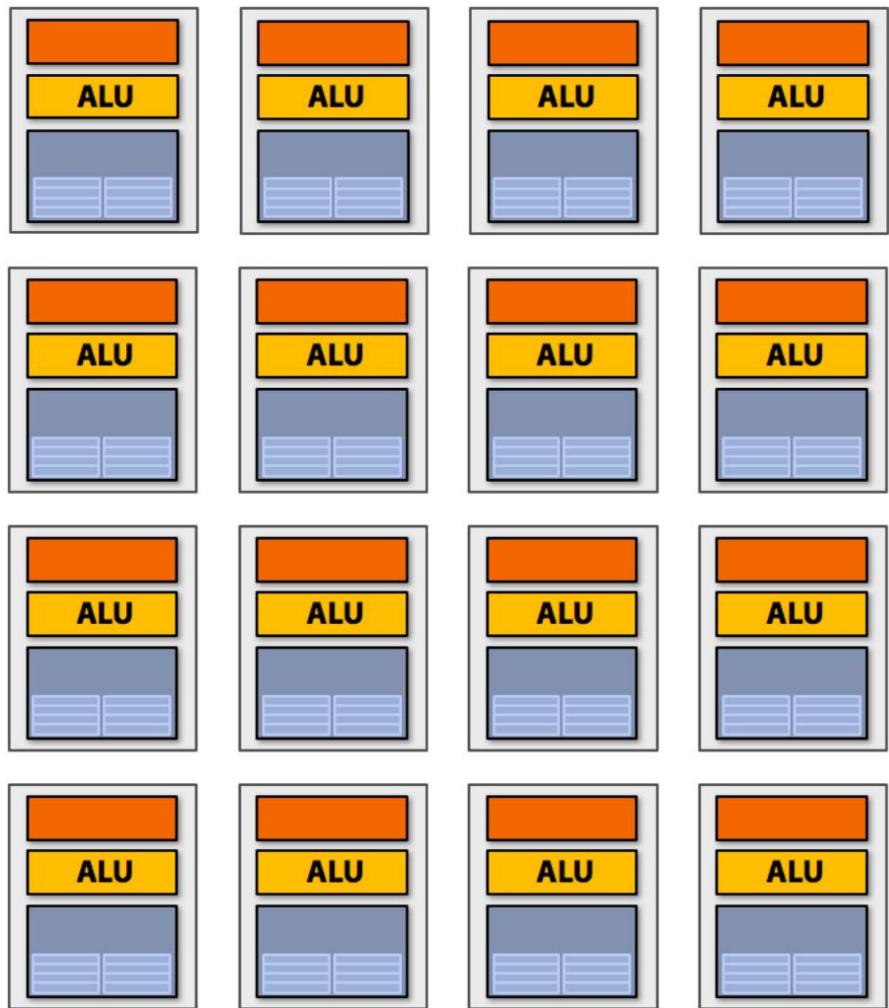
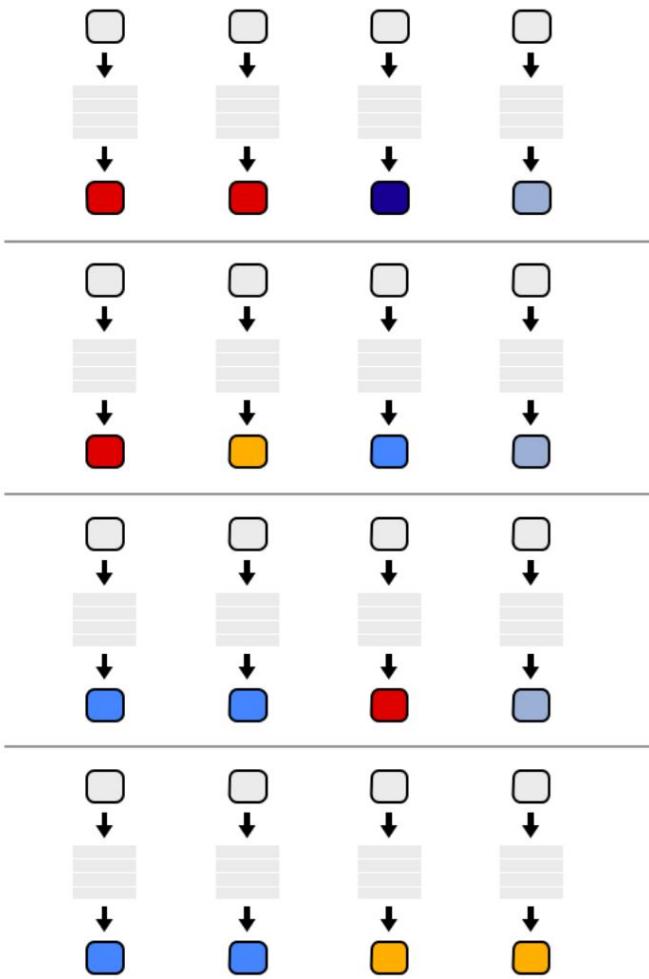
fragment 2



Four cores (four fragments in parallel)

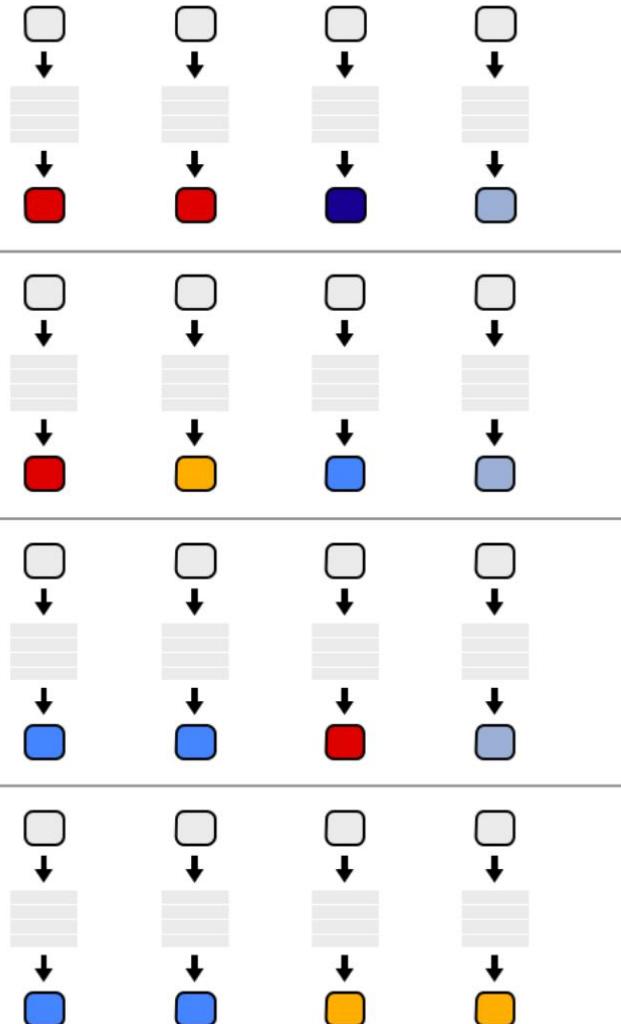


Sixteen cores (sixteen fragments in parallel)



16 cores = 16 simultaneous instruction streams

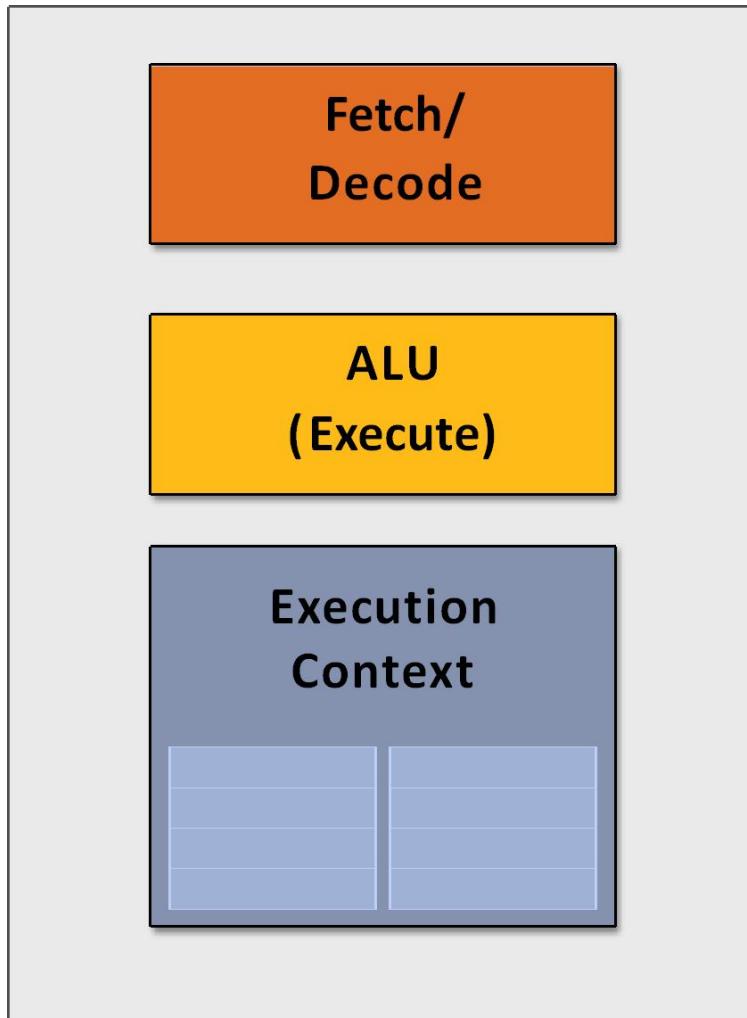
Instruction stream sharing



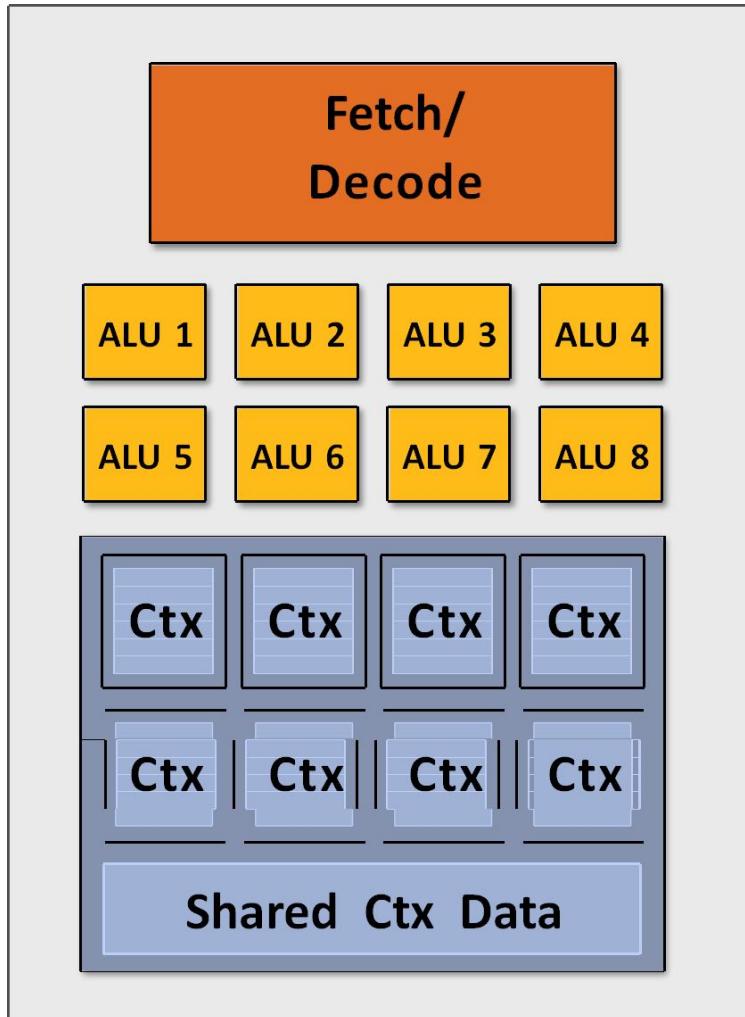
But ... many fragments should be able to share an instruction stream!

```
<diffuseShader>:  
sample r0, v4, t0, s0  
mul r3, v0, cb0[0]  
madd r3, v1, cb0[1], r3  
madd r3, v2, cb0[2], r3  
clmp r3, r3, l(0.0), l(1.0)  
  
mul o0, r0, r3  
mul o1, r1, r3  
mul o2, r2, r3  
mov o3, l(1.0)
```

Recall: simple processing core



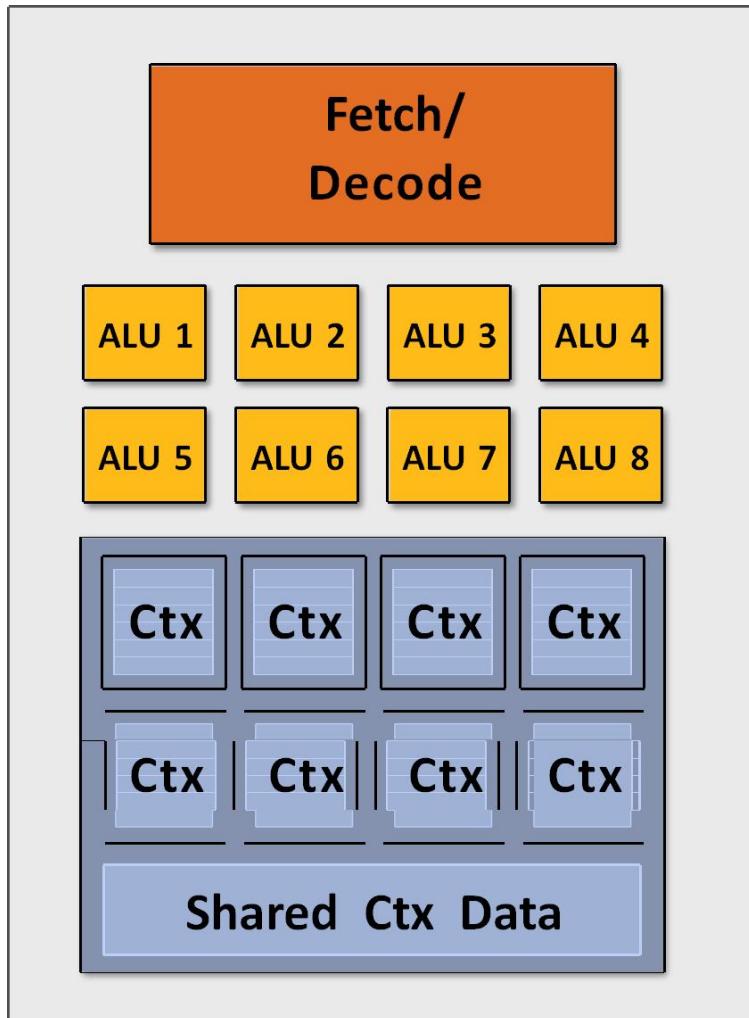
Add ALUs



Idea #2:
**Amortize cost/complexity of
managing an instruction stream
across many ALUs**

SIMD processing

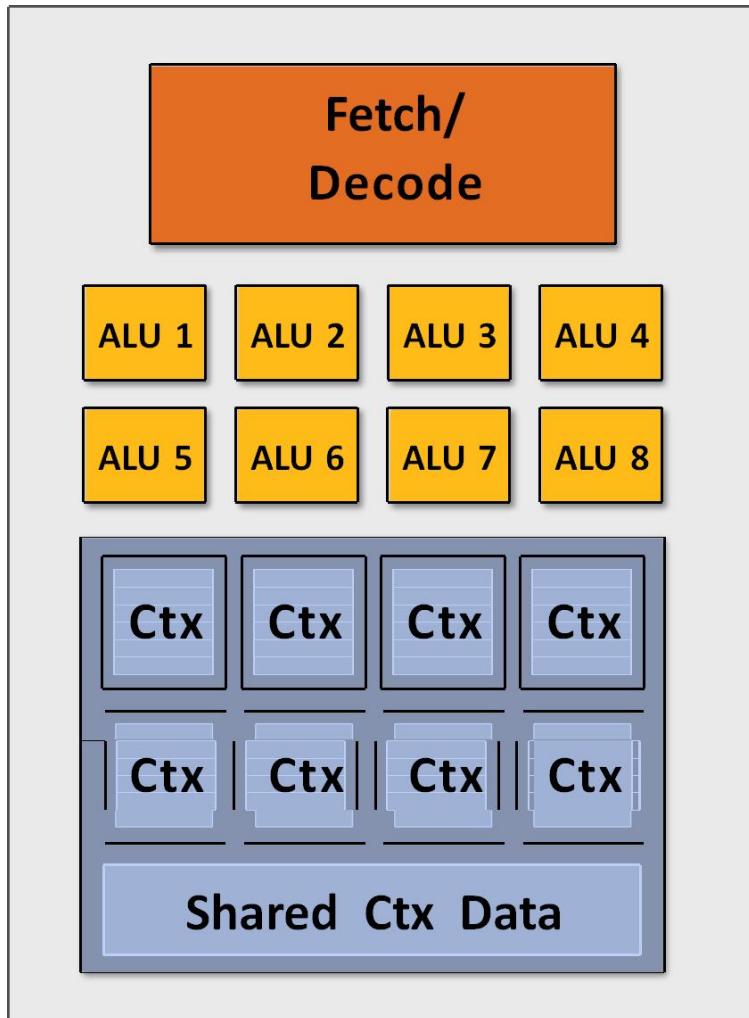
Modifying the shader



```
<diffuseShader>:  
sample r0, v4, t0, s0  
mul r3, v0, cb0[0]  
madd r3, v1, cb0[1], r3  
madd r3, v2, cb0[2], r3  
clmp r3, r3, 1(0.0), 1(1.0)  
  
mul o0, r0, r3  
mul o1, r1, r3  
mul o2, r2, r3  
mov o3, 1(1.0)
```

Original compiled shader:
Processes one fragment using scalar ops on scalar registers

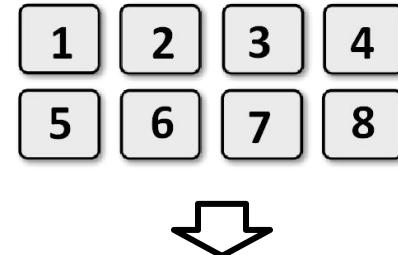
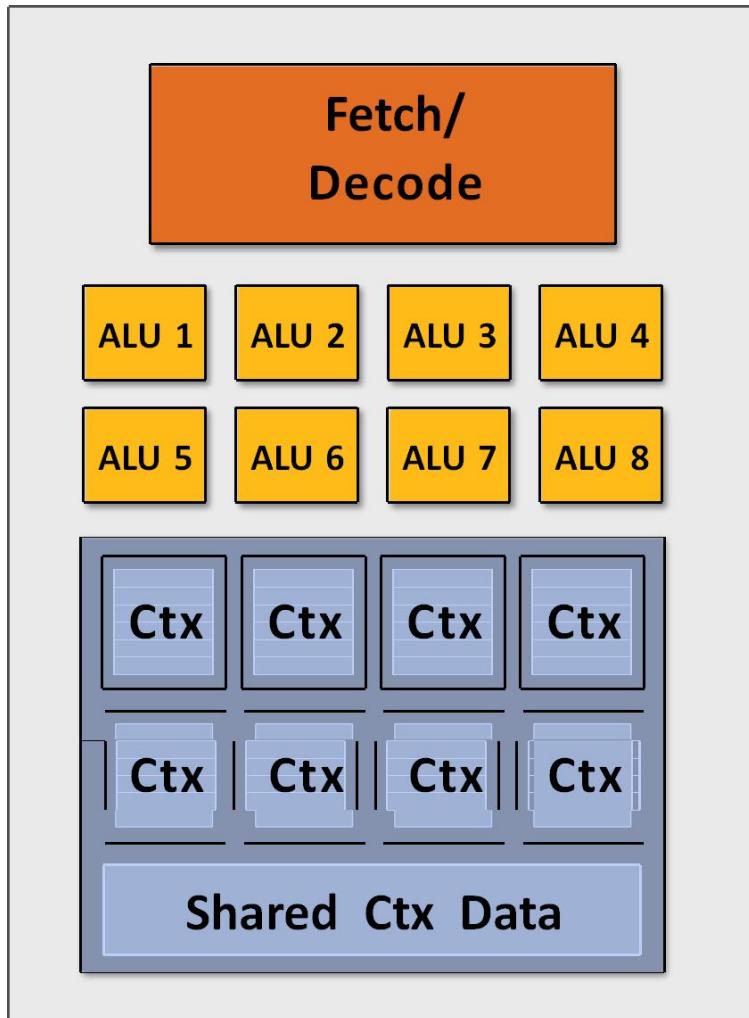
Modifying the shader



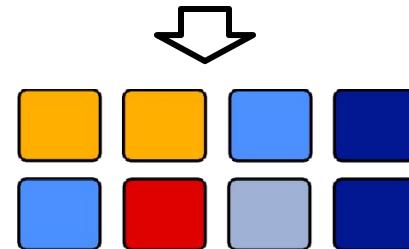
```
<VEC8_diffuseShader>:  
VEC8_sample vec_r0, vec_v4, t0, vec_s0  
VEC8_mul vec_r3, vec_v0, cb0[0]  
VEC8_madd vec_r3, vec_v1, cb0[1], vec_r3  
VEC8_madd vec_r3, vec_v2, cb0[2], vec_r3  
VEC8_clmp vec_r3, vec_r3, 1(0.0), 1(1.0)  
VEC8_mul vec_o0, vec_r0, vec_r3  
VEC8_mul vec_o1, vec_r1, vec_r3  
VEC8_mul vec_o2, vec_r2, vec_r3  
VEC8_mov o3, 1(1.0)
```

New compiled shader:
Processes eight fragments
using vector ops on vector
registers

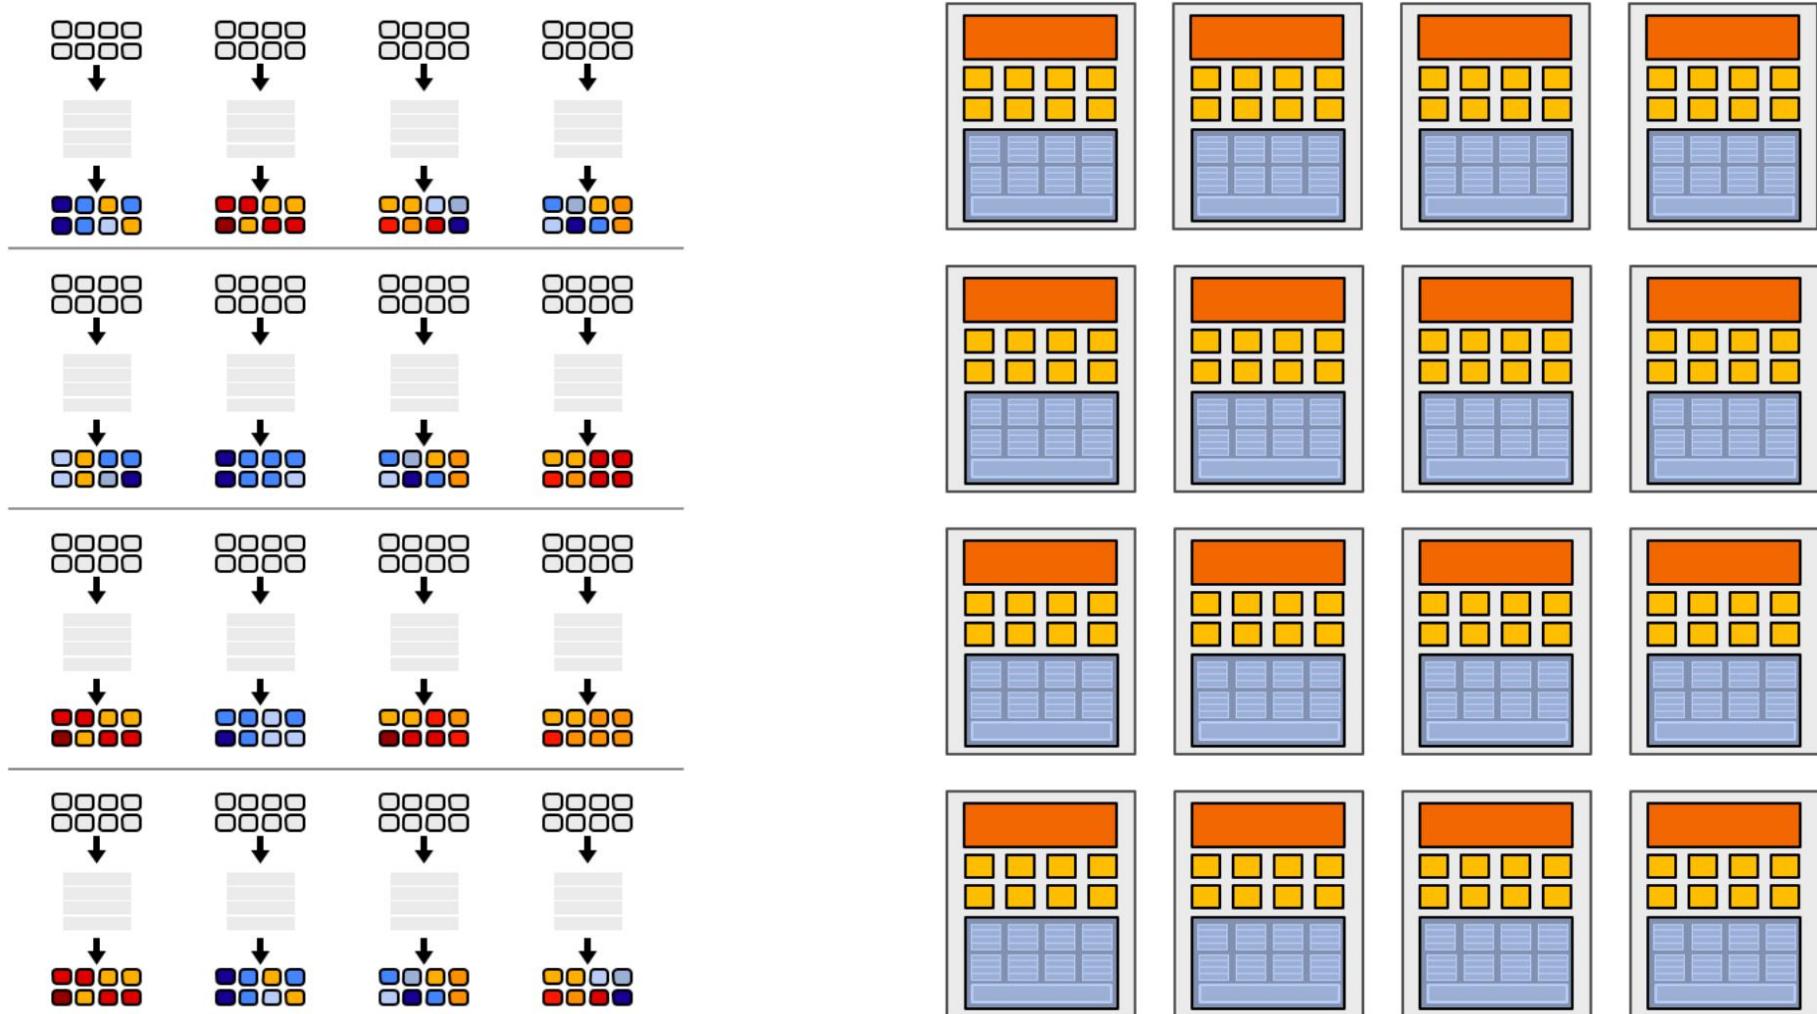
Modifying the shader



```
<VEC8_diffuseShader>:  
VEC8_sample vec_r0, vec_v4, t0, vec_s0  
VEC8_mul vec_r3, vec_v0, cb0[0]  
VEC8_madd vec_r3, vec_v1, cb0[1], vec_r3  
VEC8_madd vec_r3, vec_v2, cb0[2], vec_r3  
VEC8_clmp vec_r3, vec_r3, 1(0.0), 1(1.0)  
VEC8_mul vec_o0, vec_r0, vec_r3  
VEC8_mul vec_o1, vec_r1, vec_r3  
VEC8_mul vec_o2, vec_r2, vec_r3  
VEC8_mov o3, 1(1.0)
```

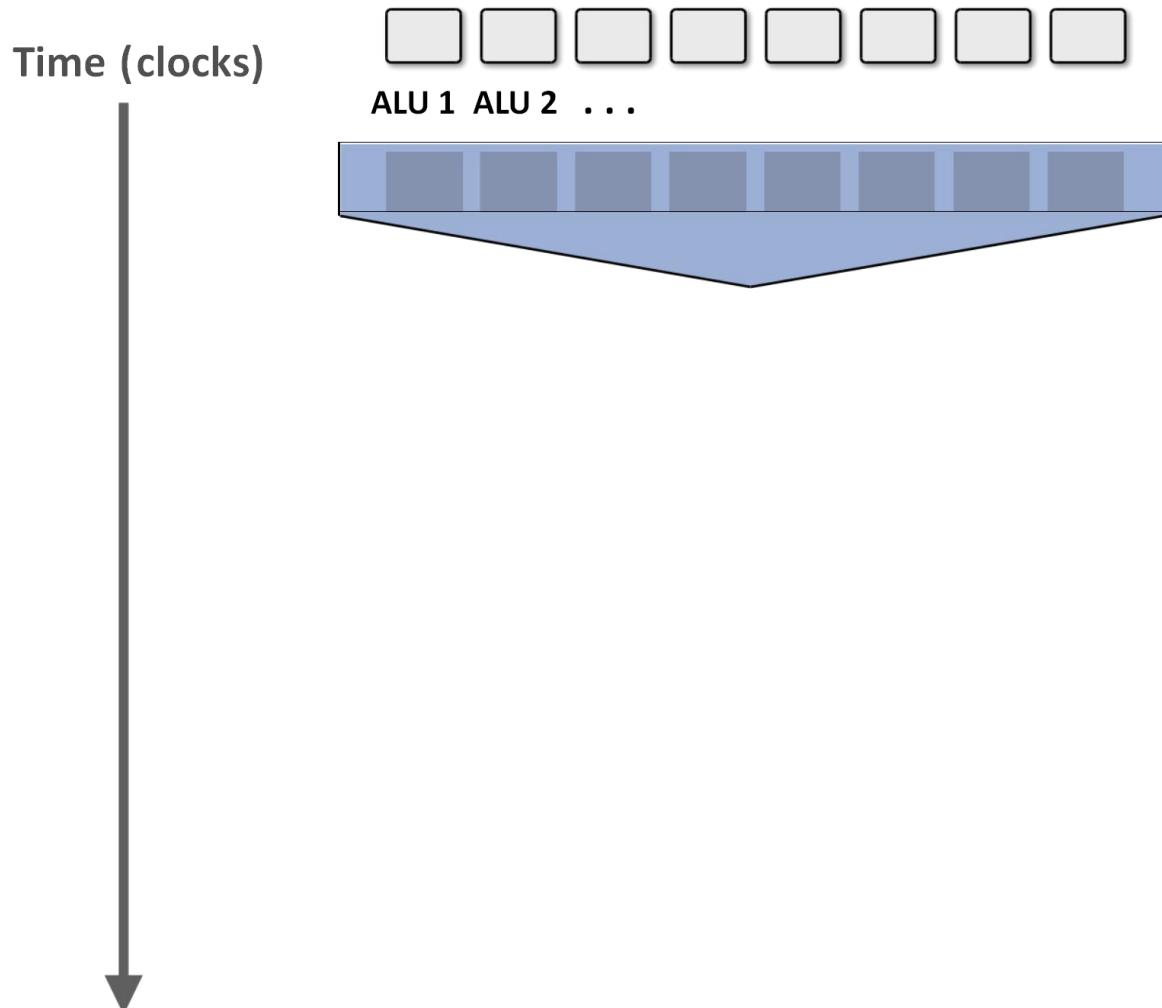


128 fragments in parallel



16 cores = 128 ALUs , 16 simultaneous instruction streams

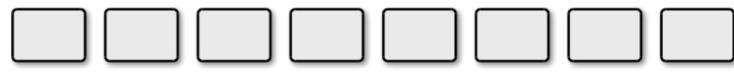
But what about branches?



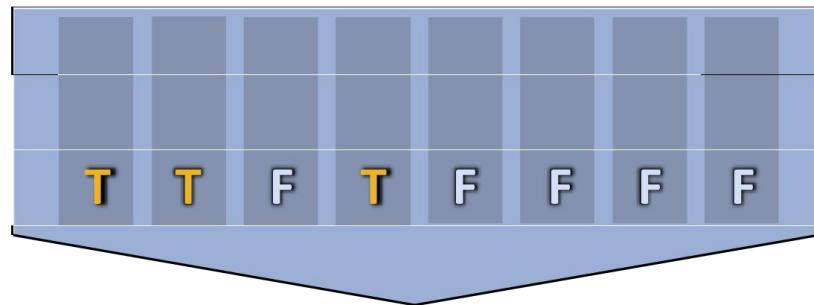
```
<unconditional  
 shader(code)>  
  
if((x(>(0)){  
    y(=(pow(x,(exp);  
    y(*=(Ks;  
    refl(=(y(+(Ka;  
}(else{  
    x(=(0;  
    refl(=(Ka;  
}  
  
<resume(unconditional  
 shader(code)>
```

But what about branches?

Time (clocks)



ALU 1 ALU 2 ...



< unconditional
shader(code)>

```
if(x > 0){  
    y = pow(x, exp);  
    y *= Ks;  
    refl = y+Ka;  
}else{  
    x=0;  
    refl=Ka;  
}
```

<resume(unconditional
shader(code)>

But what about branches?

Time (clocks)

1

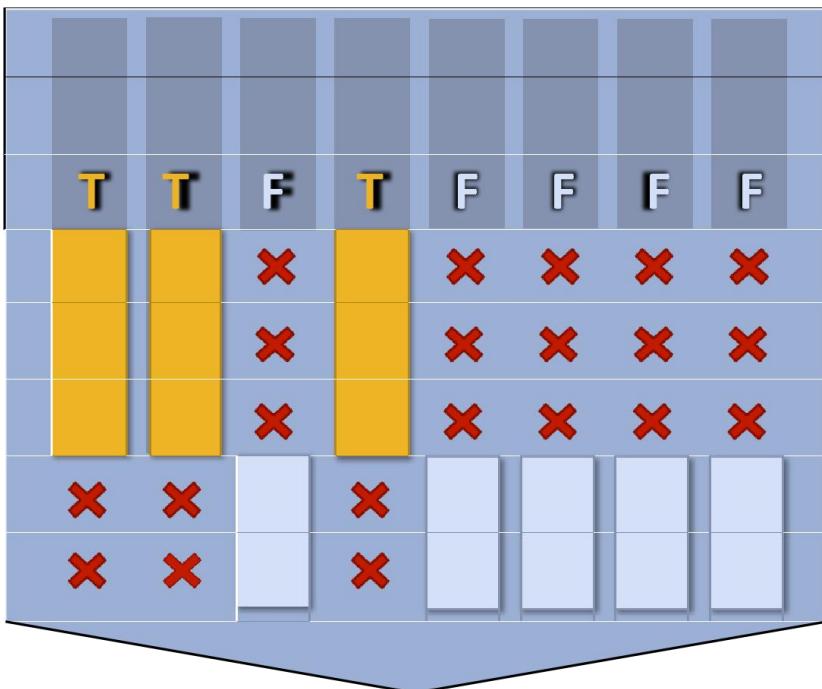
2



8

ALU 1 ALU 2 ...

ALU 8



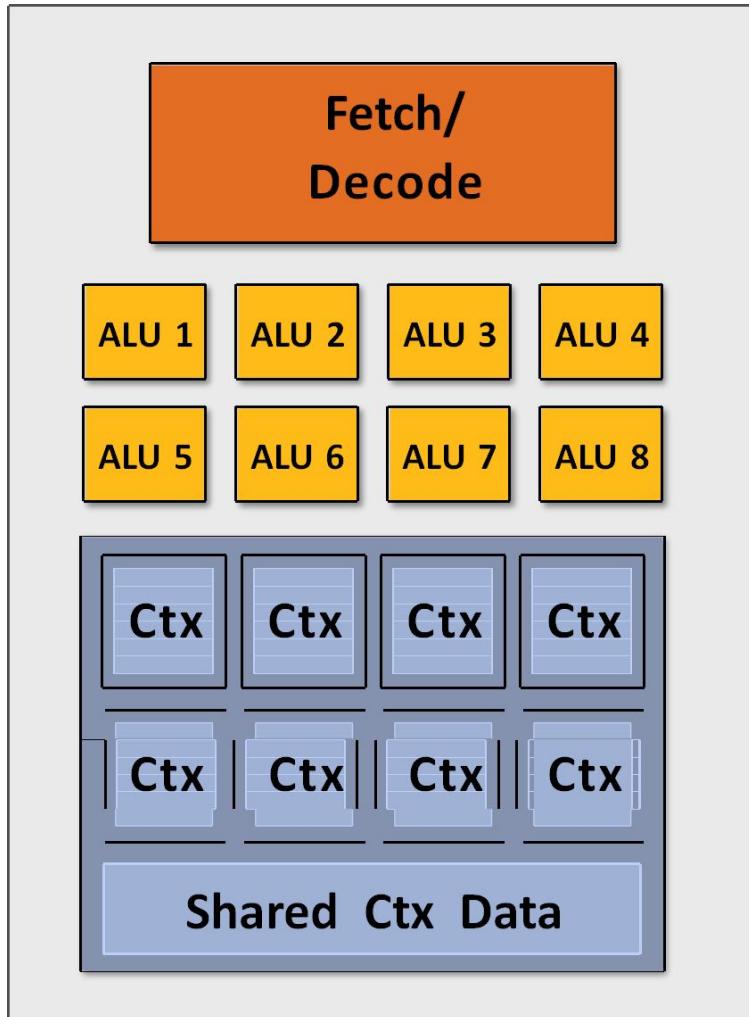
**Not all ALUs do useful work!
Worst case: 1/8 peak performance**

<unconditional
shader(code)>

```
if(x > 0){  
    y = pow(x, exp);  
    y *= Ks;  
    refl = y+Ka;  
} else{  
    x = 0;  
    refl = Ka;  
}
```

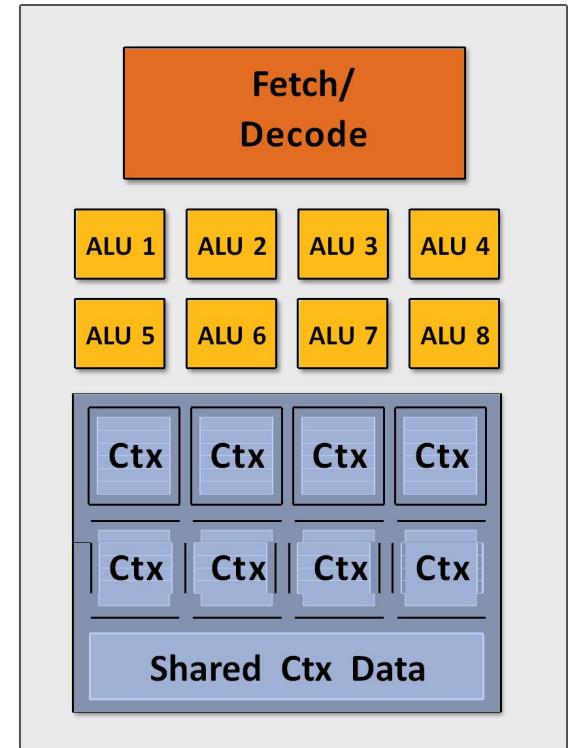
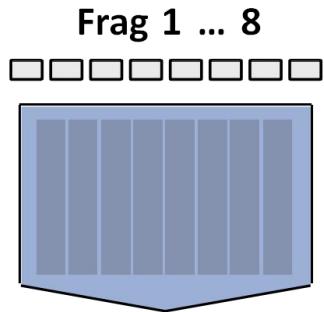
<resume(unconditional
shader(code)>

Hiding shader stalls

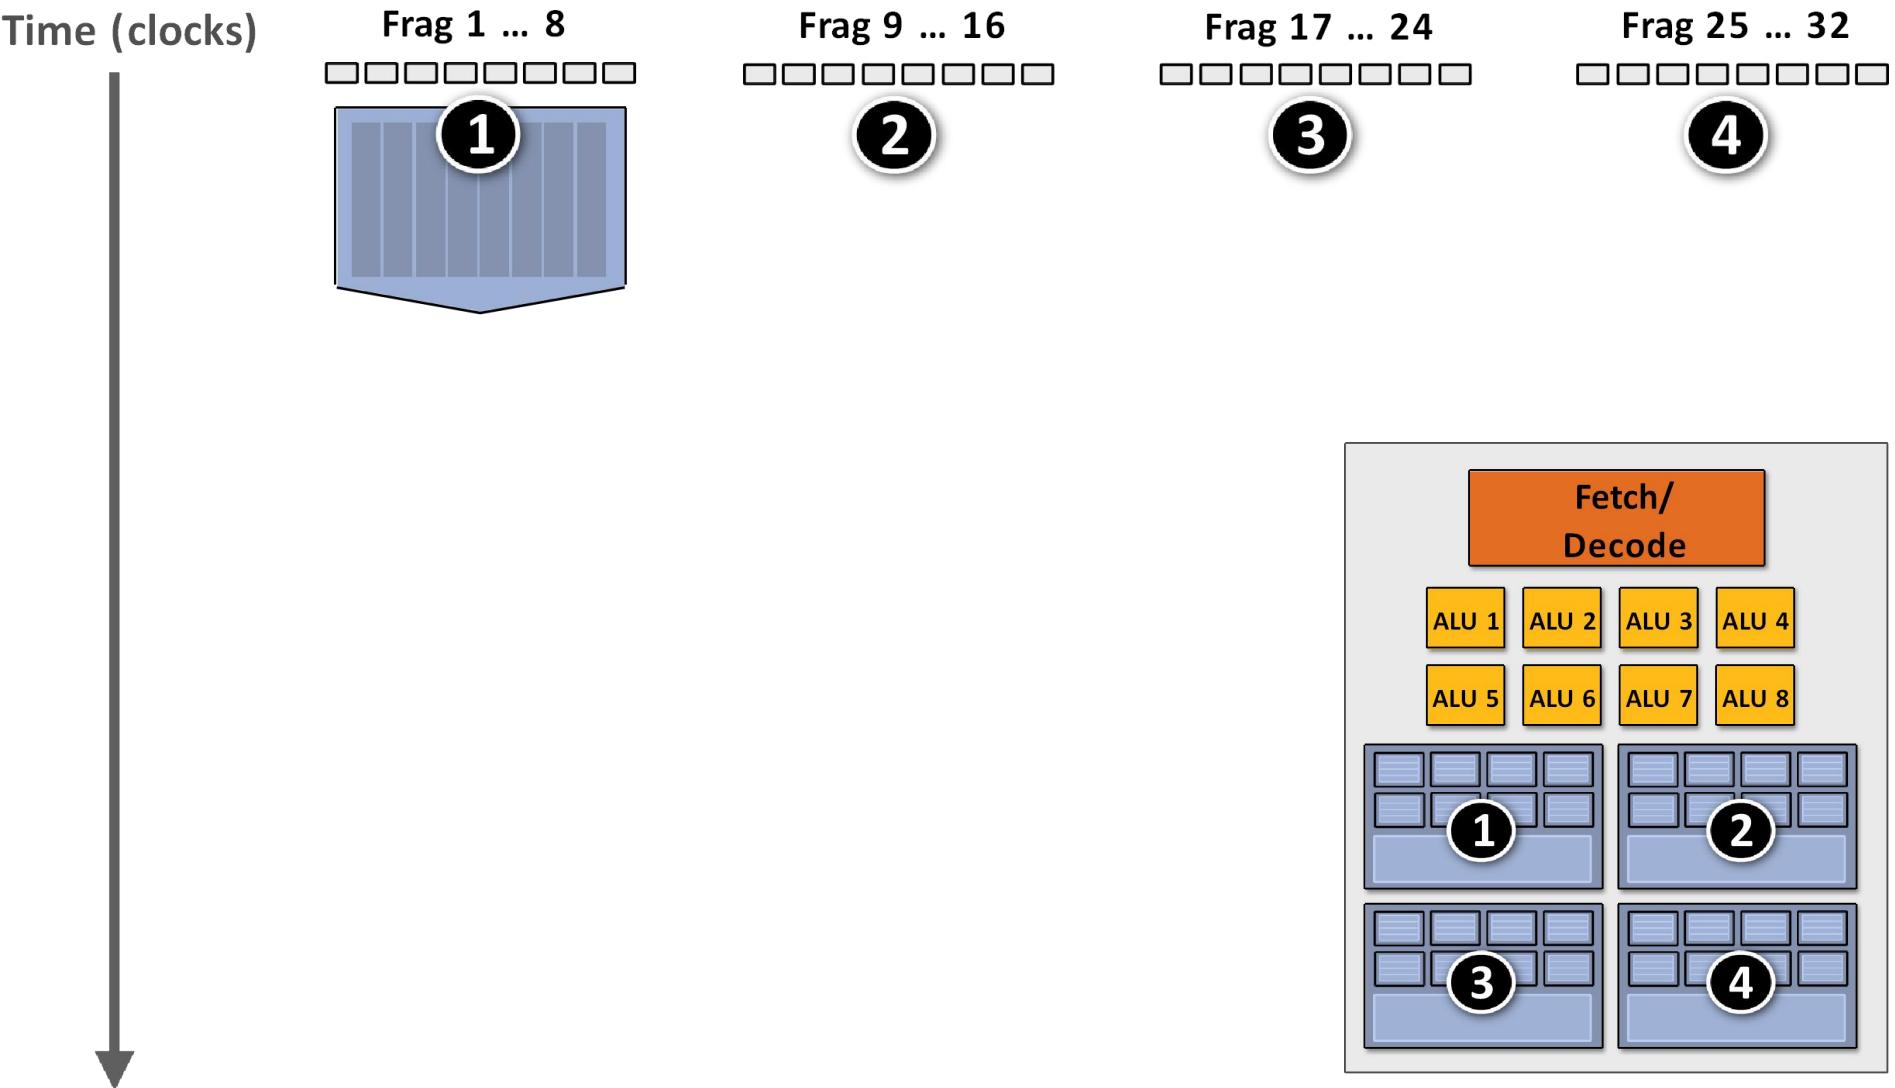


Idea #3:
Interleave processing of many fragments on a single core to avoid stalls caused by high latency operations.

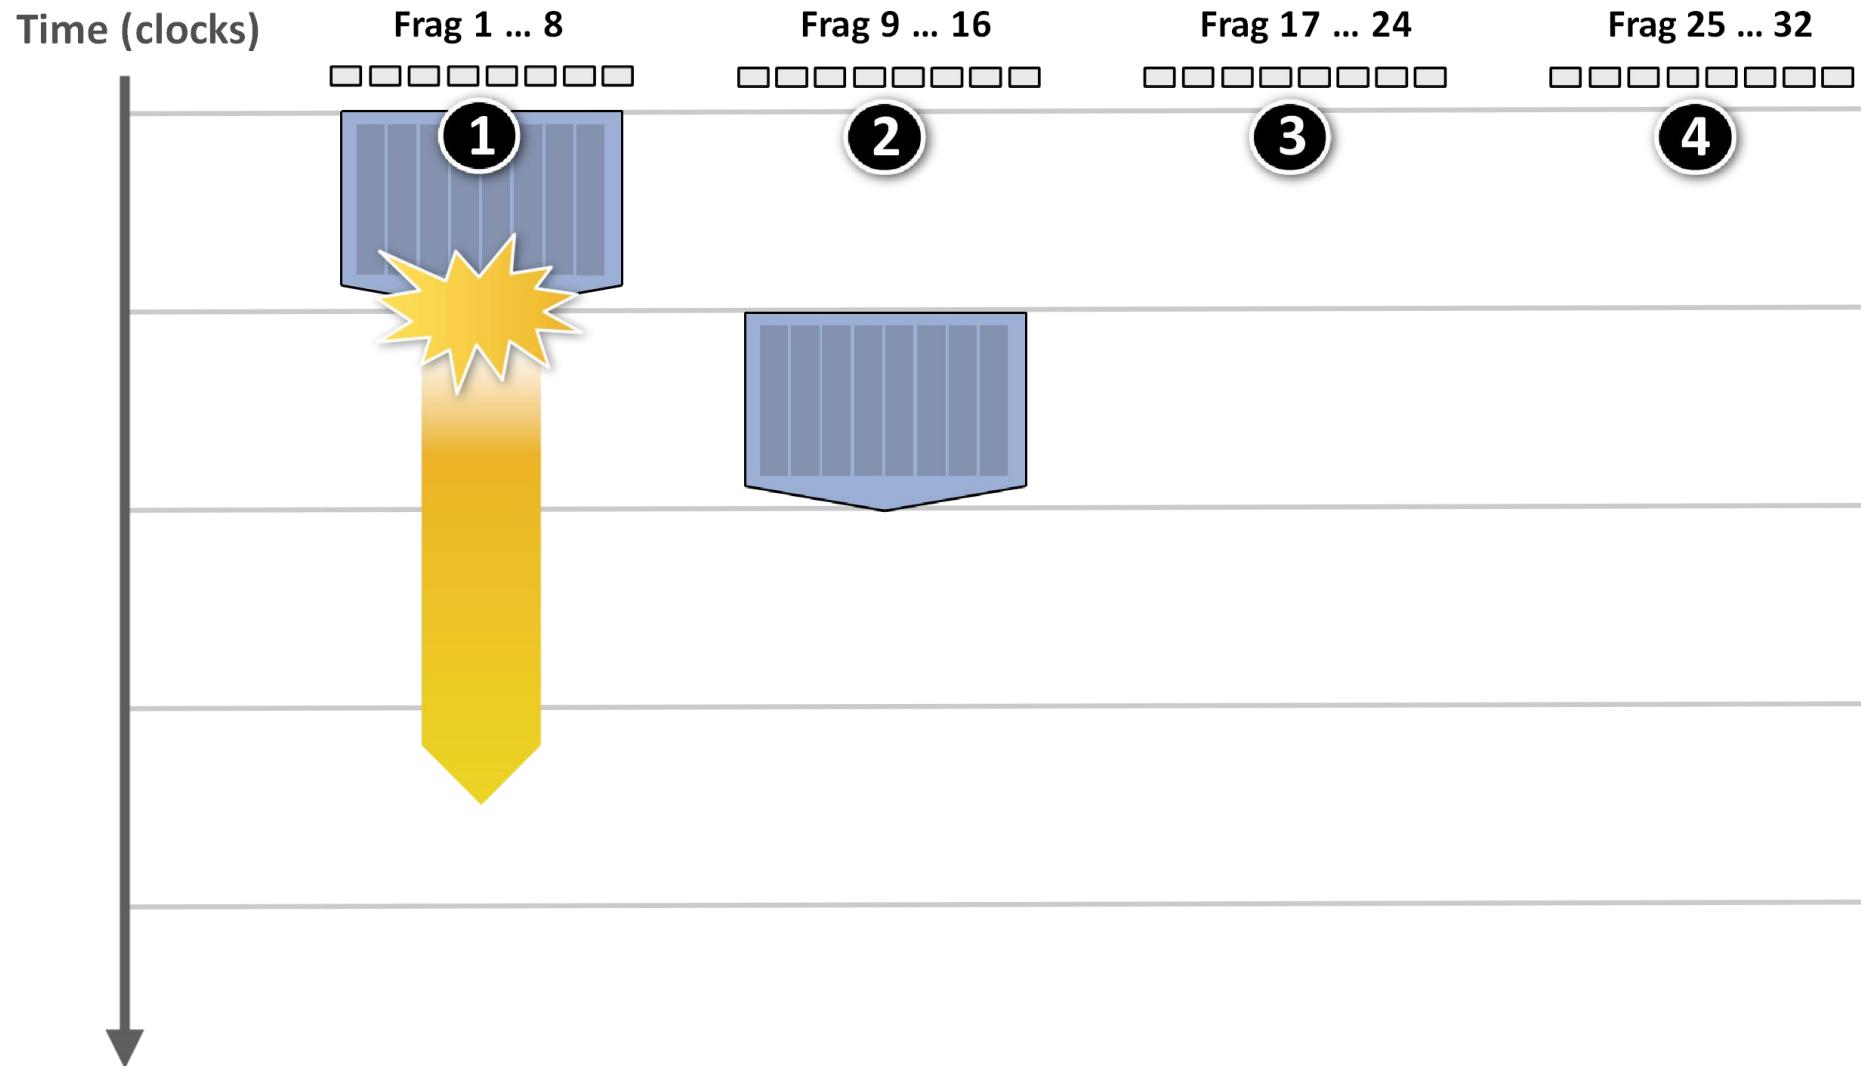
Hiding shader stalls



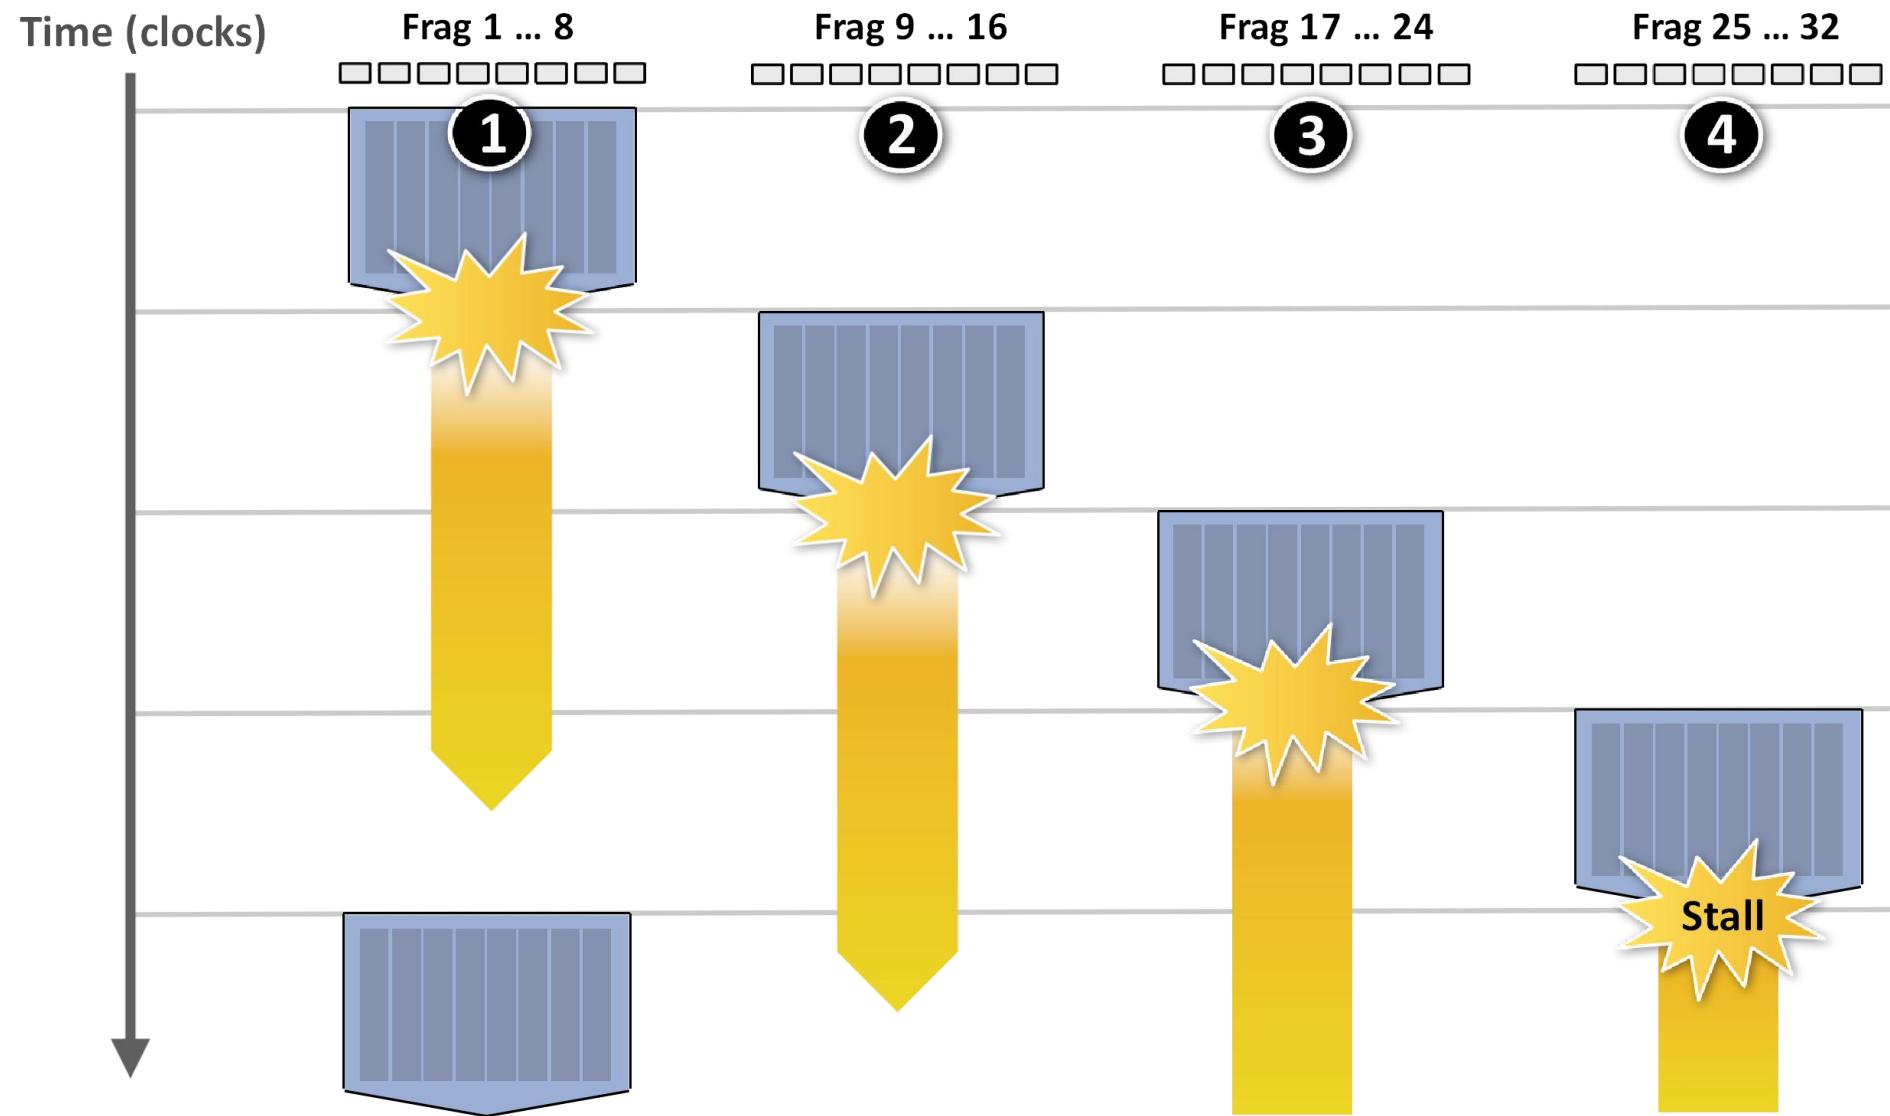
Hiding shader stalls



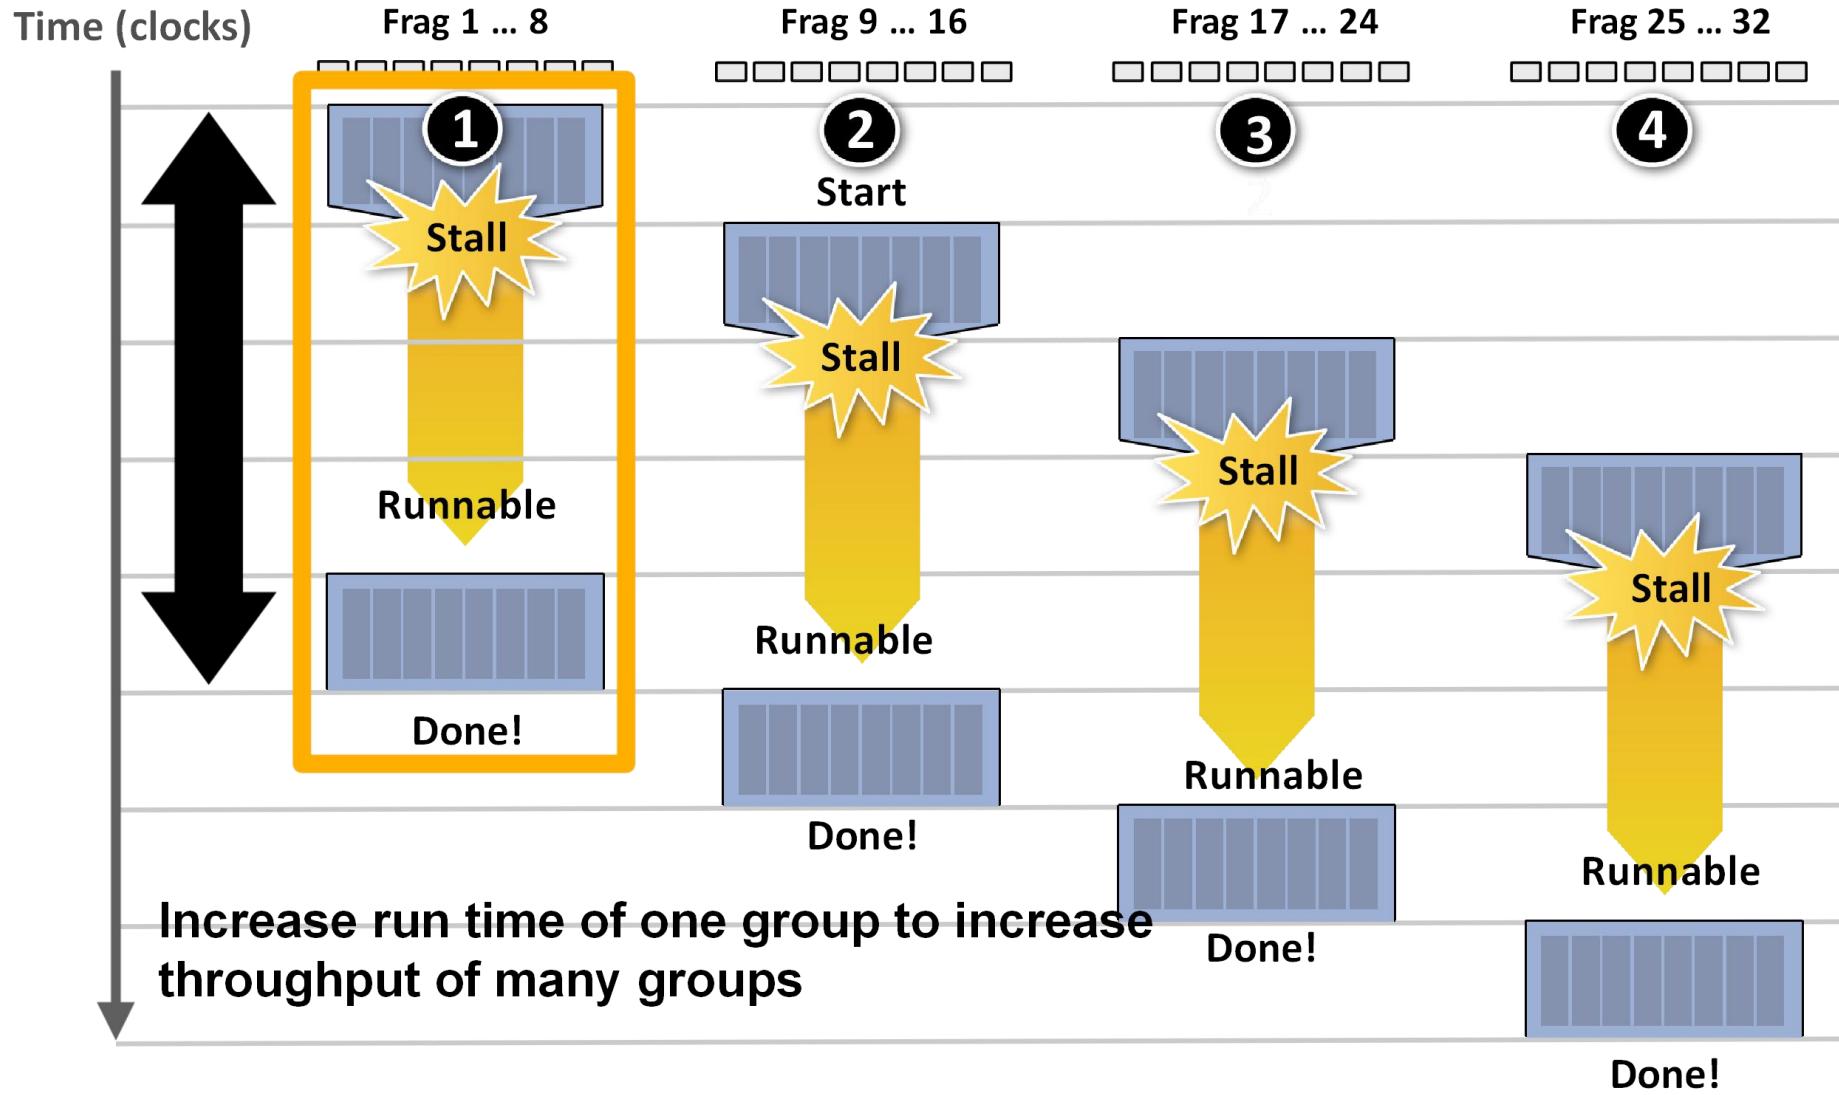
Hiding shader stalls



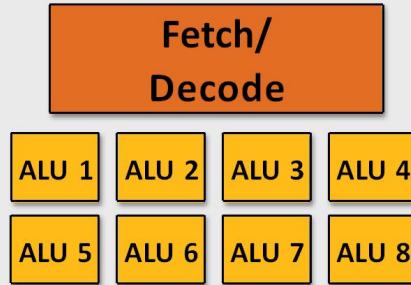
Hiding shader stalls



Throughput!

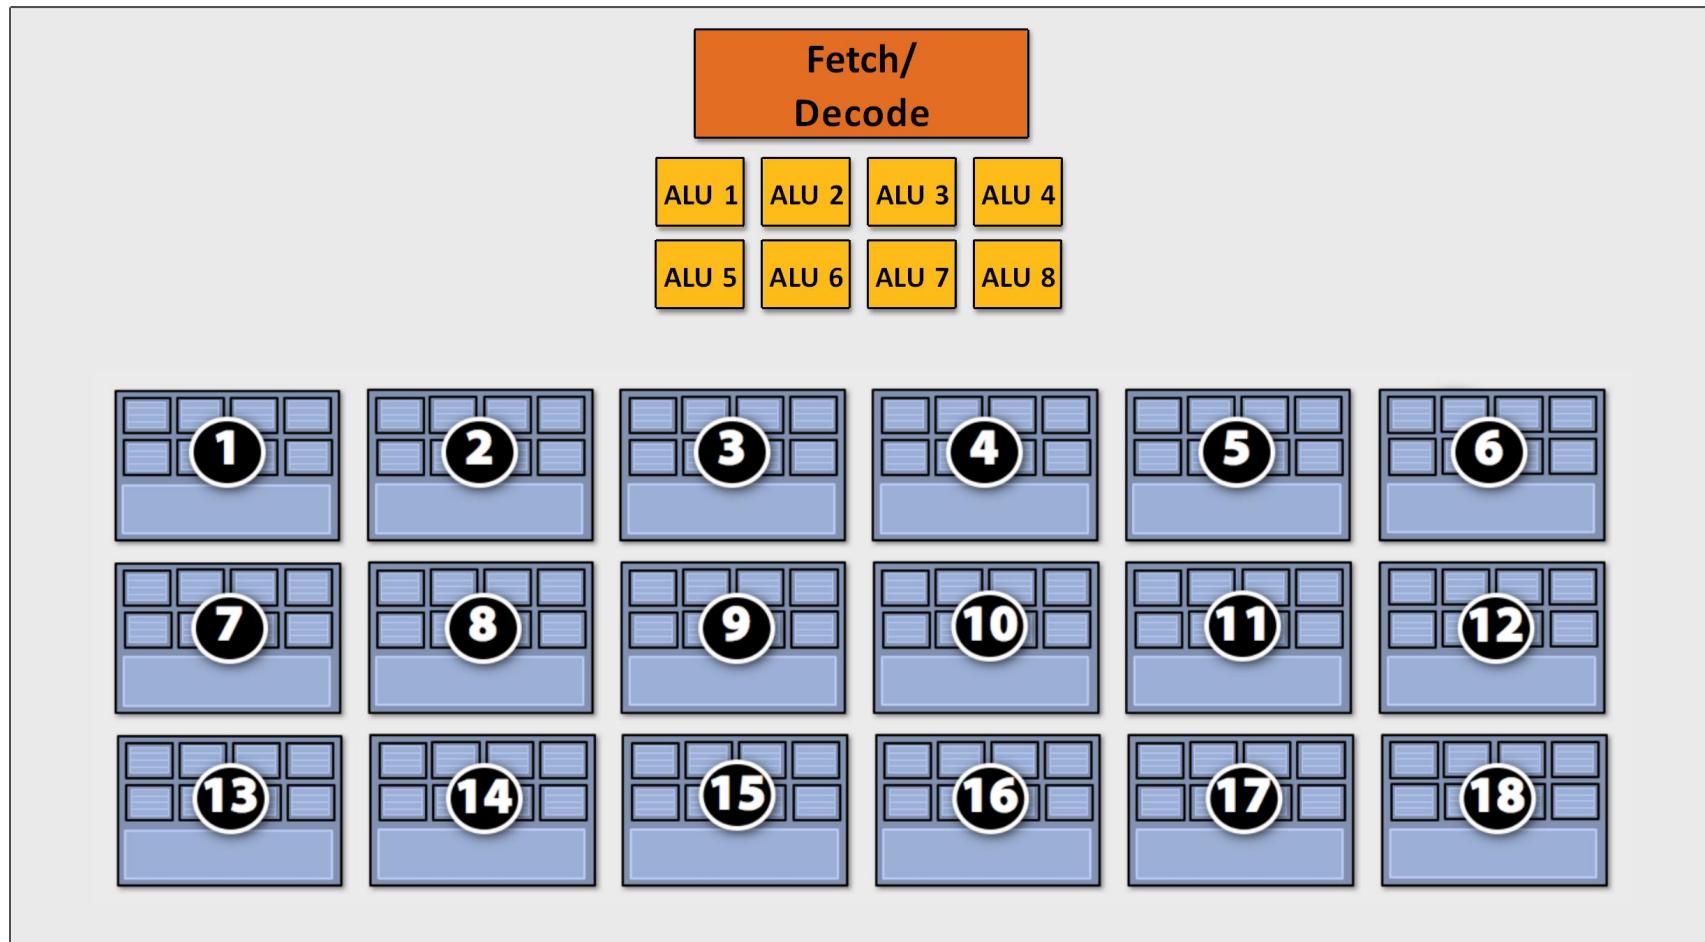


Storing contexts



**Pool of context storage
128 KB**

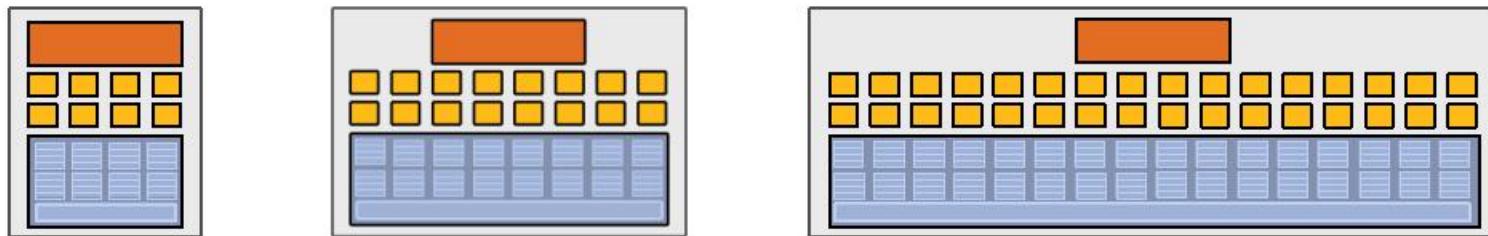
Eighteen small contexts



Clarification

SIMD processing does not imply SIMD instructions

- Option 1: explicit vector instructions
 - x86 SSE, Intel Larrabee
- Option 2: scalar instructions, implicit HW
 - HW determines instruction stream sharing across ALUs (amount of sharing hidden from software)
 - NVIDIA GeForce ("SIMT" warps), ATI Radeon architectures ("wavefronts")



In practice: 16 to 64 fragments share an instruction stream.

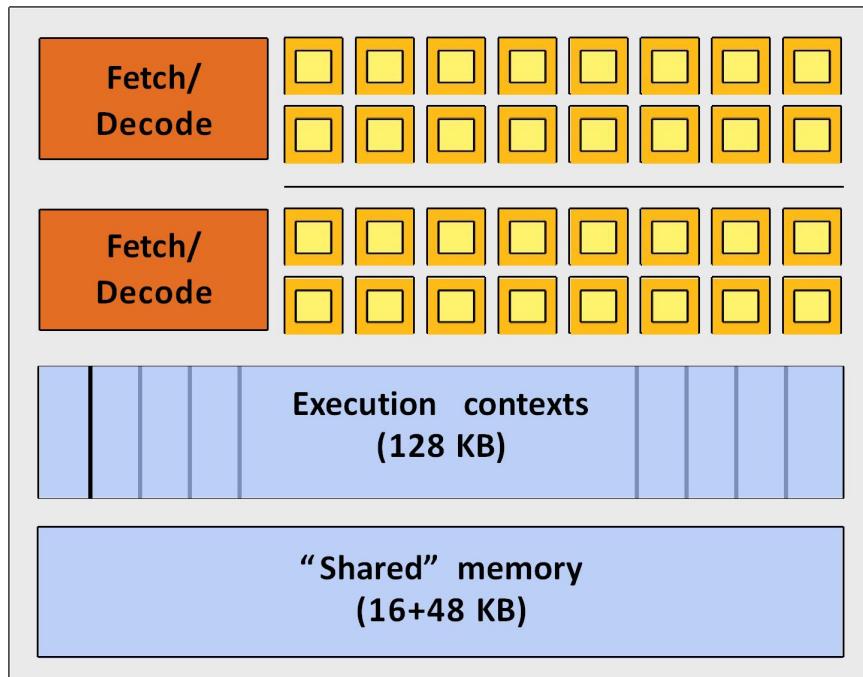
NVIDIA GeForce GTX 480 (Fermi)

Interleaving between contexts can be managed by hardware or software (or both!)

- NVIDIA-speak:
 - 480 stream processors (“CUDA cores”)
 - “SIMT execution”
- Generic speak:
 - 15 cores
 - 2 groups of 16 SIMD functional units per core



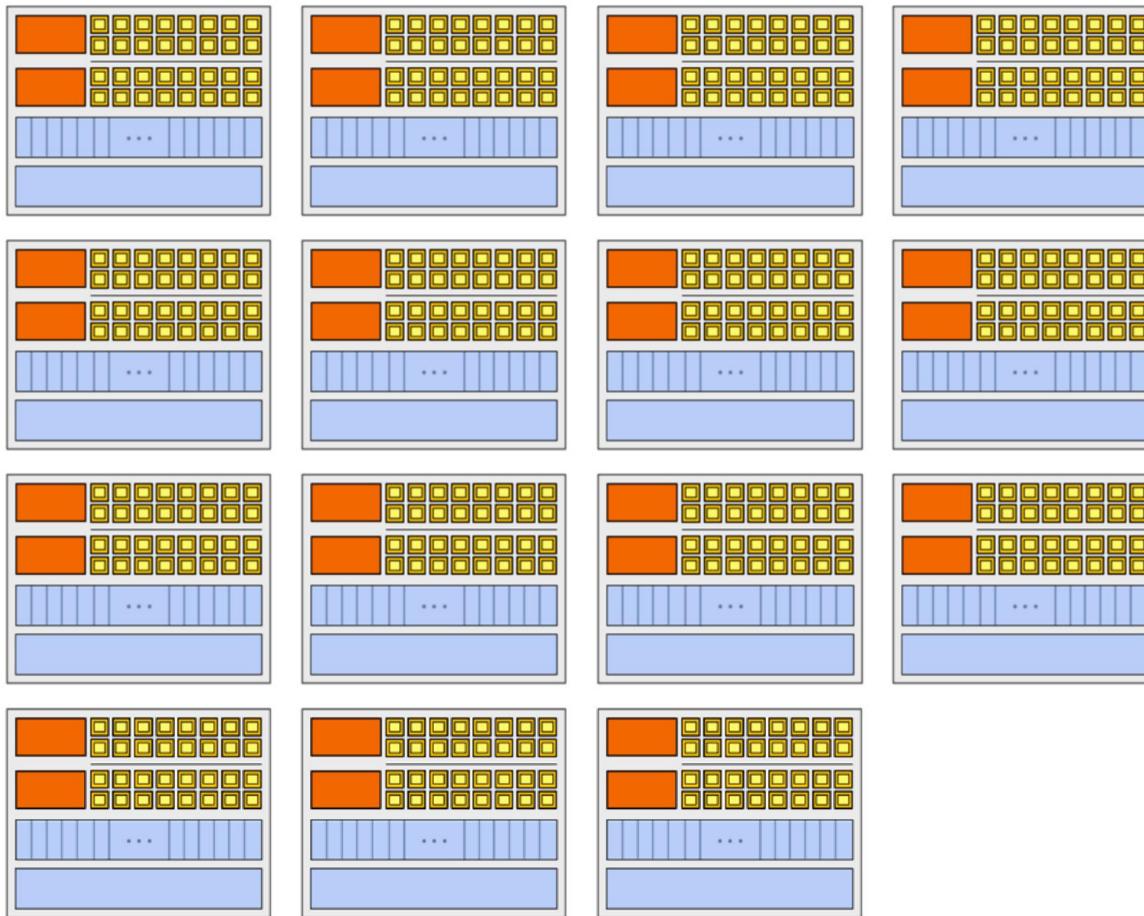
NVIDIA GeForce GTX 480 “SM”



= CUDA core
(1 MUL-ADD per clock)

- The SM contains **32** CUDA cores
- Up to **48** groups are simultaneously interleaved
- Up to **1536** individual contexts can be stored

NVIDIA GeForce GTX 480



There are **15** SMs on
the GTX 480:
That's **23,000**
fragments!
Or **23,000** CUDA
threads!

GPU Summary

- Three major ideas that all modern processors employ to varying degrees
 - Employ multiple processing cores
 - Simpler cores (embrace thread-level parallelism over instruction-level parallelism)
 - Amortize instruction stream processing over many ALUs (SIMD)
 - Increase compute capability with little extra cost
 - Use multi-threading to make more efficient use of processing resources (hide latencies, fill all available resources)
- Due to high arithmetic capability on modern chips, many parallel applications (on both CPUs and GPUs) are bandwidth bound
- GPU architectures use the same throughput computing ideas as CPUs: but GPUs push these concepts to extreme scales